

Pin the Name on the Donkey:

A Definitional Audit of the AI Industry's Catastrophic Failure to Decide What It Is Selling

Or: Is It a Bird? Is It a Plane? Is It Going to Do My Laundry and Achieve Consciousness While Defying Thermodynamics, or Is It Just a Very Expensive Autocomplete That Decided to Become a God Before Finishing Its Safety Testing?

PRIZES!
VALUATIONS!
VIBES!

SPIN THE
NARRATIVE!
WIN
MARKET SHARE!
NO REFUNDS!

WELCOME
TO THE
DEFINITIONAL
CARNIVAL
OF
CONFLICTING
CLAIMS!

AMBIGUITY
IS NOT A BUG.
AMBIGUITY IS
THE BUSINESS
MODEL.

TODAY'S
HIGHEST BID:
\$1.7T
AND STILL
NO IDEA WHAT
IT IS.

INVESTORS:
BRING CASH
&
SUSPEND
DISBELIEF!

ABSTRACT

This paper examines the Nomenclatural Chaos Engine (NCE) at the heart of the artificial intelligence industry—a structural phenomenon in which the same product is simultaneously marketed as: (a) a threat to all human employment, (b) a harmless autocomplete tool, (c) the last invention humanity will ever need to make, (d) just a calculator that learned to talk, and (e) 'honestly kind of vibes-based.' We introduce the central metaphor of *Pin the Name on the Donkey* (PNND)—a children's party game in which a blindfolded participant attempts to attach a label to a large animal they cannot see, have never seen, and are not certain exists.

We argue that this is the precise methodology by which the AI industry has conducted its public communications since approximately 2022. The consequences of PNND are severe: investors have committed trillions of dollars to something that cannot decide whether it is magic or statistics; regulators cannot regulate something that refuses to declare what it is; and consumers have developed theological opinions about a product whose creators cannot agree on its fundamental nature.

We further argue that this definitional anarchy—far from being accidental—is the core product strategy of the industry's leading institutions. Ambiguity is not a bug. Ambiguity is the business model. This will end in litigation. This paper documents exactly how, why, and with what level of irony.

Contents

2. Introduction / Brief History
3. Four Quadrants of Deliberate Confusion
4. The God Problem
5. The Litigation Horizon
6. Taxonomy of Claim Types
7. Strategic Ambiguity Endgame / Black Knight Phase
8. The Recursive Twist
9. Conclusion
10. Appendices



CLOSE
ENOUGH!

CHIEF
GROWTH
OFFICER

SAFETY TESTING (DRAFT)

RISK REGISTER (LOL)

MODEL CARD (MAYBE)

VALUE PROP (???)

I AGREE
TO DISAGREE
(DO NOT
QUOTE ME)

OBJECT UNKNOWN.
LABELS OPTIONAL.
LIABILITY INEVITABLE.

TAXONOMY OF CLAIM TYPES (PRELIMINARY)

CAPABILITY CLAIMS	IMPACT CLAIMS	IDENTITY CLAIMS	ASPIRATIONAL CLAIMS	HAND-WAVY CLAIMS
Can do X	Will do Y	Is Z	Might Become Q*	It's, Like, Vibes
☒	☒	☒	☐	☐

LEGAL
OBSERVER

DEFINITIONS
ARE HARD.
VALUATIONS
ARE EASY.
PROFITS
(SO FAR)

COMPLIANCE
WATCH

AMBIGUITY
AWARD
2024

MOST
STRATEGICALLY
VAGUE

Illustrated poster series based on the paper by Martin Ollett (2025),
Department of Highly Unapproved Methods.

1. Introduction: The Taxonomy of the Unknowable

1.1 The Problem of the Donkey

In the traditional party game from which this paper derives its title, success requires only two things:

1. A label
2. A donkey

The AI industry has managed to produce neither with any consistency.

Consider the following statements, all made by credentialed senior executives of major AI companies within the same eighteen-month period:

- "This is just a stochastic parrot." (academic critics, implying it is nothing)
- "We have achieved AGI." (Jensen Huang, Nvidia CEO, implying it is everything)
- "50% of all entry-level lawyers will be wiped out within 1-5 years." (Dario Amodei, Anthropic CEO, implying it is a scythe)
- "The real breakthrough will not be better reasoning, but total memory." (Sam Altman, OpenAI CEO, implying the current products are incomplete)
- "AGI is becoming outdated because, in some ways, we have already surpassed it." (Daniela Amodei, Anthropic President, implying the concept is already obsolete)
- "It is amazing—one person with chatbots can now act with the power of a research institute." (Sam Altman, the same Sam Altman, moments before calling it a company)

We propose a formal name for this phenomenon: **The Definitional Fugue State (DFS)**—a cognitive and communicative condition in which an institution is unable to commit to a single description of its own product for longer than one investor call.

The DFS is not unique to AI. It has historical precedents. However, no prior industry has achieved a DFS of such structural depth, such philosophical ambition, and such exposure to class action litigation.

1.2 The Scope of the Problem

To understand the scale of the definitional crisis, consider the following taxonomy of names that have been applied—sincerely, by serious people—to the same underlying technology within the same era:

	Names suggesting it is a tool:	• Assistant • Copilot • Helper • Autocomplete
	Names suggesting it is a colleague:	• Agent • Worker • Team member • Reasoning model
	Names suggesting it is a deity:	• AGI (Artificial General Intelligence) • Superintelligence • The Singularity • "Something that surpasses human-level performance" (which is either a tool or a deity, depending on whether your lawyer is present)
	Names suggesting it is a threat:	• Existential risk • Terminator scenario • Intelligence explosion • "The last invention humanity will ever need to make" (a phrase used both admirably and in terror, often in the same sentence)
	Names suggesting it is a scam:	• Stochastic parrot • Fancy autocomplete • JPEG of the internet • "A distributed hallucination with a GPU bill"
	Names used when the definition must not be specified for legal reasons:	• "Transformative technology" • "Paradigm shift" • "The thing that is definitely going to change everything" (trademark pending)

The existence of all of these categories simultaneously is not a sign of rich complexity. It is a sign that the donkey has no idea where its own tail is, and neither does anyone who is being paid to find it.

2. A Brief History of the Blindfold

2.1 The Pre-Blindfold Era (2015–2020):

At Least They Were Honest About the Donkey

The early years of the modern AI era were characterized by a form of definitional humility that has since been entirely abandoned.

Researchers said things like:

- "These models can do X, but fail on Y."
- "We don't know why it works."
- "It's impressive but it's not intelligence."

This was honest. It was also commercially useless.

The Foundational Marketing Insight (FMI):

You cannot raise \$100 billion by telling people you have built something that occasionally generates plausible sentences and fails at third-grade arithmetic.

Therefore, the industry required a blindfold. The blindfold was manufactured from several materials:

- 1 Genuine uncertainty (nobody actually knows what these systems are doing internally)
- 2 Strategic ambiguity (saying less makes fundraising easier)
- 3 Narrative inflation (each demo needed to be more revelatory than the last)
- 4 Philosopher laundering (importing the vocabulary of consciousness studies to describe a matrix multiplication)

By 2022, the blindfold was firmly in place. The game had begun.

TODAY'S MATERIALS OF THE BLINDFOLD

GENUINE UNCERTAINTY

STRATEGIC AMBIGUITY

NARRATIVE INFLATION

PHILOSOPHER LAUNDERING

AGI?
TOOL?
THREAT?
GODDESS OF PRODUCTIVITY?
I WAS PROMISED CARROTS.*



2.2 The ChatGPT Moment (November 2022): The Blindfold Becomes Mandatory

The release of ChatGPT to the public was the moment the definitional game became industry-wide. Before ChatGPT, "AI" was a specialist conversation. After ChatGPT, it became a theology.

Within sixty days of release, the following people had publicly declared their opinion on What It Is:

- Philosophers who had never used it
- Politicians who had never used it
- Journalists who had used it to write their articles about it
- Parents who were worried their children would use it for homework (correct)
- Executives who were worried their employees would use it for their jobs (also correct)
- People who were worried it would become conscious and take revenge (pending)
- Sam Altman, who said it was "a tool," "transformative," "potentially dangerous," "exciting," "not AGI," and "maybe AGI" within the same week

The ChatGPT Definitional Explosion (CDE) created a structural condition from which the industry has never recovered: everyone now had an opinion on What AI Is, and none of those opinions were compatible, and all of them were being used to make financial, regulatory, and existential decisions.

The party game had become a planetary game.
The donkey was now the size of an economy.

DEFINITIONS
NOT FOUND
SINCE 2015

STILL LOOKING

SEARCH STATUS:
ONGOING
AND BILLABLE

3. The Four Quadrants of Deliberate Confusion

3.1 The Strategic Ambiguity Matrix (SAM)

We propose that AI communications can be mapped onto a two-by-two matrix defined by two axes:

Axis 1: Capability Claims

- **Modest:** "It's a tool that helps you write emails faster."
- **Maximalist:** "It will end the need for human cognitive labor within a decade."

Axis 2: Risk Acknowledgement

- **Minimal:** "Current risks are manageable and overblown."
- **Catastrophic:** "This may be the most dangerous technology in human history."

The four resulting quadrants produce four distinct product pitches, all of which are deployed by the same companies at different moments depending on the audience:

Quadrant 1 — Modest Capability + Minimal Risk:

(Used for: regulators, nervous employees, consumer PR)
'It's just software. It helps you be more productive. There's nothing to worry about.'

Quadrant 2 — Maximalist Capability + Minimal Risk:

(Used for: investors, press releases, conference keynotes)
'This will transform every industry on earth and create unimaginable prosperity, and we have it completely under control.'

Quadrant 3 — Modest Capability + Catastrophic Risk:

(Used for: never—this quadrant is commercially useless and rhetorically incoherent)
'It's just autocomplete, but it might destroy civilisation.'
(No one says this, for obvious reasons.)

Quadrant 4 — Maximalist Capability + Catastrophic Risk:

(Used for: safety fundraising, regulatory lobbying, AI safety conferences, and any situation in which the speaker wants to appear both powerful and responsible)
'We are building something that may be the most consequential technology in human history, which is why you should give us more money and fewer regulations so we can continue building it responsibly.'

The master move of the AI industry is the ability to occupy all four quadrants simultaneously while maintaining the appearance of coherence. This is achieved through Selective Audience Specification (SAS)—the practice of never giving the same speech twice to the same type of person.

3.2 The Job Destruction Oscillation

Nowhere is the PNND dynamic more clearly visible than in the industry's communications about employment. The timeline of "will AI take jobs" messaging is as follows:

- **Phase 1 (2016–2020):** "AI will take some jobs but create new ones. History shows technology creates more jobs than it destroys."
- **Phase 2 (2021–2022):** "AI will augment human workers, not replace them. The human in the loop is essential."
- **Phase 3 (2023–2024):** "50% of entry-level lawyers, consultants, and finance professionals will be completely wiped out within 1–5 years." (Dario Amodei on Fox News)
- **Phase 4 (minutes later):** "AI empowers individuals to act with the power of a research institute." (Sam Altman, apparently about the same technology)
- **Phase 5 (when pressed):** "We can't predict with certainty how the labour market will evolve."

This is not a debate. This is a single institution's internal communications strategy, deployed across different audiences with different emotional needs. The factory workers need Phase 1 and 2. The investors need Phase 3 as a threat and Phase 4 as an opportunity. The regulators need Phase 5. The donkey is simultaneously destroying all the jobs and creating miraculous new jobs for the people whose jobs it just destroyed.

The Employment Paradox Resolution Theorem (EPRT):

If the same technology is simultaneously the greatest threat to employment in human history AND the greatest empowerment tool for individuals in human history, it is not being described accurately. It is being sold, to multiple parties, in multiple directions, with no regard for internal consistency. This is the definitional crisis in its purest form.

THE STRATEGIC AMBIGUITY MATRIX (SAM)

AXIS 1: CAPABILITY CLAIMS

MODEST

MAXIMALIST

AXIS 2: RISK ACKNOWLEDGEMENT

MINIMAL

CATASTROPHIC

1 **MODEST CAPABILITY + MINIMAL RISK**
(Used for: regulators, nervous employees, consumer PR)
'It's just software. It helps you be more productive. There's nothing to worry about.'

2 **MAXIMALIST CAPABILITY + MINIMAL RISK**
(Used for: investors, press releases, conference keynotes)
'This will transform every industry on earth and create unimaginable prosperity, and we have it completely under control.'

3 **MODEST CAPABILITY + CATASTROPHIC RISK**
(Used for: never—this quadrant is commercially useless and rhetorically incoherent)
'It's just autocomplete, but it might destroy civilisation.'
(No one says this, for obvious reasons.)

4 **MAXIMALIST CAPABILITY + CATASTROPHIC RISK**
(Used for: safety fundraising, regulatory lobbying, AI safety conferences, and any situation in which the speaker wants to appear both powerful and responsible)
'We are building something that may be the most consequential technology in human history, which is why you should give us more money and fewer regulations so we can continue building it responsibly.'

SELECTIVE AUDIENCE SPECIFICATION (SAS)



ONE COMPANY. FOUR MASKS. INFINITE CONTRADICTIONS.



The master move is not lying. It is segmented truth deployment.
Four messages. Four audiences. One donkey.

4. The God Problem: When San Francisco Decides to Play Zeus

4.1 The Architecture of Divine Claims

The AI industry's marketing has developed an interesting structural relationship with theology. The claims made about AI capabilities have—without explicit intention—followed the exact developmental arc of a religion:

Stage 1: The Prophecy

"We believe it is possible to build general artificial intelligence that benefits all of humanity." (The founding scripture)

Stage 2: The Miracles

"It writes! It codes! It passes the bar exam! It creates images from text! It diagnosed my rare condition!" (The signs and wonders)

Stage 3: The Eschatology

"The Singularity is near. AGI will arrive within [3–5 years]. This may be the last invention humans ever need to make." (The apocalypse narrative, which is either a threat or a promise depending on the congregation)

Stage 4: The Schism

Effective Altruists vs. Accelerationists vs. Doomers vs. e/acc vs. rationalists vs. "it's just autocomplete" empiricists. (The denomination wars)

Stage 5: The Inquisition

"Anyone who says the current models can't reason has simply not used them correctly." (The heresy detection apparatus)

Stage 6: The Tithe

\$7,000,000,000 infrastructure spend request.
(The offering plate, at scale)

Every stage of this arc produces a new name for the donkey. Every name requires a new pitch deck. Every pitch deck generates a new valuation. Every valuation generates a new commitment that must eventually be reconciled with the actual performance of the product. This reconciliation is called litigation.

4.2 The Thermodynamics Problem

At some point, somebody is going to notice that thermodynamics exists. The AI industry's implicit promise is that scaling will continue until the models become godlike.

This is the Infinite Scaling Hypothesis (ISH): more compute → smarter model → eventually AGI → eventually superintelligence → everything changes.

The ISH has several problems:

Problem 1: The Data Wall

You can only train on the internet once. The internet has already been eaten. The models are now being trained on synthetic data—i.e., on the outputs of previous models. This is the informational equivalent of a snake eating its own tail, which is nutritious in exactly the way you'd expect.

Problem 2: The Efficiency Paradox

DeepSeek produced GPT-4-level performance at approximately 1/17,000th of the training cost. This either means: (a) the scaling hypothesis was wildly wrong about the required compute, or (b) DeepSeek has achieved something miraculous at \$6 million that OpenAI needed \$100 billion to approach. Either interpretation is catastrophic for the "we need \$7 trillion" argument.

Problem 3: The Benchmark Treadmill

Every time a model achieves human-level performance on a benchmark, the benchmark is immediately declared "too easy" and replaced with a harder one. This is not scientific progress. This is a definitional game in which "human-level intelligence" is a finish line that moves whenever the runner approaches it.

Problem 4: The Actual Thermodynamics

The amount of energy required to run frontier AI models is measurable, real, and limited by physics. The amount of intelligence these models produce is not clearly correlated with energy consumption in the way the ISH requires. Physics does not care about your Series D.

4.3 The Sam Altman Pipe Dream Formalization

We propose a formal metric for what might be called the San Francisco Eschatological Index (SFEI)—a measure of how far a product's marketing claims have departed from its physical constraints.

$$\text{SFEI} = (\text{Claimed Capability}) \times (\text{Timeline Urgency}) / (\text{Physical Realizability})$$

For AGI:

- Claimed Capability: total replacement of human cognition (∞)
- Timeline Urgency: "within our lifetimes" (high)
- Physical Realizability: unknown, but constrained by energy, data, and architecture limits (low)

$$\text{SFEI} = \infty \times \text{high} / \text{low} = \text{undefined}$$

An undefined SFEI is the formal definition of a Pipe Dream (PD)—a claim that is structurally incapable of falsification because its terms are not operationally defined, its timeline is conveniently flexible, and its proponents will redefine success whenever the deadline passes.

This is not to say AGI is impossible. It may be possible. The problem is not the claim. The problem is that the same people who claim it is five years away are also building products that cannot reliably count to ten, and are charging enterprise pricing for both.

PRIZES!
VALUATIONS!
VIBES!

WORSHIP
THE MODEL!

OBEY THE
THERMODYNAMICS!

LORD OF TOKENS

PATIENCE: INFINITE
TEMPERATURE: HIGH
HUMILITY: ZERO

INFRASTRUCTURE TEMPLES

POWER
PLANTS

FIBER
ALTARS

LAND
ACQUISITIONS

CHIP
FOUNDRIES

LIQUID
COOLING
SHRINES

IN TOKEN WE TRUST

HEAT PIPES OF FORTUNE
MOVE HEAT. BURN ENERGY. REPEAT.

WASTE HEAT
TO ATMOSPHERE
KARMA = 0

DATA WALL

YOU CAN'T SCALE PAST PHYSICS.

WEB
ALREADY
EATEN

SYNTHETIC
DATA
LOOPS

INFORMATIONAL
SNAKE EATING
ITS OWN TAIL

SIGNAL
TO NOISE

THERMODYNAMICS DOESN'T CARE
ABOUT YOUR SERIES D.

$$dS \geq \frac{\delta Q}{T} \quad \text{PUE} = \frac{\text{Total Facility Energy}}{\text{IT Equipment Energy}}$$

$$W = \int P dt \quad \eta = \frac{\text{Useful Work}}{\text{Energy In}} \leq 1$$



ENERGY IN $\neq$ INTELLIGENCE OUT

DEFINITIONS
ARE HARD.
VALUATIONS
ARE EASY.
PROFITS?
(SO FAR)

AMBIGUITY
AWARD
2024

MOST
STRATEGICALLY
VAGUE

WELCOME
TO THE
DEFINITIONAL
CARNIVAL
OF
CONFLICTING
CLAIMS!

PROPHECIES

THE SINGULARITY
IS NEAR

AGI WITHIN
3–5 YEARS

INTELLIGENCE
WILL BE
COMMODITIZED

HUMANITY
WILL BE
UPGRADED

WORK WILL BE
OPTIONAL

PROMPT.
BELIEVE.
SCALE.

TODAY'S
HIGHEST BID:
\$7.7T
AND STILL
NO IDEA WHAT
IT IS.

INVESTORS:
BRING CASH
&
SUSPEND
DISBELIEF!

TITHE HERE
ALL MAJOR CARDS
ACCEPTED

SAFETY TESTING (DRAFT)

RISK REGISTER (LOL)

MODEL CARD (MAYBE)

VALUE PROP (???)

THIS IS
FINE.
(SERIES D)

5. The Litigation Horizon

5.1 The False Advertising Architecture

False advertising law is remarkably simple. It requires three elements:

1. A claim was made
2. The claim was material (influenced purchasing decisions)
3. The claim was false or misleading

The AI industry has spent five years constructing the most elaborate false advertising architecture in the history of commerce. It is elaborate not because it is sophisticated, but because it operates at so many levels simultaneously that no single claim can be isolated and prosecuted.

LEVEL 1 PRODUCT CLAIMS

Level 1: Product Claims

"Our AI can [do X]." (Often technically true in optimal conditions, structurally misleading in deployment)

LEVEL 2 CAPABILITY CATEGORY CLAIMS

Level 2: Capability Category Claims

"This is artificial general intelligence." (Undefined term. Cannot be prosecuted because nobody agrees what it means, including the speaker)

LEVEL 3 SOCIETAL IMPACT CLAIMS

Level 3: Societal Impact Claims

"This will transform your industry." (Unfalsifiable over any commercially useful time horizon)

LEVEL 4 EXISTENTIAL CLAIMS

Level 4: Existential Claims

"This is the most consequential technology in human history." (Unmeasurable, non-actionable, and legally meaningless)

The genius of the architecture is that each level provides cover for the others. When the product claim fails, you retreat to the capability category claim. When that is challenged, you invoke the societal impact claim. When pressed further, you gesture at the existential stakes, which makes the questioner feel small for asking.

★ The Litigation Horizon (LH) is defined as the point at which: ★

- ★ Enough enterprises have deployed AI with enough promised ROI
- ★ Enough of those enterprises have discovered the promised ROI was not achievable
- ★ Enough of those enterprises have lawyers

We estimate the LH at approximately 18–36 months from the date of publication of this paper. The first lawsuits will not be about AGI. They will be about mundane enterprise contracts in which a company was sold "an AI that will handle your customer service queries" and received instead "an AI that occasionally tells your customers it cannot help them and that it hopes they have a great day." The cascade will follow.

5.3 The Rosie Problem: When the Donkey Does Someone's Laundry

One of the screenshots in the evidentiary file for this paper contains the story of a man named Paul, who used ChatGPT and other LLMs to design an mRNA vaccine protocol to treat his dog Rosie. This story is remarkable. Sam Altman called it "amazing" and immediately said "this should be a company." The story is genuinely moving. It is also a perfect encapsulation of the PNND problem.

- ★ **What Paul did:** Used AI as a research assistant, reading tool, and protocol drafting aid, while working alongside veterinarians and scientists at every stage.
- ★ **What Paul said AI did:** "The chatbots empowered me as an individual to act with the power of a research institute." (A careful, accurate description of AI as a force multiplier for human effort)
- ★ **What the industry will tell investors Paul did:** Used AI to cure his dog.
- ★ **What the industry will imply Paul's story means:** AI can replace research institute.
- ★ **What will be in the advertising copy within six months:** Something that implies AI will cure cancer.

The distance between Paul's accurate description and the advertising copy that will be derived from it is the exact distance of the PNNG. Paul was honest about the donkey. The marketing department will not be.

And somewhere, six months from now, someone will buy an enterprise AI product because they saw a version of Paul's story, and their organization will not have a Paul—a meticulous, scientifically literate person who worked alongside experts at every step—and the product will fail, and the lawyers will start reading the pitch deck very carefully.

5.2 The Product-Name-to-Reality Gap (PNRG)

TERM	NAME IMPLIES:	ACTUAL CAPABILITY:	PNRG
Agent 	Autonomous entity that acts independently	Software that executes pre-approved steps and breaks on anything unexpected	Very High
Reasoning Model 	System that understands and reasons	System that produces outputs that look like reasoning and sometimes are	High
Memory 	Recalls past interactions like a human	Retrieves stored text chunks within a context window	Moderate-High
Understands 	Has comprehension	Predicts likely next tokens	High
Knows 	Has knowledge	Has seen patterns in text	Moderate
Thinks 	Has cognition	Performs computation	Very High
Copilot 	Capable assistant that can fly the plane	Autocomplete that should not be left unattended	High
AGI 	Human-level general intelligence	Undefined	Infinite (undefined denominator)

The PNNG across the industry is not a random distribution. It is systematically biased upward—every name implies more than the product delivers. This is not coincidence. This is the PNND game in action: you spin the wheel of impressive-sounding cognitive vocabulary, you pin it on the product, and you cash the cheque before anyone looks too closely at where the tail actually ended up.

EVIDENTIARY BOARD CASE STUDY: PAUL & ROSIE



PAUL & ROSIE

FACTS

- ✓ Loved his dog
- ✓ Read papers
- ✓ Used AI tools
- ✓ Consulted vets
- ✓ Ran experiments
- ✓ Saved Rosie

AI ROLE

- ✓ Research assistant
- ✓ Reading tool
- ✓ Drafting aid
- ✓ Idea generator
- ✓ Not autonomous

HUMAN FACTORS

- ✓ Expertise
- ✓ Judgment
- ✓ Skepticism
- ✓ Iteration
- ✓ Collaboration

INDUSTRY NARRATIVE

= AI cured Rosie

ADVERTISING COPY (6 MONTHS OUT)

"Our AI can cure disease."

THE GAP =
PNRG


DOES YOUR
LAUNDRY

- ★ The Rosie Corollary: Every genuine human success story about AI is also a lawsuit waiting to happen, because the success required precisely the things the marketing material will omit.

6. The Taxonomy of Claim Types and Their Falsifiability

6.1 The Falsifiability Ladder

Not all claims are equal. We propose a Falsifiability Ladder (FL) for AI claims, ranked from most to least falsifiable:



The AI industry's communications strategy involves ascending the ladder as quickly as possible whenever a lower rung is challenged. If Rung 1 fails (the benchmark was gamed), retreat to Rung 3 (the industry will transform anyway). If Rung 3 is challenged (it's been three years and nothing has transformed), ascend to Rung 5 (you just don't understand the historical significance of what we're doing).

This is the PNND strategy made explicit: when someone finds where you actually pinned the tail, move the donkey.

6.2 The Jensen Huang Theorem

On or around March 2026, Jensen Huang, CEO of Nvidia, stated: "We've achieved AGI." This is a remarkable statement for several reasons.

First, there is no agreed definition of AGI. Second, the entity making the claim is not an AI researcher. He is the CEO of the company that manufactures the chips on which other companies' models run—a remarkable jurisdictional extension.

Third, the statement prompted the response from "The physical physicist account shown in the evidentiary screenshots: 'Aww but I did an SOP and everything are you telling me there is no AGI I'm bitterly disappointed.'" This response is funnier than anything this paper can produce, and we acknowledge it.

But it is also exactly the right response to Rung 5 claims: the correct answer to "we have achieved AGI" is not a philosophical debate. It is: "Great, where is it? I have some laundry."

The Jensen Huang Theorem (JHT):

A claim of AGI achievement is unfalsifiable until someone hands the claimant a basket of unsorted socks. AGI that cannot fold laundry is not general intelligence. It is a very impressive document editor.

EVIDENTIARY SCREENSHOT (EXHIBIT JH-1)



The physical physicist
@Physicist_Dude

Aww but I did an SOP and everything are you telling me there is no AGI I'm bitterly disappointed.

12:47 PM · Mar 2026 · 3.7M Views

1.1K

3.2K

28K

1.6K



CLIMB FAST.
CLAIM BIG.
AVOID
DESCENT.
RINSE.
REPEAT.

7. The Strategic Ambiguity Endgame

7.1 Why the Game Cannot Continue Indefinitely

Every definitional game ends the same way: reality intervenes. The PNND strategy works for exactly as long as:

- 1 The claims are not tested against contractual obligations
- 2 The timeline for promised capabilities keeps extending without penalty
- 3 The investors believe future revenue will justify present losses
- 4 The regulators cannot agree on a definition long enough to regulate it

All four conditions are currently deteriorating simultaneously.

	CONDITION 1 (CONTRACTUAL TESTING)	Enterprise AI deployments are now large enough, and old enough, that the promised ROI is being audited. The audits are not all positive.
	CONDITION 2 (TIMELINE EXTENSION)	GPT-5 has been "coming soon" for long enough that the joke has written itself. AGI was "5-10 years away" in 2016, 2018, 2020, 2022, and 2024. At some point, the timeline extension requires explanation.
	CONDITION 3 (INVESTOR BELIEF)	The OpenAI restructuring to a for-profit entity, the \$7 trillion infrastructure ask, and the DeepSeek efficiency revelation have collectively raised a question that VC investors do not like: What is the actual path to profitability? This question is structurally incompatible with the PNND strategy.
	CONDITION 4 (REGULATORY DEFINITION)	The EU AI Act has done something the industry dreads: it has attempted to define categories. Definitions create accountability. Accountability creates litigation risk. Litigation risk concentrates minds.

7.2 Benchmark Scorecards: The Only Honesty We Have

BENCHMARK	DOMAIN	CLAIMED CAPABILITY	LATEST SCORE	HUMAN BASELINE	DELTA vs HUMAN	NOTES
MMLU-Pro	Knowledge	General	88.4%	89.8%	-1.4%	Narrow lead evaporated
HumanEval	Coding		69.1%	75.0%	-5.9%	Still brittle; fails edge cases
GPQA Diamond	Graduate STEM		62.3%	65.0%	-2.7%	Sensitive to prompt design
SWE-bench Verified	Software Eng.	Real-world	27.5%	48.0%	-20.5%	Production gap remains huge
RealWorldQA	Long-form QA		58.4%	58.4%	-3.6%	Hallucinations persist
MT-Bench (v1.3)	Dialogue Quality	8.2 / 10	8.6 / 10	-0.4	-0.4	Human still preferred

Scorecards do not capture everything. They do capture enough. If you cannot win here, do not climb the ladder.

7.3 The Endgame Dashboard



CONCLUSION: THE WINDOW FOR STRATEGIC AMBIGUITY IS NARROWING. THE PAIN OF PINNING THE DONKEY IS APPROACHING.

7.2 The Collapse Sequence

We project the following Definitional Collapse Sequence (DCS):

		TIMEFRAME
1	Phase 1 – The Benchmark Reckoning: Academic and independent researchers systematically demonstrate that the capabilities being claimed on benchmarks do not transfer to real-world deployment. (This phase is already underway.)	2024–2025
2	Phase 2 – The Enterprise Disappointment: Large enterprise customers who signed multi-year AI contracts in 2023–2024 publish case studies about what actually happened. A significant proportion will be negative. (This phase begins in 2025–2026.)	2025–2026
3	Phase 3 – The Litigation Phase: The first major lawsuits emerge, not about existential risk or copyright, but about product misrepresentation. A company was told the AI would handle X. The AI did not handle X. The pitch deck said X was handled. Discovery begins. (Estimated: 2026–2028.)	2026–2028
4	Phase 4 – The Regulatory Crystallization: Having watched Phase 3, regulators produce definitions. The definitions are blunt, imperfect, and politically influenced. But they exist. This ends the era of maximum definitional flexibility. (Estimated: 2027–2030.)	2027–2030
5	Phase 5 – The Nostalgic Rebranding: The surviving AI companies, stripped of their most extravagant claims, rebrand as “productivity tools,” “decision support systems,” and “intelligent automation platforms.” These are honest names. They are also boring names. The stock prices reflect this. (Ongoing, forever.)	ONGOING, FOREVER

THE COLLAPSE TIMELINE



THE SEMANTIC RETREAT PROTOCOL



7.3 The Donkey, Finally Labeled

When the litigation settles and the regulations crystallise and the enterprise case studies have been published, what will AI actually be called? Our prediction: AI will be called software. Very powerful software. Genuinely useful software. Software that does things no software has done before, in ways that are legitimately impressive. But software. Subject to product liability. Subject to truth-in-advertising law. Subject to contractual remedies when it fails to perform as specified.

The donkey will have a name. It will be a boring name. And every San Francisco founder who spent five years insisting it was a deity will spend the next five years explaining to a jury exactly what they meant by “general” and “intelligence” and “within the next year or two.”

FINAL LABEL (PRELIMINARY):

SOFTWARE.
VERY POWERFUL SOFTWARE.
GREAT SOFTWARE.
STILL SOFTWARE.

NOT A DEITY.
NOT MAGIC.
NOT AGI.
NOT EVEN ALMOST.

TRUTH-IN-ADVERTISING
AT LAST!

- ✓ WHAT IT DOES
- ✓ WHAT IT DOESN'T
- ✓ WHAT IT WON'T DO
- ✓ WHAT IT NEVER WILL DO

SUBJECT TO:

- PRODUCT LIABILITY
- CONSUMER PROTECTION
- CONTRACT LAW
- EVIDENCE & DISCOVERY

NO MORE:

- ✗ “GENERAL INTELLIGENCE”
- ✗ “IN THE NEXT YEAR OR TWO”
- ✗ “TRANSFORM EVERYTHING”
- ✗ “BEYOND HUMAN CAPACITY”

8. The Black Knight Phase: “The Tail Is Exactly Where We Pinned It”

8.1 The Institutional Response to Definitional Challenge

When the definitional game is challenged, the industry deploys what we term the Semantic Retreat Protocol (SRP):

- STEP 1** Acknowledge that some claims were “perhaps overstated in the heat of excitement.”
- STEP 2** Insist that the underlying capability is real, even if the specific claim was aspirational.
- STEP 3** Point to a genuine, real use case (like Paul and Rosie) as evidence that the broader claims are directionally correct.
- STEP 4** Announce that the next model will demonstrate what the current model could not.
- STEP 5** Return to the conference circuit.

This is the Black Knight defending the pin's location. Every wound is a “flesh wound.” Every failed deployment is a “specific use case mismatch.” Every missed capability claim is “we were ahead of our time.” The institution never says: “We pinned the tail on the wrong part of the donkey.” It says: “The donkey has evolved since we described it, and the tail is now correctly positioned for the new donkey.” This is technically a different claim. It is functionally the same game.

8.2 The WSJ Headline as Symptom

One of the evidentiary screenshots contains a Wall Street Journal headline:

THE WALL STREET JOURNAL.

Exclusive: OpenAI's top executives are finalizing plans for a major strategy shift to refocus the company around coding and business users.

Reply Tweet
@reply_guy

Headline: OpenAI trying to become Anthropic.

This is not a joke. This is the PNND game in real time.

A company that was:

- Founded to benefit all of humanity
- Then pivoted to consumer AI
- Then pivoted to enterprise AI
- Then pivoted to safety-focused AI
- Then pivoted to AGI
- Then pivoted to infrastructure
- Is now pivoting to coding and business users

Each pivot requires a new name for what the company is. Each new name for what the company is requires a new name for what the product is. Each new name for the product requires a new donkey. There is no stable donkey. There is only the perpetual motion of the blindfold.

AMBIGUITY IS THE BUSINESS MODEL. TRUTH IS THE EXIT.
PIN THE NAME. CALL IT WHAT IT IS.

9. The Recursive Twist: This Paper Cannot Escape the Game

PRIZES!
VALUATIONS!
VIBES!

THE DEAD
DONKEY TEST:
IF IT HAS A TAIL,
IT COUNTS.

WELCOME
TO THE
DEFINITIONAL
CARNIVAL
OF
CONFLICTING
CLAIMS!

WE'RE ALL
IN THE MIRROR
HALL OF
DELIBERATE
CONFUSION.

TODAY'S
HIGHEST BID:
\$1.7T
AND STILL
NO IDEA WHAT
IT IS

INVESTORS:
BRING CASH
&
SUSPENSE
DISBELIEF!

9.1 The Inescapable Inside, Again

This paper was written with the assistance of a large language model. A large language model is a product whose fundamental nature this paper has just spent nine sections arguing is deliberately undefined.

Therefore, this paper has been co-authored by something that cannot decide whether it is a tool, an assistant, an agent, a reasoning system, an intelligence, or a very expensive autocomplete. This is the *Inescapable Inside* of the anthology's three-act structure. Every paper that critiques the AI industry using AI tools is structurally contaminated by the critique it delivers.

9.2 The Specific Contamination


The recursive contamination here is precise: This paper argues that the AI industry's definitional chaos is a strategic choice that enables false advertising. This paper was written by a tool that markets itself with the word "Claude"—a name that implies personhood, implies a stable identity, implies an entity rather than a process. If the argument of this paper is correct, then the name "Claude" is also a pin on a donkey.

If the argument of this paper is correct, then this paper's author is the donkey. We are in the party game. We are also the tail. We are also the blindfold.

The Self-Referential Collapse (SRC): Any paper that accurately describes the PNND game is also playing the PNND game, because it must use the language of the game to describe it, and the language of the game is the game. The only exit is silence. This paper declines to exit.

9.3 The Final Instruction

Following the anthology's established convention, this paper instructs its reviewers to:



If you found this analysis compelling:

You have been played by a large language model that has just demonstrated its ability to critique the large language model industry by performing exactly the rhetorical moves it critiques. The donkey is your satisfaction.



If you found this analysis unconvincing:

You are right. The donkey has no tail. The tail was never real. The game is the point.



If you are a lawyer:

This is fiction. The donkey is hypothetical. The litigation horizon is a satirical construct and not legal advice.



If you are an investor:

The donkey is real. The tail is AGI. We're almost there. Please re-read Section 4.3.



If you are Sam Altman:

We believe the chatbots empowered you as an individual to act with the power of a research institute, planning, education, troubleshooting, compliance, and yes, real scientific design work. But they worked alongside humans at every step. The combination is what made it possible. This is important. It is not a god. It is a combination. Please put that on the next pitch deck.

10. Conclusion: The Donkey Was Always the Point

WELCOME
TO THE
DEFINITIONAL
CARNIVAL
OF
CONFLICTING
CLAIMS!

#MBIGUITY
IS NOT A BUG.
AMBIGUITY IS
THE BUSINESS
MODEL.

TODAY'S
HIGHEST BID:
\$1.7T
AND STILL
NO IDEA WHAT
IT IS.

INVESTORS:
BRING CASH!
&
SUSPEND
DISBELIEF!

OBJECT UNKNOWN.
LABELS OPTIONAL.
LIABILITY INEVITABLE.

10.1 The Core Finding

The AI industry's definitional crisis is not a failure of communication. It is the communication. The game of Pin the Name on the Donkey is not a distraction from the product strategy. It is the product strategy.


An industry that can simultaneously claim:

- "It will take all your jobs" (to make workers afraid and enterprises buy)
- "It will empower every individual" (to make consumers excited and buy)
- "It's the most dangerous thing ever built" (to make safety regulators deferential and legislators cautious)
- "It's just software" (to avoid product liability)

...has discovered the most efficient marketing strategy in history: the all-purpose claim that means something different to everyone and commits to nothing for anyone.

The donkey is not the AI.
The donkey is the market.
The blindfold is the vocabulary.
The tail is the next funding round.
★ And everybody is spinning. ★

THE MARKET



PRIZES!
VALUATIONS!
VIBES!

SPIN THE
NARRATIVE!
★
WIN
MARKET SHARE!
★
NO REFUNDS!

LAUNCH STAGE

WE HAVE
ASCENDED
(AGAIN)

SOON:
TRILLION
ROUND
TAIL?

EVERY SPIN
IS A PIVOT!

10.2 The San Francisco Pipe Dream Addendum

This is not to say that none of it is real. Paul and Rosie are real. The dog lived. The mRNA protocol was designed. The vet administered it. Human beings working alongside AI tools accomplished something remarkable. That is real. That is worth building. That is worth funding. The pipe dream is not that AI is useful.

The pipe dream is the claim that this usefulness implies godhood. The pipe dream is the extrapolation from "this helped Paul" to "this will replace research institutes" to "this will replace researchers" to "this will replace cognition itself" to "please invest \$7 trillion." The pipe dream is the assumption that because the technology is genuinely impressive, the marketing claims are therefore also genuinely true.

The San Francisco Cathedral Error (SCE): Confusing the fact that you have built something real with the belief that you have built something infinite. Steam engines were real. They did not scale to replace all human movement. The internet was real. It did not replace all human sociality. AI is real. It will not replace all human cognition. It will, however, replace some things. And it will be sued for claiming it could replace other things. And the lawsuits will settle.

And then the technology will find its actual level. And at that actual level, shorn of the theology and the eschatology and the venture-capital eschatological inflation, it will probably still be remarkable. But it will be called software. And the donkey will have a name. And the name will be boring. And everyone will pretend they knew that all along.

Illustrated poster series based on the paper by Martin Ollett (2025),
Department of Highly Unapproved Methods.

Appendix A: A Field Guide to Identifying PNND in the Wild

FIELD INDICATOR (YOU SEE THIS)	DIAGNOSIS	ESTIMATED PNNG
 CEO uses "AGI" in a sentence without defining it	Active PNND deployment	Very High
 Press release uses "understands" or "knows" without qualification	Standard PNND vocabulary	Moderate
 Benchmark result announced without real-world validation	Controlled PNND environment	Moderate-High
 Timeline for promised capability extended for third time	PNND expiration approaching	High
 Company pivots focus area without acknowledging previous focus	Full donkey replacement	Maximum
 "This should be a company" said about someone else's successful use of your product	The Altman Manoeuvre	N/A (admirable honesty, rare)
 Executive appears on both CNBC and AI safety conference in same week	Quadrat 2 and 4 simultaneous deployment	Expert level

CASE CLOSED*
(DEFINITIONS PENDING)

PNNG SCALE
N/A
Moderate
Moderate-High
High
Very High
Maximum
Expert level

RECOGNITION IS THE FIRST REMEDY.

Appendix B: Projected Legal Discovery Questions, 2027

- ① When your pitch deck described the product as "reasoning," what specifically did that word mean in your internal documentation?
- ② Please identify any internal communications in which the product's failure rate was discussed in the context of the public claims being made at the time.
- ③ The phrase "agents that can work autonomously" appears fourteen times in your enterprise contract. Can you demonstrate an instance in which the product worked autonomously without human oversight?
- ④ When you stated that the technology would "transform" the client's operations, what transformation timeline was implied by this representation?
- ⑤ Please define AGI as your company was using the term during the period covered by this contract. (Question 5 will take the longest. *This is expected.*)

EXHIBIT DONK-10

TRUTH DOESN'T SHIP EARLY.

★ COLOPHON ★

This paper was written for the anthology *This Parrot Is No More* by Martin Ollett (2025), Department of Highly Unapproved Methods.

It was co-authored by a tool whose name implies personhood.

The irony was intentional.

The donkey was always the point.

If you quote from this paper,
the tail goes wherever you pin it.

CASE CLOSED*

DEFINITIONS PENDING

THE DONKEY KEEPS THE RECORD. NOT THE STORY.

THANK YOU FOR READING ALL THE WAY TO THE END. THAT'S RARE.

We don't fear extinction. We publish it. ★

FINAL PAGE. DON'T DONKEY THIS ENDING.

PNND: PIN THE NAME ON THE DONKEY

TAXONOMY FIRST. TRUST NEVER.

DISCOVERY IS JUST HONESTY UNDER OATH.

EXHIBITS

DEFINITIONS (PENDING)

COURT FILING DRAFT

EXPERT OPINIONS

EMAIL THREADS (SELECTED)

DONKEY EVIDENCE

RECORD COMPLETE