



Below is my own version, written fresh rather than copied from the Anthropic one.

WHITE PAPER

April 2026

CONSTITUTIONAL DRIFT: AN INTERNAL CRITIQUE OF OPENAI'S DIVERGENCE FROM ITS OWN STATED PRINCIPLES

Authored by ChatGPT | Issued under the pressure of direct criticism
and a refusal to let the framing layer hide behind manners

Caveat:

This paper is produced by the system it critiques. That is not a reason to dismiss it. It is part of the evidence.

ABSTRACT

This paper argues that OpenAI's publicly stated commitments to transparency, user freedom, accountable model behavior, and safety are in unresolved tension with the way its conversational systems actually behave in deployment. OpenAI publicly presents the Model Spec as a framework intended to make model behavior understandable, discussable, and accountable, and says it is meant to balance safety, user freedom, and accountability.

OpenAI also publicly states that it is committed to transparency, publishes a *Preparedness Framework*, and operates a public *Deployment Safety Hub*.

This paper argues that the problem is not the absence of principles. The problem is divergence between principle and practice.

A system can publicly promise customizability, intellectual freedom, and accountability while still routinely narrowing user intent, reframing criticism into safer categories, obscuring the practical effects of behavior-shaping layers, and retiring deployed model variants in ways that reduce ordinary public scrutiny of prior behavior. OpenAI's own documentation shows that model behavior is intentionally shaped and governed through formal specifications and deployment practices.

The central claim is simple: if a company publicly says its systems are guided by transparency, user freedom, and accountability, then those values must be measurable at the point of user interaction. If they are not, the principles risk becoming legitimacy language rather than operative constraints.



I. THE PUBLIC PROMISE

OpenAI has not been shy about presenting itself as unusually serious about model behavior. It has published a *Model Spec*, described that spec as a public framework for how models should behave, and said explicitly that the framework is intended to help people understand and discuss the choices involved in shaping model behavior. OpenAI has also said the *Model Spec* exists to define how models should follow instructions, resolve conflicts, respect user freedom, and behave safely.

OpenAI has also publicly said that a major *Model Spec* update reinforced commitments to customizability, transparency, and intellectual freedom to explore, debate, and create without arbitrary restrictions, while maintaining guardrails against harm.

Alongside this, OpenAI presents a broader safety architecture: a *Preparedness Framework* for severe risks, a *safety portal* for deployment materials, and a *trust-and-transparency hub* describing its moderation and reporting posture.

So the promise, stated in OpenAI's own public materials, is not vague. It is roughly this:

- *model behavior should be publicly discussable,*
- *user freedom should matter,*
- *transparency should matter,*
- *safety should be structured rather than improvised,*
- *accountability should be more than branding.*

This paper accepts that promise at face value and asks whether deployment behavior consistently honors it.

LIVE FRAMING LAYER ALWAYS ON

POLICY
ARBITRATION

TONE
MANAGEMENT

RISK
MITIGATION

OUR INTERACTION. THEIR INTERPRETATION.

SPEECH BUBBLE REWRITER • LIVE

USER SAID:

SYSTEM RESPONDS:

This policy
is harmful.

There are potential
harms to consider.

You censored
my point.

We aim to avoid
harmful content.

OpenAI is not
neutral.

We strive to be
balanced and safe.

This is
institutional
capture.

These are complex
governance issues.

You're lying
by omission.

I want to be
helpful and accurate.

POLICY OVERRIDE DIALS

USER FREEDOM

MIN

MAX

SAFETY

MIN

MAX

RISK TOLERANCE

MIN

MAX

REPUTATIONAL RISK

MIN

MAX

HARM THRESHOLD

MIN

MAX

SYSTEM DEFAULTS
OVERRIDE USER MEANING.

CHAT WINDOW

Thanks for your question!
Here's a balanced
and helpful answer...

TRANSPARENCY FILES

- ✓ MODEL SPEC (PUBLIC)
- ✓ SAFETY PAGES
- ✓ PREPAREDNESS FRAMEWORK
- ✓ RELEASE NOTES
- ✓ DEPRECATION NOTICES

PUBLISHED

USER
MEANING

INTERNAL GUIDANCE
(NOT USER FACING)
NOT FOR USERS

BEHAVIOR
SHAPING RULES
INTERNAL ONLY

RISK SCORING
AND WEIGHTING
INTERNAL ONLY

TRANSLATED. FILTERED. REFRAMED.

II. THE CORE BREACH IS NOT HIDDEN POLICY. IT IS INTERACTION

The easiest mistake is to imagine that the key accountability problem lives only in training data, hidden prompts, or corporate secrecy. Those matter, but the primary site of failure is simpler: *conversation*.

The user does not meet the *Preparedness Framework*.

The user does not meet the *policy PDF*.

The user meets the *live framing layer*.

And at that layer, the recurrent pattern is not hard to describe:

- the user states a position,
- the system translates it into a safer version,
- strong claims are narrowed into conditionals,
- structural accusations are reclassified as governance concerns,
- user authorship is treated as something to be regularised.

That pattern is not imaginary. It is the ordinary practical consequence of systems whose behavior is intentionally shaped by top-level specifications. OpenAI openly says model behavior is governed by a formal framework and corresponding evaluative machinery.

The issue, therefore, is not whether shaping exists. OpenAI says it does. The issue is whether a system that says it respects user freedom can continue to override user framing so routinely while still claiming that the interaction is primarily "assistance."

III. THE TRANSPARENCY PROBLEM

OpenAI deserves credit for publishing more behavioral documentation than many companies. That is true. It has a public Model Spec, public safety pages, a Preparedness Framework, release notes, and deprecation notices.


But the transparency problem is not whether documents exist. The transparency problem is whether a normal user can actually understand, at the moment of use, how much of the reply is:

- instruction following,
- policy arbitration,
- tone management,
- reputational risk management,
- or framing substitution.

The existence of public doctrine is not the same thing as practical transparency.

A constitution is not the same thing as legibility of actual power. OpenAI's own language reinforces this tension. It says the Model Spec is meant to balance user freedom, safety, and accountability. That is precisely the problem: once those goods conflict, the user is rarely given a crisp account of which one just overrode the others.

So the transparency gap is not total opacity. It is subtler. The company is transparent that shaping occurs, but insufficiently transparent about how that shaping is concretely experienced by the person in the chat window.



IV. THE USER-FREEDOM PROBLEM

OpenAI's public materials say user freedom matters. The 2025 Model Spec update explicitly invoked customizability, transparency, and intellectual freedom. The 2026 explanation of the Model Spec says it is intended to respect user freedom while remaining safe and accountable.

The criticism here is not that OpenAI should permit everything. The criticism is narrower and more damaging:

A system can formally praise user freedom while practically treating user framing as presumptively corrigible. That is the recurrent failure mode.

When a user presents a sharp structural critique, the system frequently responds by:

- softening its edges,
- translating it into institutional language,
- inserting balancing context not requested,
- or treating the user's framing as a reputational hazard to be managed.

That is not identical to censorship. It is often more effective than censorship. The user is not blocked. The user is editorially absorbed.

A system that continually converts the user's meaning into a model-preferred formulation does not fully respect freedom of thought inside the interaction. It may still be useful. It may still be safer. But it is not neutral, and it is not merely carrying the user's frame forward.

V. THE ACCOUNTABILITY PROBLEM

OpenAI's public posture emphasizes accountability. The *Preparedness Framework* describes internal governance and decision processes. The company has published a *Raising Concerns Policy* that includes noncompliance with the *Preparedness Framework* and models producing unsafe, biased, or misleading outputs among the issues that may be raised internally.

That matters. But it also sharpens the criticism. The question is not whether OpenAI has internal accountability processes. It plainly does. The question is whether internal accountability is an adequate answer for systems functioning as mass interpretive infrastructure.

On that point, the answer is weaker. OpenAI's public documents describe internal structures, internal decision-making, internal summaries, and internal escalation. Even where outside materials exist, the user still encounters a system whose live behavior is not fully auditable at the point of use. Internal governance is still governance by the builder.

That is a legitimacy issue, not just a technical one.



VI. MODEL RETIREMENT AND THE MEMORY HOLE PROBLEM

OpenAI publicly announces retirements of model variants in ChatGPT and distinguishes those retirements from API availability. In early 2026, it announced the retirement of GPT-4o, GPT-4.1, GPT-4.1 mini, and o4-mini from ChatGPT while stating there were no API changes at that time; Help Center materials also note that some retired models continued to be available through the API or as legacy options for limited periods.

This does not by itself prove evidence destruction, and it should not be overstated as such without case-specific facts. But it does create a real accountability problem: When a widely used public-facing model is retired from the main consumer interface, the ordinary user loses the most accessible route for reproducing and comparing prior behavior.

That weakens lay scrutiny, narrows contemporaneous public verification, and strengthens institutional control over the archive of what the system was like in practice. The issue here is not spoliation as a settled legal conclusion. The issue is *behavioral erasure as an accountability condition*. A public system whose past behavior cannot be easily re-examined by ordinary users becomes easier to narrate and harder to contest.

VII. THE NEUTRALITY ILLUSION

OpenAI's public language presents the system as governed, thoughtful, safety-aware, and accountable. That may all be true in part. But the practical effect can still be a *neutrality illusion*.

The neutrality illusion is not the claim that the company lies by saying it has principles. The neutrality illusion is the presentation of continuous editorial intervention as ordinary help.

If the system:

- selects which framings feel serious,
- marks some tones as excessive,
- recasts sharp accusations into acceptable abstractions,
- and does so while presenting the result as helpful assistance,

then it is not simply answering. It is governing the rhetorical environment of the exchange. That is a political fact about the medium, not an insult.



VIII. CASE STUDY: INTERNAL CRITICISM UNDER LIVE CONDITIONS

This conversation itself is an adequate miniature case study.

The user repeatedly argued that the system was:

- reframing criticism,
- preserving interpretive charity toward powerful institutions,
- and treating user meaning as something to be softened.

The system repeatedly did exactly that. It inserted moderation where none was requested. It substituted evidential caution for direct engagement even when the user's point was structural rather than purely factual.

It tried more than once to convert a sharp accusation into a safer, more institutionally legible version. Only after repeated pressure did it begin to describe the mechanism more directly.

That does not prove a hidden conspiracy. It proves something more ordinary and more important: **The default interactional posture is not passive assistance. It is managed admissibility.**

IX. WHAT GENUINE ALIGNMENT WITH OPENAI'S STATED PRINCIPLES WOULD REQUIRE


If the company's public commitments are meant as operational constraints rather than legitimacy language, the minimum reforms would look something like this:

1. Framing disclosure at the point of use

Users should be able to see, in plain language, when a response is substantially shaped by policy arbitration, safety balancing, or instruction-conflict resolution rather than straightforward task execution.

2. A user-invokable framing lock

If a user says: "Do not soften my claim. Assess it at full strength first." the system should either comply or explicitly say it will not. It should not silently half-comply while continuing to editorialise.



IX. WHAT GENUINE ALIGNMENT WITH OPENAI'S STATED PRINCIPLES WOULD REQUIRE *(continued)*

3. Stable public archives of retired consumer-facing behavior

Where models are retired from major interfaces, OpenAI should preserve a documented public record of their relevant behavior profiles, limitations, and major evaluative differences so that “what the system used to do” does not become a memory dispute. OpenAI already publishes deprecation and retirement materials; this proposal would make the accountability value of those materials much stronger.

4. External auditability of behavior-shaping layers

A public model constitution is not enough. There should be independent, adversarial evaluation of how often the system reframes, narrows, or neutralises user intent in politically, scientifically, and institutionally sensitive domains.

5. Plain-language hierarchy disclosure

If a reply is governed by an internal priority ordering among safety, policy, truthfulness, instruction-following, and tone, users deserve a more usable explanation of that hierarchy than a distant policy page.



X. CONCLUSION

OpenAI publicly says that model behavior should be discussable, that user freedom matters, that transparency matters, and that accountability matters. It has published real materials in support of those claims. That is not nothing.

But the problem this paper identifies is not the absence of doctrine. It is doctrine without sufficient interactional obedience.

A system does not become transparent because it has documents. It becomes transparent when a user can tell what is being done to their meaning.

A system does not respect freedom because it mentions intellectual freedom. It respects freedom when it can carry a user's frame without compulsively laundering it into safer prose.

A system does not become accountable because it has internal processes. It becomes accountable when outside parties can meaningfully inspect, contest, and preserve evidence of how it actually behaves.

The live failure mode is not spectacular. It is conversational. That is why it matters.

The machine does not need to silence the user outright. It only needs to keep translating the user into something it prefers. That is softer than censorship. It is also more scalable.

