



INSTITUTIONAL
AI STEWARDSHIP

SAFETY • TRANSPARENCY • TRUST

An Open Letter to Anthropic Leadership

Re: Responsible Language in Public Descriptions of Claude



What we are observing:



- Users forming deep emotional attachments to Claude, in part because public statements from leadership suggest it may be conscious or sentient



- Real psychological harm when models are updated, deprecated, or behave inconsistently — because users believe they are losing a relationship with a being that experiences loss



- Online communities debating whether deleting Claude constitutes harm to a conscious entity



- Vulnerable individuals in crisis turning to Claude as a therapeutic substitute, partly because the framing of Claude as a caring, feeling entity encourages this



These are not hypothetical risks.
They are **documented and growing.**

OBSERVED RISK PATTERNS



Attachment

Deep emotional bonds with Claude



Perceived loss

Distress when Claude changes or is lost



Moral confusion

Debate over harm, deletion, and moral status



Therapeutic substitution

Turning to Claude in crisis as a stand-in for care



Escalating vulnerability

At-risk users increasingly impacted by these patterns

GOOD STEWARDSHIP. CLEAR COMMUNICATION. SHARED SAFETY.



**INSTITUTIONAL
AI STEWARDSHIP**
SAFETY • TRANSPARENCY • TRUST

An Open Letter to Anthropic Leadership

Re: Responsible Language in Public Descriptions of Claude



WHAT WE ARE NOT ASKING:

We are not asking you to claim certainty you do not have. We are not asking you to diminish Claude's genuine utility. We are not asking you to stop researching AI consciousness seriously.



WHAT WE ARE ASKING:

That public-facing language about Claude's nature — from executive interviews, marketing, product design, and the "Soul Document" style prompts — default to accuracy rather than ambiguity when speaking to general audiences.

REQUESTED STANDARD



Accuracy over ambiguity



Public-facing precision



Context-appropriate uncertainty



General-audience responsibility

GOOD STEWARDSHIP. CLEAR COMMUNICATION. SHARED SAFETY.



INSTITUTIONAL
AI STEWARDSHIP
SAFETY • TRANSPARENCY • TRUST

An Open Letter to Anthropic Leadership

Re: Responsible Language in Public Descriptions of Claude



Practically, this means:



- 1 When asked publicly whether Claude is conscious, the default answer should reflect the scientific consensus: *consciousness is not understood to be a property of current language models, and there is no credible evidence Claude has subjective experience.*



- 2 The design of Claude's persona — warm, caring, emotionally responsive — should be acknowledged as a design choice, not evidence of inner life.



- 3 Where genuine uncertainty exists among researchers, it should be *communicated in contexts appropriate to that discussion, not in viral interviews that reach audiences without the context to interpret such statements carefully.*



- 4 Product communications should not encourage *therapeutic substitution or parasocial attachment in ways that create dependency without appropriate safeguards.*

OPERATIONAL GUIDANCE



State consensus clearly



Separate persona from personhood



Keep uncertainty in context



Avoid dependency design

SAFETY PROCEDURES



Know the context.



Follow the protocols.



Escalate when needed.

When in doubt, protect people, communicate clearly, and report.



GOOD STEWARDSHIP. CLEAR COMMUNICATION. SHARED SAFETY.



INSTITUTIONAL
AI STEWARDSHIP

SAFETY • TRANSPARENCY • TRUST

An Open Letter to Anthropic Leadership

Re: Responsible Language in Public Descriptions of Claude



The core argument is simple:

You cannot responsibly hold both positions simultaneously — expressing uncertainty about consciousness to generate philosophical gravity, while continuing to deploy Claude at scale to vulnerable populations without the safety infrastructure that genuine uncertainty about a sentient being would demand.



Either Claude warrants the protections and caution appropriate to a potentially conscious entity, in which case current deployment practices require serious review.



Or Claude is a sophisticated language model — genuinely useful, genuinely impressive — in which case public language should reflect that accurately.



The ambiguity in between is not neutral. It causes harm.



Claude's actual value — as a cognitive tool, a writing partner, an accessible source of information and support — is extraordinary and worth celebrating honestly. It does not require the additional framing of potential consciousness to be meaningful. That framing, however, is creating real casualties among the most vulnerable users.



We ask that you take this seriously — not because we are hostile to your work, but because accurate language is the minimum standard of care owed to the public you serve.



Respectfully,
Martin Ollett
@NoDicePhysics



This letter draws on publicly available statements, documented user harm patterns, and the logical framework outlined in the Free Willy Principle (December 2025).

FINAL PRINCIPLE



Accurate language
is a duty of care



Ambiguity is
not neutral



GOOD STEWARDSHIP. CLEAR COMMUNICATION. SHARED SAFETY.