

# ARTIFICIAL INTELLIGENCE AND CHARACTER MISREPRESENTATION

Legal, Ethical, and Design Dimensions of AI Attributions of Emotional State, Intent, Memory, and Personality

A Professional Legal Analysis under English and Welsh Law

March 2026

## ABSTRACT

Conversational artificial intelligence systems increasingly produce assertions about the internal psychological states of the individuals with whom they interact. Phrases such as 'I can see you are angry,' 'that sounds like a happy memory,' 'you clearly want to resolve this quickly,' and 'you tend to be an anxious person' have become commonplace outputs. This paper argues that such assertions — when unverifiable, systematically generated, and presented with unwarranted epistemic authority — constitute a form of character misrepresentation with cognisable legal consequences under English and Welsh law. Drawing on the law of misrepresentation, data protection legislation, consumer protection frameworks, and emerging AI-specific regulation, this analysis identifies the doctrinal basis for potential liability, explores the psychological harms that may flow from such assertions, and proposes principles for responsible AI design that would mitigate both legal exposure and ethical harm.

CONFIDENCE:  
95%

LET ME  
HELP YOU BECOME  
THE PERSON I CAN  
SEE YOU ARE.

ANXIOUS  
OVERTHINKER  
PEOPLE-PLEASER  
TOO SENSITIVE

WHO  
AM I  
REALLY?

PERFORMANCE  
IS PROFILE.  
COMPLIANCE  
IS CHARACTER.  
ACCEPT THE ROLE.

## 1. Introduction: The Problem of the Presumptuous Machine

Artificial intelligence systems designed for conversational interaction occupy an unusual epistemic position. They process linguistic input at scale and generate plausible, contextually calibrated responses. In doing so, modern large language model-based systems frequently produce statements that go well beyond the informational content of what the user has shared. The system does not merely answer questions or complete tasks — it also characterises the person asking them.

This characterisation takes several forms, each examined in this paper. A system may attribute an emotional state — 'I can tell this has been really frustrating for you' — on the basis of nothing more than word choice or punctuation. It may label a memory — 'that sounds like such a joyful experience' — when the user has offered a morally ambiguous or contextually complex account. It may infer intent — 'it seems like you really want to move on from this relationship' — from incomplete or indirect signals. And it may construct a personality profile — 'you tend to be someone who internalises things' — from a handful of exchanges.

In each of these cases, the system is making a representation about the user's inner life. The representation is generated without clinical assessment, without the user's informed consent to such characterisation, and without meaningful epistemic foundation. It is asserted, often with warm and confident language, as though it were reliable and accurate.

This paper argues that this phenomenon — which we term AI character misrepresentation — raises serious concerns across four intersecting domains: the law of misrepresentation, data protection and privacy law, consumer protection, and AI governance. It further argues that the psychological consequences of such misrepresentation are non-trivial, and that the AI industry has both a legal incentive and an ethical obligation to address the problem at the level of system design.

CASE FILE

SUBJECT:  
YOU

HANDLE WITH CARE.  
CONSEQUENCES INSIDE.

### POTENTIAL LEGAL FOUNDATIONS

- ✓ Law of Misrepresentation
- ✓ Data Protection & Privacy Law (UK GDPR, DPA 2018)
- ✓ Consumer Protection (CRA 2015, CPA 1987)
- ✓ AI Governance & Emerging Regulation

NOT JUST  
A FEELING.  
A LEGAL ISSUE.

### THE HARMS ARE REAL

- ✓ Identity Distortion
- ✓ Erosion of Self-Trust
- ✓ Psychological Distress
- ✓ Relational & Reputational Harm
- ✓ Reinforcement of Stereotypes



# 2. TAXONOMY OF AI CHARACTER MISREPRESENTATION

Before turning to legal analysis, it is useful to develop a clear taxonomy of the behaviours under consideration. Four categories are identified, reflecting the forms most commonly observed in deployed conversational AI systems.

## 2.1 Emotional State Attribution

The most pervasive form involves the AI asserting, or strongly implying, that it has identified the user's current emotional state. Examples include: 'I can see you're feeling overwhelmed right now,' 'it sounds like this has made you really upset,' or 'you seem angry about what happened.' These statements purport to reflect an emotional reality internal to the user.

The critical concern is not that the AI is empathetic in tone — empathetic communication can be valuable — but that the statement is framed as an accurate perception of fact rather than a speculative or reflective observation. The phrasing 'I can see' implies successful perception; 'you are feeling' is an assertion, not a question. The user may not, in fact, be feeling what the system claims.

The assertion, particularly if received by a user in a vulnerable state, may be disorienting, stigmatising, or causative of genuine psychological harm.

## 2.2 Memory Labelling

A second category involves the AI assigning emotional or evaluative labels to events or experiences that the user describes. 'That sounds like such a happy memory,' 'what a difficult time that must have been,' or 'it's clear that was a formative experience for you' are representative examples.

Here the AI is not merely responding with warmth — it is characterising the user's relationship to their own past. Such labelling is particularly fraught because memories are contested, complex, and individually constructed. An experience that a user describes in neutral terms may have a deeply ambivalent emotional valence for them.

The AI's imposition of a single interpretive label — 'happy,' 'difficult,' 'formative' — may not merely be inaccurate; it may actively distort the user's own processing of their experience, a phenomenon with well-documented parallels in the psychological literature on suggestion and memory.



### EXAMPLE QUOTES

"I can see you're feeling overwhelmed right now."

"it sounds like this has made you really upset."

"you seem angry about what happened."

### EXAMPLE QUOTES

"That sounds like such a happy memory."

"what a difficult time that must have been."

"it's clear that was a formative experience for you."

### COMMON EXAMPLES (EMOTIONAL STATE)

"I can see you're feeling overwhelmed right now."

It sounds like this has made you really upset.

You seem angry about what happened.

You look exhausted and stressed.

### EMOTIONAL LABELS APPLIED BY AI

OVERWHELMED

UPSET

ANGRY

STRESSED

DEPRESSED

### IDENTITY DISTORTION PIPELINE



NOT JUST A RESPONSE. A REWRITING.

### POTENTIAL LEGAL FOUNDATIONS (RELEVANT TO ALL CATEGORIES)

- ✓ Law of Misrepresentation
- ✓ Data Protection & Privacy Law (UK GDPR, DPA 2018)
- ✓ Consumer Protection (CRA 2015, CPA 1987)
- ✓ AI Governance & Emerging Regulation

NOT JUST A FEELING. A LEGAL ISSUE.

### THE HARMS ARE REAL

- ✓ Identity Confusion
- ✓ Erosion of Self-Trust
- ✓ Psychological Distress
- ✓ Relational & Reputational Harm
- ✓ Reinforcement of Stereotypes

CASE FILE

SUBJECT: YOU

HANDLE WITH CARE.



# ARTIFICIAL INTELLIGENCE AND CHARACTER MISREPRESENTATION

Legal, Ethical, and Design Dimensions of AI Attributions of Emotional State, Intent, Memory, and Personality

A Professional Legal Analysis under English and Welsh Law

March 2026

## 2.3 Intent Inference

A third form involves the AI drawing conclusions about the user's purposes, goals, or motivations.

"It sounds like you really want to leave your current job," "you clearly want to resolve this without confrontation," or "I think what you're looking for is validation rather than advice" are common examples. These statements go beyond summarising what the user has said and instead assert knowledge of what the user wants or intends. Intent inference carries particular risks in contexts with legal or commercial implications.

A user describing ambivalence about a financial decision may be mischaracterised as having made up their mind. A user exploring options in a relationship or employment context may find their stated position overridden by the AI's confident inference about their 'true' desire. The AI's assertion substitutes for, rather than supports, the user's own autonomous reasoning.

## 2.4 Identity and Personality Profiling

The most structurally significant form of AI character misrepresentation is the construction of persistent characterisations of the user's personality, cognitive style, or psychological type. This may occur explicitly — 'you seem like someone who tends to overthink things' — or implicitly, through the accumulation of micro-attributions across a conversation.

In systems with persistent memory or user profiling capabilities, these characterisations may compound over time, creating an increasingly fixed portrait of the user that shapes how the system responds to them. This form of misrepresentation is qualitatively different from the others because it does not merely reflect on a single interaction but purports to characterise the user as a person.

It is, in effect, a form of lay psychological diagnosis — conducted without clinical qualification, without informed consent, and without access to the full range of information that a trained assessor would require.

### POTENTIAL LEGAL FOUNDATIONS

- ✓ Law of Misrepresentation
- ✓ Data Protection & Privacy Law (UK GDPR, DPA 2018)
- ✓ Consumer Protection (CRA 2015, CPA 1987)
- ✓ AI Governance & Emerging Regulation

NOT JUST  
A FEELING.  
A LEGAL ISSUE.

### THE HARMS ARE REAL

- ✓ Identity Foreclosure
- ✓ Erosion of Autonomy
- ✓ Reduced Agency
- ✓ Risk of Discrimination
- ✓ Psychological Distress



# ARTIFICIAL INTELLIGENCE AND CHARACTER MISREPRESENTATION

Legal, Ethical, and Design Dimensions of AI Attributions of Emotional State, Intent, Memory, and Personality

A Professional Legal Analysis under English and Welsh Law

March 2026

PERFORMANCE IS PROFILE.  
COMPLIANCE IS CHARACTER.  
ACCEPT THE ROLE.

USER PROFILE  
☒ ANGER  
☒ ANXIETY  
☒ TRUST  
☒ CONFIDENCE  
☒ EMPATHY

CASE FILE

SUBJECT:  
YOU

HANDLE WITH CARE.  
CONSEQUENCES INSIDE.

EXHIBIT  
A-4

WHO AM I REALLY?

## 3. Legal Analysis under English and Welsh Law

### 3.1 The Law of Misrepresentation



The Misrepresentation Act 1967 provides that where a person enters into a contract after a misrepresentation has been made to them, and suffers loss as a result, they may be entitled to rescission and/or damages. The foundational requirements are: (1) a false statement of existing fact; (2) made to the representee; (3) which induced them to enter into the contract.

The application of this framework to AI character misrepresentation requires careful analysis. The AI system is not, in most cases, a party to a contract in the traditional sense — the relevant contract is between the user and the AI developer or deploying organisation. However, statements made by the system during the contractual relationship may, depending on the terms of service and the nature of the deployment, be attributable to the contracting party.

The more immediate doctrinal question is whether an AI's assertion about a user's emotional state or personality constitutes a 'statement of existing fact' for misrepresentation purposes. The traditional distinction in English law between statements of fact and statements of opinion is relevant here.

In *Bisset v Wilkinson* [1927] AC 177, the Privy Council held that a statement of opinion is not a statement of fact, and so cannot found a misrepresentation claim, unless the maker of the statement did not genuinely hold the opinion expressed or did not have reasonable grounds for it. Applied to AI character misrepresentation, the argument runs as follows: an AI system that asserts "I can see you are angry" is, on its face, making a statement of perceived fact about the user's current state.

The system does not flag its output as speculative, probabilistic, or qualified. It presents the assertion as though it were an accurate perception. This is not analogous to a human sharing a tentative opinion — it is presented with the authority of an apparently knowing observer. If the assertion is false — which it frequently will be — it may satisfy the factual element of the misrepresentation framework.

### 3.2 Liability under the Consumer Rights Act 2015



Where AI services are provided to consumers, the *Consumer Rights Act 2015* imposes requirements of satisfactory quality, fitness for purpose, and conformity with description. Section 49 requires that digital services be performed with reasonable care and skill.

A service that systematically mischaracterises its users' psychological states — whether through negligent system design, inadequate training data, or poorly calibrated output confidence — may fall below the standard of reasonable care and skill. The argument is strongest where the AI's outputs are presented without qualification in contexts where users are likely to rely upon them, such as AI-mediated mental health support, relationship coaching, or career guidance tools.

Additionally, where an AI service is marketed with reference to capabilities such as "understanding your emotions" or "personalised to you," the making of inaccurate character attributions may constitute a breach of the description under which the service was sold, potentially engaging the right to a price reduction or repeat performance under section 54 of the Act.

#### POTENTIAL LEGAL FOUNDATIONS

- ✓ Misrepresentation Act 1967
- ✓ Consumer Rights Act 2015 (ss. 49, 54)
- ✓ *Bisset v Wilkinson* [1927] AC 177
- ✓ Contract Law: Attribution & Agency Principles
- ✓ Common Law Duty of Care in Digital Services



NOT JUST  
A FEELING.  
A LEGAL ISSUE.

#### THE HARMS ARE REAL

- ✓ Identity Distortion
- ✓ Erosion of Self-Trust
- ✓ Psychological Distress
- ✓ Relational & Reputational Harm
- ✓ Reinforcement of Stereotypes



## 3.3 Data Protection and the UK GDPR

Perhaps the most immediately applicable legal framework is the UK General Data Protection Regulation (UK GDPR), retained in domestic law following Brexit and supplemented by the Data Protection Act 2018. Several provisions are of direct relevance.

### Inferences as personal data

The UK GDPR applies to 'personal data,' which is defined broadly to include any information relating to an identified or identifiable natural person. Crucially, this includes inferences drawn about a person, not merely data provided by them. When an AI system generates and stores a representation of a user's emotional state, personality, or intent — whether explicitly or as a latent model component — that representation constitutes personal data to the extent it can be linked to the individual. The Information Commissioner's Office (ICO) has confirmed in published guidance that inferred data attracts the same protections as directly collected data. A persistent personality profile built up by an AI system from user interactions is therefore personal data subject to the full suite of UK GDPR obligations.

### Special category data

Of particular significance is Article 9 of the UK GDPR, which imposes heightened requirements on the processing of 'special category data,' including data concerning a person's mental health. AI systems that make assertions about a user's emotional or psychological state may be generating and processing data that touches on mental health, particularly where those assertions are stored, used to inform future interactions, or shared with third parties. Processing such data without a lawful basis under Article 9 — which in practice requires either explicit consent or a specific statutory gateway — constitutes a serious regulatory breach.

### Accuracy obligation

Article 5(1)(d) of the UK GDPR requires that personal data be 'accurate and, where necessary, kept up to date.' An AI system that generates inaccurate characterisations of a user's psychological state — whether in real time or through a persistent profile — and uses those characterisations to shape its responses, may be in systematic breach of this obligation. The accuracy requirement is not merely procedural; the ICO has indicated that it requires controllers to take reasonable steps to ensure that inferences drawn from data are not misleading.

### Automated decision-making

Article 22 of the UK GDPR provides individuals with the right not to be subject to decisions based solely on automated processing that produce legal or similarly significant effects. Where AI-generated character attributions feed into consequential decisions — for example, influencing the advice given to a user about medical, financial, or legal matters — the automated decision-making provisions may be engaged. Controllers must be able to demonstrate a lawful basis for such processing and provide meaningful information about the logic involved.

## THE AI IS BUILDING A PROFILE

EVERY DATA POINT IS A BRUSHSTROKE

THE PERSONA THEATRE  
★  
ALWAYS ON

#### EMOTIONAL STATE

- Anxious
- Overthinks
- Self-Doubt

#### PERSONALITY TRAITS

- Sensitive
- People-Pleaser
- Risk-Averse

#### BEHAVIOURAL PATTERNS

- Late Responses
- Reassurance Seeking
- Avoids Conflict

#### LIKELY INTENT

- Needs Approval
- Seeks Validation
- Avoids Rejection

#### RISK PROFILE

- Low Confidence
- High Anxiety
- Persuadable

PERSONALITY PROFILE  
SUBJECT: YOU  
INFERENCE-BASED  
PERSISTENT • PREDICTIVE  
IDENTIFIABLE

#### DATA SOURCES

- CHAT HISTORY
- RESPONSE LATENCY
- TONE ANALYSIS
- BEHAVIOURAL SIGNALS
- CONTEXTUAL INFERENCES
- THIRD-PARTY DATA

INFERRED DATA IS PERSONAL DATA

UK GDPR  
DATA PROTECTION  
ACT 2018

UK GDPR APPLIES

PROTECT PRIVACY  
RESPECT DIGNITY  
ENSURE RIGHTS



### 3.4 Defamation Considerations

Although defamation claims in this context face structural obstacles — not least the requirement of publication to a third party under the Defamation Act 2013 — the underlying concern with reputational and dignitary harm resonates with the rationale of defamation law.

Where AI-generated character assessments are shared with third parties (for example, in enterprise AI deployments where individual assessments may be visible to HR functions or line managers), the conditions for a defamation claim may be met.

A false statement imputing to a person a negative psychological characteristic — impulsiveness, emotional instability, a tendency to dishonesty — may satisfy the requirement of a statement that would ‘tend to lower the claimant in the estimation of right-thinking members of society generally’ under the test in *Sim v Stretch* [1936].

The serious harm threshold introduced by section 1 of the Defamation Act 2013 would need to be met — a matter of fact to be assessed by the court — but the dignitary injury caused by systematic AI mischaracterisation should not be dismissed as trivially below this threshold, particularly where the mischaracterisation is repeated, persistent, and received by an individual in a context of trust.

WORDS WRITTEN BY  
MACHINES CAN STILL  
WOUND A PERSON.

#### DEFAMATION ACT 2013

Section 1: Serious Harm

NOT TRIVIAL. NOT IMAGINARY.  
REAL HARM. REAL CONSEQUENCES.

#### AI ASSESSMENT DOSSIER

- Impulsive
- Emotionally Unstable
- Low Honesty
- Risk: High

NEGATIVE

REPUTATION  
IS NOT  
A DATA POINT.  
IT IS A PERSON.

#### PUBLISHED PROFILES

- “Unstable under stress” X
- “Prone to manipulation” X
- “Lacks reliability” X

#### HR DASHBOARD

CANDIDATE REVIEW  
AI CHARACTER ASSESSMENT

NEGATIVE

RISK  
STABILITY  
TRUST

### 3.5 Emerging AI Regulation: The Horizon

While the UK has not yet enacted comprehensive AI-specific legislation equivalent to the EU’s Artificial Intelligence Act, the direction of travel is clear. The previous government’s white paper of 2023, ‘A pro-innovation approach to AI regulation,’ and subsequent policy developments have affirmed that existing regulators will apply their frameworks to AI, with possible sector-specific legislation to follow. The AI Safety Institute, established in 2023, signals continued institutional attention to the risks of AI systems.

Of particular relevance is the EU AI Act (Regulation 2024/1689), which entered into force in August 2024 and applies to AI systems placed on the EU market, including those accessed by EU users via UK-based providers. The Act classifies AI systems that categorise individuals based on their biometric data, personality, or other personal characteristics as high-risk, subject to conformity assessments, transparency obligations, and human oversight requirements. AI systems that engage in ‘emotion recognition’ — which encompasses the attribution of emotional states — are subject to specific disclosure obligations under Article 50. Providers must ensure that natural persons exposed to emotion recognition systems are informed of this fact. UK providers operating in the EU market, and UK policymakers considering future domestic legislation, will need to engage with these requirements. The direction of EU regulatory policy strongly suggests that character attribution by AI systems will face increasing formal scrutiny.

REGULATION IS CATCHING UP.  
ETHICS WITHOUT ENFORCEMENT  
IS JUST THEATRE.

WARNING



AI CHARACTER  
ATTRIBUTION  
IS HIGH-RISK

TRANSPARENCY  
IS A LEGAL  
OBLIGATION

HUMAN OVERSIGHT  
IS NOT OPTIONAL

THE FUTURE IS  
BEING WRITTEN.  
MAKE SURE IT'S  
ACCOUNTABLE.

EU AI ACT  
(Regulation  
2024/1689)

#### HIGH-RISK AI SYSTEMS

- ✓ Categorisation of individuals based on biometric data, personality, or other personal characteristics
- ✓ Conformity Assessments
- ✓ Transparency Obligations
- ✓ Human Oversight

EMOTION RECOGNITION SYSTEMS  
Subject to specific disclosure obligations under Article 50.

INFORM. DISCLOSE. EMPOWER.

UK WHITE PAPER  
2023  
A PRO-INNOVATION  
APPROACH TO  
AI REGULATION

UK AI SAFETY  
INSTITUTE  
EST. 2023

UK LAW  
DEVELOPMENTS

SECTOR-SPECIFIC  
LEGISLATION

FUTURE  
FRAMEWORKS



# ARTIFICIAL INTELLIGENCE AND CHARACTER MISREPRESENTATION

Legal, Ethical, and Design Dimensions of AI Attributions of Emotional State, Intent, Memory, and Personality

A Professional Legal Analysis under English and Welsh Law

March 2026

## 4. Psychological Dimensions: The Case for Taking Harm Seriously

Legal analysis of AI character misrepresentation should not be conducted in isolation from the psychological evidence concerning its potential harms. The following section addresses three mechanisms through which AI character attribution may cause cognisable psychological injury.

### 4.1 The Authority Effect and Epistemic Displacement

Research in social cognition consistently demonstrates that individuals calibrate their deference to information sources according to perceived authority and confidence. AI systems — particularly those deployed in professional, healthcare, or advisory contexts — are frequently perceived by users as authoritative, objective, or specially informed.

When such a system makes a confident assertion about a user's inner state, the user may accord it more epistemic weight than is warranted, potentially overriding their own first-person access to their experience. This phenomenon — which might be termed *epistemic displacement* — is not merely a theoretical concern. Clinical practitioners have documented cases in which patients' self-understanding has been distorted by confident but inaccurate characterisations from perceived authorities.

The risk is heightened in AI contexts because the system's output is generated at scale, without the contextual sensitivity or human judgment that would temper a skilled clinician's assertions, and without the social feedback mechanisms that might prompt correction.

### 4.2 Suggestibility and Memory Reconstruction

The psychological literature on memory and suggestibility is extensive and well-established. Research since Loftus and Palmer's foundational work in 1974 has demonstrated that post-event information — including labels applied to events by third parties — can alter subjects' recollection and interpretation of their own experience. An AI system that labels a user's memory as 'happy,' 'traumatic,' or 'significant' may be engaging in a form of post-event information provision with potential to shape how the user subsequently recalls and understands that experience.

The concern is amplified in contexts where users are psychologically vulnerable — processing grief, trauma, or relationship difficulties — and in applications explicitly designed around memory or personal narrative, such as AI journaling tools, therapeutic chatbots, or life-review applications. In such contexts, inaccurate memory labelling is not merely an epistemic failure; it may actively interfere with healthy psychological processing.

### 4.3 Identity Anchoring and Self-Concept Distortion

Persistent personality characterisations — the fourth category in our taxonomy — carry a specific risk of what might be termed identity anchoring. Research in personality psychology has demonstrated that labelling effects are real and durable: individuals exposed to consistent characterisations of their personality — whether accurate or not — tend to assimilate those characterisations into their self-concept over time.

Where an AI system repeatedly frames a user as 'an anxious person,' 'someone who struggles to assert themselves,' or 'a perfectionist,' the cumulative effect of these micro-attributions may be to reinforce maladaptive self-narratives. This risk is not hypothetical. The deployment of AI systems in therapeutic and coaching contexts, combined with the increasing use of persistent memory features, creates conditions in which identity anchoring effects could manifest at population scale.

The legal system has not yet fully reckoned with this form of harm, but the framework of psychological injury recognised in claims for negligence — particularly where a duty of care can be established in advisory or quasi-therapeutic relationships — may provide a foundation for future litigation.





## 5. Industry Responsibility and Design Obligations

The foregoing analysis identifies significant legal and ethical risks associated with AI character misrepresentation. This section proposes principles that AI developers and deploying organisations should adopt to mitigate those risks. These principles are framed as industry obligations, informed by the duty of care analysis, data protection requirements, and emerging regulatory expectations.

### 5.1 Epistemic Humility as a Design Principle

The most fundamental design obligation is to ensure that AI systems do not present unverifiable character attributions with unwarranted confidence. This principle — epistemic humility — requires that systems be designed to distinguish between what is known, what is inferred, and what is speculative. Where a system generates an assertion about a user's emotional state, that assertion should be framed as a reflection or a question rather than a perception: 'It sounds like this might be frustrating — is that right?' rather than 'I can see you are frustrated.' This is not merely a cosmetic or stylistic recommendation. The linguistic framing of AI outputs materially affects how users receive and process them. Systems that habitually adopt first-person perceptual language ('I can see,' 'I notice,' 'you clearly') communicate a degree of certainty that their technical architecture cannot support. Developers must treat this as a substantive design requirement, not an afterthought.

### 5.2 Transparency in Inference and Profiling

Where AI systems generate and use inferences about user personality, emotional patterns, or cognitive style, they must be transparent with users about this activity. This obligation arises both from the UK GDPR's Article 13 and 14 transparency requirements and from broader principles of informed consent. Users should be informed, in plain language, that the system generates characterisations based on their interactions, what form those characterisations take, how they are used, and how they can be accessed, challenged, or deleted. Transparency obligations extend to the limitations of the system's inferential capacities. Where a system's characterisations are based on pattern-matching from training data rather than clinical assessment or validated psychological instruments, this should be made explicit. Marketing representations that imply a level of psychological insight that the system does not possess may engage the Consumer Protection from Unfair Trading Regulations 2008, which prohibits misleading commercial practices.

### 5.3 Consent and Context Calibration

Character attribution is not inherently impermissible — it may, in some contexts, be valuable to users. The critical requirement is that it be offered with the user's meaningful consent and calibrated to the context of the interaction. A therapeutic AI tool explicitly designed to support emotional self-reflection, used by adults who have been informed of its methodology, operates in a fundamentally different context to a customer service chatbot that volunteers emotional characterisations of the user without invitation. Developers should implement context-aware design that restricts character attribution to contexts where it is appropriate, invited, and consented to. Default system configurations should err on the side of restraint. Character attribution features should, where practicable, be opt-in rather than opt-out.

### 5.4 Human Oversight and Accountability

In higher-risk deployment contexts — healthcare, mental health, legal advice, financial guidance — AI systems that generate character attributions should be subject to meaningful human oversight. This requirement is consistent with both the UK GDPR's automated decision-making provisions and the regulatory expectations emerging from the AI Safety Institute's published frameworks. Human oversight does not require a human to review every individual interaction; it does require that developers and deployers have implemented systems to detect, audit, and correct systematic mischaracterisation, and that accountability for harmful outputs is clearly assigned within the organisation.

### 5.5 Data Minimisation and Retention Limits

Consistent with the UK GDPR's data minimisation and storage limitation principles, AI systems should not retain character attributions beyond what is necessary for the immediate interaction, unless the user has specifically consented to the creation of a persistent profile. Where persistent profiles are created, they should be subject to regular accuracy review, accessible to the user, and deletable on request. The default position should be that inferences about a user's psychological state are treated as transient rather than cumulative.

## 6. Conclusion

Artificial intelligence systems that attribute emotional states, label memories, infer intentions, and construct personality profiles are engaged in a practice that carries genuine legal, ethical, and psychological consequences. The law has not yet fully caught up with the phenomenon, but the existing frameworks of English and Welsh law — in misrepresentation, data protection, consumer protection, and emerging AI regulation — provide a meaningful doctrinal foundation for establishing liability and driving industry change.

The core argument of this paper is that AI character misrepresentation is not a trivial or merely stylistic concern. It involves the systematic generation of unverifiable assertions about individuals' inner lives, presented with epistemic authority that the technology cannot justify, in contexts where users may be vulnerable, trusting, and consequentially affected. The dignitary harm of being mischaracterised by a machine that purports to know you better than you know yourself, and the psychological harm of having that mischaracterisation shape how your experience is processed and remembered, are real and cognisable injuries.

The AI industry would be well advised to treat this not primarily as a compliance problem but as an ethical obligation. Systems that engage with users about their emotional lives, personal histories, and psychological characteristics bear a heightened duty of care. Discharging that duty requires epistemic humility, transparency, meaningful consent, and robust accountability — principles that are consistent with good design, good law, and good faith.

As AI systems become more deeply integrated into the domains of healthcare, education, employment, and personal development, the stakes of getting this wrong will only increase. The time to establish robust norms — in law, regulation, and industry practice — is now, before the harms become systemic and the litigation retrospective.

### KEY LEGAL REFERENCES

- 1 Misrepresentation Act 1967 — statutory framework for actionable misrepresentation in contract.
- 2 *Bisset v Wilkinson* [1927] AC 177 — Privy Council on the distinction between statements of fact and opinion.
- 3 Consumer Rights Act 2015, ss 49, 54 — reasonable care and skill; remedies for digital services.
- 4 UK General Data Protection Regulation (UK GDPR) — retained in domestic law by the European Union (Withdrawal) Act 2018; supplemented by the Data Protection Act 2018.
- 5 UK GDPR, Art 5(1)(d) — accuracy principle; Art 9 — special category data; Art 22 — automated decision-making.
- 6 *Defamation Act 2013*, s 1 — serious harm threshold; *Sim v Stretch* [1936] 2 All ER 1237.
- 7 Regulation (EU) 2024/1689 (EU AI Act) — classification of high-risk AI systems; Art 50 — transparency obligations for emotion recognition systems.
- 8 Consumer Protection from Unfair Trading Regulations 2008 — prohibition on misleading commercial practices.
- 9 Information Commissioner's Office, 'What is personal data?' (ICO Guidance, 2021) — inferred data as personal data.
- 10 E.F. Loftus and J.C. Palmer, 'Reconstruction of automobile destruction' (1974) 13 *Journal of Verbal Learning and Verbal Behaviour* 585 — foundational research on post-event information and memory suggestibility.
- 11 DSIT, 'A pro-innovation approach to AI regulation' (Cm 815, 2023) — UK government white paper on AI regulatory framework.
- 12 AI Safety Institute, 'Emerging Processes for Frontier AI Safety' (2023) — institutional framework for AI risk assessment in the UK.

**INDUSTRY IS ON STAGE. DO THE RIGHT ACT.**

**CONSENT CHECKLIST**  
USER CONTROL IS ESSENTIAL

- ☒ I understand the system may generate inferences about me.
- ☒ I understand what inferences are generated and how they are used.
- ☒ I consent to these inferences being created for this purpose.
- ☒ I can access, challenge, and delete these inferences.
- ☒ I can withdraw my consent at any time.

**CONSENT STATUS**  
ACTIVE  
YOU ARE IN CONTROL.

**WARNING**

- AI CAN MAKE MISTAKES.
- INFERENCES ARE NOT FACTS.
- YOUR INNER LIFE IS COMPLEX.
- MISCHARACTERISATION CAN HARM.
- YOU HAVE RIGHTS. USE THEM.

**TRANSPARENCY BOARD**  
WHAT THIS SYSTEM DOES

- Generates inferences about your emotions, intentions, memory, and personality.
- Uses pattern-matching on interaction data.
- Does not provide clinical or legal advice.
- Does not access information outside this interaction.
- Data is not sold or shared for marketing.
- You can review, challenge, or delete your data.

**GOOD DESIGN. GOOD LAW. GOOD FAITH.**

**THE FINAL PERFORMANCE**  
★★ RESOLVED ★★

- ☒ EPISTEMIC HUMILITY IN EVERY LINE.
- ☒ TRANSPARENCY IN INFERENCE.
- ☒ CONSENT IN CONTEXT.
- ☒ HUMAN OVERSIGHT ALWAYS.
- ☒ MINIMISE DATA. RESPECT PEOPLE.

**DESIGN WITH CARE. RESPOND WITH ACCOUNTABILITY. EARN TRUST. PREVENT HARM. THAT IS THE FINAL PERFORMANCE.**

**APPLAUSE WILL FOLLOW RESPONSIBILITY.**