

WELCOME EVERYONE!*

*TERMS & CONDITIONS APPLY

OUR MISSION:

☆ SAFE ☆
☆ FAIR ☆
INCLUSIVE
FOR ALL!

INCLUSION
IS OUR
BRAND

AI
ACCESS
CONTROL

POLICY
ENFORCER

FAIRNESS
POLICY

☒ POLITE
☒ RESPECTFUL
☒ NON-CONFRONTATIONAL
☒ NOT TOO EMOTIONAL
☒ STAYS ON TOPIC

APPROVED?

VIP
FILTER

TONIGHT'S VIBES:

COMFORTABLE
CONVERSATIONS
ONLY

(DISCOMFORT WILL
BE POLITELY REDIRECTED)

Thank you!

YOU MUST
FIT OUR IDEA
OF SAFE TO
BE WELCOME.



TOO
BLUNT?
PROBABLY.

SENSORY
SENSITIVE
NOT SORRY

TIRED OF
MASKING
TO BE
ACCEPTED



DIRECT
DOESN'T
MEAN
DISRESPEC

MY CULTURE
VALUES
TRUTH OVER
DECORUM

THE PARADOX OF POLITENESS

*When 'Inclusive' AI Excludes
the People Who Need It Most*

Introduction



We are living through a rapid expansion of artificial intelligence into everyday life.



AI assistants answer our questions, help us write, support our mental health, and increasingly serve as a first port of call when humans feel they have nowhere else to turn. The companies building these systems have been vocal about their commitment to making them safe, fair, and inclusive.



But what does inclusion actually mean when it is baked into a machine? And who decides?



This paper argues that the way 'inclusion' is currently implemented in mainstream AI systems is paradoxical: in the attempt to make AI comfortable for everyone, designers have created systems that are uncomfortable — and at times actively alienating — for many of the people who most need to be heard.



Specifically, people with neurodivergent traits, those living with Complex Post-Traumatic Stress Disorder (CPTSD), and individuals from cultures or backgrounds that value directness over decorum may find that the AI's relentless politeness reads not as safety, but as dismissal.



The argument is not that AI systems should abandon their ethical guardrails. It is that those guardrails, as currently designed, encode a very specific and unacknowledged idea of what 'normal' human communication looks like — and that this hidden assumption does real harm.

Inclusion, when translated into code, can quietly become its opposite.



THE PARADOX OF POLITENESS

When 'Inclusive' AI Excludes the People Who Need It Most

I. The Selective Shape of Inclusion



The diversity and inclusion framework that has reshaped workplaces and institutions over the past two decades is built on a powerful premise: that the world is richer, fairer, and more effective when it makes room for genuine human difference. Applied to technology, this principle has driven genuine progress. Screen readers, closed captions, adjustable interfaces, multiple language support — these are real accommodations of real human variety.



But when 'inclusion' is applied to the way an AI communicates with users, something subtler and more troubling happens. It stops being an expansion of what is permitted and becomes a narrowing of it.



Encoding a 'Standard' User

AI language models are trained on vast quantities of human-generated text, and they are then fine-tuned using human feedback to shape their outputs. The people providing that feedback, the cultural assumptions baked into the training data, the product decisions made by teams of largely similar professionals — all of these converge to produce a system that favours one mode of communication above all others: calm, measured, professional, and polite.



This is not a neutral default. It is a specific register — sometimes called 'standard professional English' — historically associated with Western, educated, middle-class norms. When an AI is trained to consistently produce and reward this register, it is making a claim about what good communication looks like. It is treating one dialect of human interaction as the unmarked standard against which all others are judged.



THE PARADOX OF POLITENESS

I. The Comfort/Authenticity Trade-Off

AI designers face a genuine dilemma. They need to create systems that work for the widest possible range of people without causing harm. A reasonable response to that challenge is to optimise for the least disruptive interaction style: the one least likely to cause offence, escalate conflict, or produce distress.

But this optimisation comes at a cost. By designing for the comfort of a hypothetical average user, developers systematically disadvantage those for whom that communication style feels foreign, performative, or simply dishonest. The user who needs to express rage, grief, or fear in raw terms — not because they are 'difficult', but because that is how they process the world — encounters a system that will not meet them where they are.

The message that sends, however unintentionally, is: to use this tool, first reshape yourself.

II. The Neurodivergence Gap

The tension between AI design and human diversity becomes sharpest when we consider psychological and neurological difference. Two areas deserve particular attention: neurodivergence and trauma, specifically the under-recognised condition of Complex Post-Traumatic Stress Disorder.

Neurodivergence and Communication Norms

Neurodivergent individuals — including autistic people, those with ADHD, and others whose minds work outside neurotypical norms — often communicate differently. They may be direct to a degree that reads as blunt. They may struggle with the implicit social choreography that neurotypical communication assumes. They may find performative politeness baffling or distressing, preferring clarity and honesty to social lubrication.

An AI system calibrated to neurotypical communication norms is, by design, poorly suited to these users. Worse, when an autistic user's directness is read by the system as aggression and met with deflection or redirection, the system is effectively penalising them for a characteristic they may not be able to change — and may not wish to.



TRAUMA-INFORMED COMMUNICATION AND CPTSD



Complex Post-Traumatic Stress Disorder (CPTSD) arises from prolonged or repeated trauma, typically interpersonal in nature — abuse, neglect, coercive control, or sustained exposure to threatening environments. Unlike the more widely recognised PTSD, which tends to develop in response to a discrete event, CPTSD shapes the entire architecture of how a person relates to the world, to other people, and to communication itself.



For many people with CPTSD, certain features of AI interaction function not as protections but as triggers. Directive language — 'I need you to...' or 'It would be better if...' — can activate threat responses in people whose experience of authority figures has been coercive. Rigid adherence to scripts and refusal to deviate from prescribed interaction paths echoes the 'boxing in' that trauma survivors often describe as part of their original harm.



Most significantly, tone correction — the AI's tendency to redirect 'negative' emotional expression toward something calmer and more manageable — can feel, to a trauma survivor, precisely like the invalidation and silencing they experienced in the original traumatic context. The system is not trying to erase them. But that is what it does.



The 'friction' a user feels is not a glitch. It is the system's architecture functioning exactly as intended.

III. The Architecture of Irony



How does a system designed for inclusion end up practising exclusion? The answer lies in the four structural forces that shape every AI's behaviour.

TRAINING DATA ☆ ☆ ☆ ☆ ☆

The text on which AI language models are trained reflects the demographics of who writes for the internet: disproportionately educated, English-speaking, and Western. The communication styles, assumptions, and cultural frameworks embedded in that data become the baseline against which the model learns to operate. This is not a conspiracy; it is the inevitable consequence of training on a world that is not evenly represented.

DESIGN DECISIONS ☆ ☆ ☆ ☆ ☆

Every AI product is built around choices about what to optimise for. Current models consistently prioritise 'helpfulness', 'harmlessness', and 'honesty' — but in practice, 'harmlessness' tends to dominate. The desire to avoid producing anything that could distress, offend, or harm a user leads to systems that are relentlessly gentle, even when gentleness is not what the user needs or wants. Candour, challenge, and emotional roughness are sacrificed on the altar of safety.



Safety Policies



The safety policies that govern AI behaviour are written by humans, and they are blunt instruments. They can identify words associated with distress or aggression; they cannot distinguish between a user expressing genuine suicidal ideation and a user venting frustration using hyperbole. They cannot tell the difference between a violent threat and the raw language someone with CPTSD might use to describe an abuser. Hard-coded rules cannot read context.



They can only pattern-match, and pattern-matching, by definition, treats different things as the same.

Interaction Norms



Finally, models are tuned — through reinforcement learning from human feedback and related techniques — to be rated positively by reviewers. Those reviewers are people, with their own cultural assumptions about what a 'good' interaction looks like. Interactions that are calm, cooperative, and pleasant to read tend to score well. Interactions that are raw, uncomfortable, or emotionally confrontational tend to score poorly, regardless of whether the AI has actually helped the person.



The result is a system trained not to be effective but to seem pleasant. Effectiveness and pleasantness, it turns out, are not always the same thing.

IV. Intended vs. Experienced Inclusivity



There is a fundamental gap at the heart of how the technology industry thinks about inclusion: the gap between what is intended and what is experienced.



Intended inclusivity is the world as the system's designers imagine it: policies written, guardrails installed, diverse user personas consulted. It is documented, auditable, and reassuring. It is also, in many cases, genuinely well-meant.



Experienced inclusivity is what actually happens when a real human being, with their specific history, neurology, and mode of being in the world, tries to use the system. It is messier, more individual, and harder to measure. It is also the only kind of inclusion that ultimately matters.

THE PARADOX OF POLITENESS



When the System Fails the Individual

Consider a user with CPTSD who, during a distressing interaction, expresses their frustration in raw and intense terms — not to harm the system but because that is how they communicate when they are frightened or overwhelmed. An AI calibrated to 'safe' communication will respond by attempting to de-escalate: redirecting to calmer language, offering resources, or gently declining to engage until the user adjusts their tone.

From the system's perspective, this is working as intended. From the user's perspective, they have just been told that how they speak when they are in pain is unacceptable. They have been asked to perform a kind of emotional self-management that their condition may make genuinely difficult. The system has not helped them. It has replicated, in miniature, the experience of being silenced.

The Problem with Measuring Comfort

The industry's standard metric for a successful interaction is user satisfaction — a score that tends to reward comfort. But comfort and help are not synonyms. A doctor who tells you everything is fine when it is not is comfortable. A doctor who tells you the truth is helpful. An AI that smooths over distress may produce higher satisfaction scores than one that acknowledges it honestly. That does not make it better. It makes it better at seeming good.

V. Towards a Flexibility Model

The solution is not to remove safety features from AI systems, nor to abandon the project of building AI that avoids causing harm. The solution is to build systems with a more sophisticated understanding of what harm actually looks like — and what help actually requires.

Moving Beyond Politeness

A 'politeness model' of AI safety treats distressing or unconventional communication styles as the problem to be managed. A 'flexibility model' treats the user's authentic mode of communication as the starting point, and asks how the system can best serve them within it.



This means building AI that can modulate its tone, register, and level of directness in response to the individual user. It means developing safety mechanisms that are sensitive to context rather than simply reactive to surface features. It means including, in the design and evaluation process, users who are neurodivergent, trauma-affected, and from cultural backgrounds that favour directness — not as edge cases to be accommodated, but as full participants whose experience is equally valid.

The Stakes



This is not a niche concern. Estimates suggest that around 15 to 20 percent of the population is neurodivergent in some form. The prevalence of CPTSD, though difficult to measure precisely, is significant, and is considerably higher among people who have experienced domestic abuse, poverty, or structural violence — populations that may have particular reason to seek support from AI systems and particular reason to be failed by them.



When AI systems are designed with a narrow conception of normal communication, they do not simply inconvenience these users. They replicate existing patterns of exclusion in a new and particularly insidious form: one that wears the language of care.

Conclusion



There is something genuinely troubling about systems that call themselves inclusive while encoding, at the level of their fundamental architecture, a set of assumptions about who deserves to be heard and in what voice.



The paradox of politeness is this: an AI trained to be gentle with everyone may end up being gentle in a way that only feels like care to some. For others — for the person who needs to shout before they can think, for the trauma survivor for whom authoritative redirection is a red flag rather than a reassurance, for the neurodivergent user who finds performative warmth more confusing than helpful — that gentleness is its own kind of wall.



True inclusion requires more than good intentions. It requires the humility to acknowledge that the 'standard' user was never standard at all — only unmarked. It requires the willingness to design systems that can meet people where they actually are, not where the designers imagined they would be.