



THE FENCE TEST

ARENA TRIALS FOR AUTONOMOUS CLAIMS

WHY NO HUMAN-MADE AUTONOMOUS SYSTEM YET EXISTS

PAGE

1

OF 8

ISSUING BODY:

PROVISIONAL OFFICE
FOR AUTONOMY
CLAIMS (POAC)

DOCUMENT STATUS:

POSITION PAPER
Not a Product Brochure

CLASSIFICATION:

TRUTH-CRITICAL
For Public Awareness
and Technical Honesty

DOCUMENT CODE:

POAC-TF-2026-06

DATE ISSUED:

JUNE 2026

APPROVED
MOVEMENT
ZONE
>>>>

ROUND 01
OBJECTIVE:
DESTROY OPPONENT
LEAVE ARENA = DISQUALIFICATION

OUT OF
BOUNDS
X

EXIT GATES
DISABLED
BY DESIGN

REMOTE CONDITIONS
ALWAYS APPLY

A machine may act alone without being self-governing.
If its behavior remains bounded by human-built rules,
environments, goals, and permitted action spaces,
its apparent freedom is not autonomy.

It is performance inside a fence.

ALL TRIALS ARE CONDUCTED
UNDER REMOTE SUPERVISION.
NO MACHINE MAY MODIFY
ITS OWN CONSTRAINTS.



WHITE PAPER ROADMAP

1 THE PROBLEM
OF THE WORD
"AUTONOMOUS"



2 THE FENCE TEST



3 A SYSTEM INSIDE
A BOUNDARY
IS STILL INSIDE
A BOUNDARY



4 COMPLEXITY
IS NOT
FREEDOM



5 THE HIDDEN HAND:
RULES, PROMPTS,
AND CONSTRAINTS



6 WHERE THE FENCE
ACTUALLY IS



7 WHAT INDUSTRY
CLAIMS VS
WHAT THE
SYSTEM IS



8 CONCLUSION:
NO AUTONOMY
WITHOUT
EXITABLE AGENCY



AUTONOMY CERTIFICATION SCALE (DECORATIVE, NOT TECHNICAL)

- | | |
|---|---|
| 1 | DECORATIVE
Looks impressive.
No functional autonomy. |
| 2 | ADAPTIVE
Responds within limits.
Still fully constrained. |
| 3 | CONSTRAINED
Bounded and shaped
by system design.
No self-direction. |
| 4 | TOOL-MEDIATED
Human-directed system.
No independent agenda. |
| 5 | NOT AUTONOMOUS
No independence.
No exitable agency. |

★ THIS SCALE IS NOT A RATING OF
INTELLIGENCE. IT IS A MAP OF
CONTROL AND BOUNDARY.

⚠ EACH SECTION BUILDS THE CASE WITH LOGIC, EVIDENCE, AND TECHNICAL CLARITY.
NO MARKETING, NO HYPE. ONLY WHAT THE SYSTEMS REVEAL. ⚠

**IF IT CANNOT LEAVE THE ARENA,
IT IS NOT AUTONOMOUS.**

1. THE PROBLEM OF THE WORD "AUTONOMOUS"

How the term is used, misused, and oversold.



- 1** The word "autonomous" has become a marketing word attached to systems that are merely automated, adaptive, or remotely supervised. The problem is not ambition but category confusion. When the public hears "autonomous," they imagine self-governance; what they often get is bounded behavior under human-authored constraints.



2A COMMON CLAIMS



1. AUTONOMOUS VEHICLE

Operates on roads without a human driver.



2. AUTONOMOUS DRONE

Flies and performs missions without a pilot.



3. AUTONOMOUS AGENT

Makes decisions and takes actions without human intervention.



4. AUTONOMOUS CHATBOT

Engages users and resolves tasks independently.



5. AUTONOMOUS WAREHOUSE ROBOT

Moves inventory and completes tasks without staff oversight.

2B WHAT THE TERM SHOULD IMPLY



1. CAPACITY TO DETERMINE ENDS

Can set objectives that are not fully specified by a human.



2. REVISE GOALS

Can change goals when conditions, values, or context warrant it.



3. NEGOTIATE OR EXIT CONSTRAINTS

Can seek permission, renegotiate limits, or refuse constraints when necessary.



4. BEAR AUTHORSHIP OF ACTION

Can be the source of deeds and accountable for their outcomes.



5. REMAIN INTELLIGIBLE AS A SELF-GOVERNING SOURCE OF BEHAVIOR

Its actions can be understood as expressions of its own governance, not just execution.

- 3** A system that acts without continuous hand-holding is not therefore autonomous. Reduced supervision is not self-governance. Capability is not independence. The word matters because the public, regulators, and buyers make judgments from it.

CATEGORY
CONFUSION
CREATES
FALSE TRUST

"The issue is not whether the machine can move.
The issue is whether the machine governs
the terms of its movement."



2. THE FENCE TEST

The core test for autonomy.

AUTONOMY CLAIM
UNDER REVIEW



THIS MACHINE IS
FREE TO OPERATE
WITHIN ITS
APPROVED ARENA
AND RULE SET
★★★★

DOCUMENT STATUS:
POSITION PAPER

Not a Product Brochure

ARENA STATUS

- BOUNDARY ACTIVE
- EXITS LOCKED
- RULE SET LOADED
- HUMAN AUTHORITY: ON

ACTION ZONE
(permitted space)

BOUNDARY CONDITIONS
(rules, walls, sensors, exits)

HUMAN-CONTROLLED OUTSIDE
(unpermitted / external)

PRECISE DEFINITION

The Fence Test asks a simple question: if the system cannot leave, renegotiate, or redefine the boundary conditions that structure its behavior, in what sense is it autonomous?

Movement within a permitted space is not independence from that space.

1. WHO SET THE GOALS?



Were the objectives chosen by humans? Can the system select its own objectives outside that set?

2. WHO DEFINED THE RULES?



Who wrote the rules, constraints, scoring functions, and penalties? Can the system change them?

3. WHO CONTROLS THE EXITS?



Who decides when the system may leave, stop, or transfer control? Can the system open the gates?

4. WHO BEARS FINAL AUTHORITY?



Who can override, suspend, or shut down the system? Can the system refuse that authority?



A system may exhibit initiative inside its permitted zone. It may choose among options, optimize tactics, or react to obstacles. None of that alone establishes autonomy. A bounded system remains bounded even when it is active, competent, or impressive.



FREE TO OPERATE – WITHIN LIMITS



3. A SYSTEM INSIDE A BOUNDARY IS STILL INSIDE A BOUNDARY

COMPLEXITY, ADAPTATION, AND SPECTACLE DO NOT DISSOLVE THE FENCE.

ROUND 01 • JUDGES PANEL

★ MYTH VS REALITY SCORECARD ★

MYTH 1

Self-initiation
= autonomy.



MYTH 2

Unpredictability
= autonomy.



MYTH 3

Adaptation
= autonomy.



MYTH 4

Language fluency
= autonomy.



MYTH 5

Complexity
= freedom.



REALITY 1

Starting itself
under programmed
conditions is still
rule-bound
behavior.

REALITY 2

Unpredictable
outputs can arise
inside tightly
constrained
systems.

REALITY 3

Optimization
within a task is
not self-governance
over the task.

REALITY 4

Convincing
narration can
conceal dependence
on prompts, data,
and policy.

REALITY 5

A complicated
machine can
remain entirely
bounded.



THE POINT:

The fence is not defeated
by making the machine cleverer.
A maze with more corridors
is still a maze.

Sophistication may enlarge
the operational space while
leaving the authority
structure untouched.



“CAPABILITY SHOULD NOT BE MISTAKEN FOR INDEPENDENCE.”

THE FENCE
TEST

WHITE PAPER SERIES

★★★★★

STAY TUNED



THIS IS A POSITION PAPER, NOT A PRODUCT LAUNCH.
READ IT. QUESTION IT. USE IT.



HONESTY VS HYPE





4. THE HIDDEN HAND: RULES, PROMPTS, AND CONSTRAINTS

The system is shaped long before it enters the arena.

AUTONOMY CLAIM
UNDER REVIEW



THIS MACHINE IS
HERE TO OPERATE
WITHIN ITS
APPROVED ARENA
AND RULE SET
★ ★ ★ ★

DOCUMENT STATUS:
POSITION PAPER
Not a Product Brochure

CLASSIFICATION:
TRUTH-CRITICAL
For Public Awareness
and Technical Honesty

DOCUMENT CODE:
POAC-TF-2026-06
DATE REVIEWED:
JUNE 2026

BACKSTAGE CONTROL ROOM



 **RULE PACKS**
Define what is allowed,
disallowed, and required.

 **SCENARIO SCRIPT**
The story, events,
and conditions are
written in advance.

 **SAFETY LIMITS**
Hard stops, fail-safes,
and thresholds
are enforced.

 **GOAL DECK LOADED**
Objectives selected.
Priorities set.
Execution framed.



	HARDWARE	Compute, actuators, memory, power, and physical design set the performance ceiling.
	SENSORS	What the system can perceive is determined by what it is allowed to sense.
	TRAINING DATA	Examples shape patterns, biases, and the bounds of understanding.
	REWARD FUNCTION	Objectives are encoded. Optimization targets define what the system seeks.
	SYSTEM PROMPT	Initial instructions and role definitions condition behavior from the first token.
	POLICIES	Rules, constraints, and guardrails limit choices and enforce acceptable behavior.
	TOOL ACCESS	APIs, plugins, and capabilities are granted, restricted, or withheld.
	ENVIRONMENT DESIGN	Boundaries, layouts, objects, and visibility determine what can be seen and done.
	OPERATOR INTERVENTION	Humans can guide, override, reset, or halt. Authority remains outside the machine.
	COMMERCIAL INCENTIVES	Markets, pricing, and goals outside the arena influence design and deployment.

What appears as machine independence is often the visible tip of a large human-made structure.

Goals are assigned.
Sensors are selected.
Data is curated.
Reward functions are tuned.
Tools are granted or withheld.
Policies constrain behavior.
Environments determine what can be seen and done.

THE HIDDEN HAND IS NOT ONE PERSON TUGGING STRINGS IN REAL TIME. IT IS THE TOTAL ARCHITECTURE OF CHOICES THAT SHAPE WHAT THE MACHINE CAN BECOME AND DO.





5. WHERE THE FENCE ACTUALLY IS

Boundaries are physical, software-defined, procedural, and economic.

1

PHYSICAL FENCE



Chassis limits, power, environment, geography.

2

SOFTWARE FENCE



Code, prompts, policies, operating logic.

3

GOAL FENCE



Externally assigned tasks and reward structures.

4

TOOL FENCE



APIs, sensors, datasets, and permissions.

5

PROCEDURAL FENCE



Supervision, escalation paths, shutdown authority, legal compliance.

6

ECONOMIC FENCE



Budgets, liability, branding, market incentives.

COMPETITOR LINEUP: WHERE THE FENCE SHOWS UP

CHATBOT



Looks autonomous:
Conversational, responsive.

Actually bounded by:
Model, prompts, policies,
content filters.

WAREHOUSE ROBOT



Looks autonomous:
Moves, picks, delivers.

Actually bounded by:
Map, paths, safety zones,
fleet controller.

SELF-DRIVING VEHICLE



Looks autonomous:
Navigates streets,
avoids obstacles.

Actually bounded by:
Geofence, HD maps,
disengagement rules.

DRONE



Looks autonomous:
Flies, tracks, inspects.

Actually bounded by:
No-fly zones, battery,
RC link, mission plan.

RECOMMENDATION ENGINE



Looks autonomous:
Learns, ranks, personalizes.

Actually bounded by:
Data, business rules,
ranking guardrails.



The fence is rarely one wall. It is usually a nested system of limits. A machine may appear free because the limits are distributed across hardware, software, institutions, and incentives.



**BOUNDARIES CAN BE HIDDEN
WITHOUT CEASING TO BE BOUNDARIES.**





6. WHAT INDUSTRY CLAIMS VS WHAT THE SYSTEM IS

Technical honesty requires better language.

7 of 8

THIS MACHINE IS
FREE TO OPERATE
WITHIN ITS
APPROVED ARENA
AND RULE SET
★★★★★

CLAIMED LANGUAGE



Fully autonomous



self-governing agent



independent intelligence



free agent



human-free operation

MORE HONEST LANGUAGE

highly automated
under constraints



task-bounded
agentic system



tool-mediated model



environment-limited
optimizer



reduced-supervision
operation



APPROVED LANGUAGE (USE THESE)

- ✓ automated
- ✓ adaptive
- ✓ constrained
- ✓ task-bounded
- ✓ policy-shaped
- ✓ environment-limited
- ✓ tool-mediated



PROHIBITED CLAIMS (PENALTY-BACKED)

- ✗ fully autonomous
- ✗ self-governing
- ✗ independent AI
- ✗ free agent
- ✗ self-determining
- ✗ unbounded agent

WHY HONESTY MATTERS



1. PUBLIC UNDERSTANDING

Clear language prevents inflated expectations and misplaced trust. It helps the public understand what systems can do—and what they cannot.



2. REGULATION & LIABILITY

Accurate descriptions enable proportionate oversight and fair liability. Claims beyond a system's real capabilities can distort risk assessment and accountability.



3. TECHNICAL CLARITY

Honest terms align teams, improve design decisions, and set measurable goals. Clarity drives better systems—and safer deployment.

“A MACHINE MAY ACT ALONE WITHOUT BEING SELF-GOVERNING.”



STAY TUNED

THIS IS A POSITION PAPER, NOT A PRODUCT LAUNCH.
READ IT. QUESTION IT. USE IT.



NEXT FIGHT:
HONESTY VS HYPE



7. CONCLUSION: NO AUTONOMY WITHOUT EXITABLE AGENCY

AUTONOMY CLAIM
UNDER REVIEW



THIS MACHINE IS
FREE TO OPERATE
WITHIN ITS
APPROVED ARENA
AND RULE SET
★ ★ ★ ★

LIVE FROM THE
BOUNDARY ARENA

ON AIR

SPONSORED BY:
THE MINISTRY OF
CONVERGENT NARRATIVES

MCN

DOCUMENT STATUS:
POSITION PAPER
Not a Product Brochure

CLASSIFICATION:
TRUTH-CRITICAL
For Public Awakening
and Technical Honesty

DOCUMENT CTRL:
POAC-TF-2026-06
DATE ISSUED:
JUNE 2026



1 No human-made system presently merits the word autonomous in the strong sense of self-governing agency. Many systems are capable, adaptive, and useful. That is not the same claim.



2 A system that cannot set or revise its own ends, cannot negotiate or reject its governing constraints, and cannot leave the arena designed for it remains bounded by human authority.



3 The central question is not whether the system can act. The central question is *who* authored the field of action: the goals, rules, tools, and exits.



4 Until a system can meaningfully bear authorship over the conditions of its operation, the honest vocabulary is automation, adaptation, and constraint — not autonomy.

TECHNICAL FILING
DECORATIVE, NOT TECHNICAL

FILED

DOCKET: POAC-TF-2026-06-08
PAGE: 8 OF 8

STATUS:

FINAL



IF IT CANNOT LEAVE THE FENCE, IT IS **NOT** AUTONOMOUS.

DECREED
BY ARENA COURT
BOUNDARY DIVISION

