



The Marked Bench

Public benchmark-only repository

Martin123132 / The-Marked-Bench-

EXHIBIT A-001

PRIMARY SUBJECT

BENCH ID: MB-0001

WITNESS NOTE #17

It's always been here."

- morning walker 06:12



WITNESS NOTE #42

No one sits for long.

Feels like it remembers.

- anonymous 18:09

OBSERVATION LOG

- ☒ Unattended
- ☒ Weathered
- ☒ Public
- ☒ Persistent

STATUS: UNDER REVIEW

EVIDENCE TAG

MB-0001

ITEM: PARK BENCH

LOCATION: RIVERSIDE PARK

COLLECTED: 2025-05-04

COLLECTED BY: PB-STD UNIT

PBSTD-250504-0001



SUBJECT: MB-0001

CLASS: PUBLIC INFRASTRUCTURE

ROLE: SILENT WITNESS

EVIDENCE MB-0001 DO NOT REMOVE

1

2

3

PERIMETER

DO NOT CROSS

2.01 m

1.72 m

SERLED

EVIDENCE LINE - DO NOT CROSS

SURVEILLANCE SEQUENCE (STILLS)

20:11:02
CAM_07

21:47:31
CAM_07

00:13:58
CAM_07

04:59:17
CAM_07

CHAIN OF CUSTODY		
☒ COLLECTED	PB-STD UNIT	2025-05-04
☒ RECEIVED	FORENSIC OPS	2025-05-04
☒ REGISTERED	DATA VAULT	2025-05-04

HANDLING INSTRUCTIONS

Handle with care.

Do not alter.

Document everything.

Canonical source: This is the public benchmark-only repository for The Marked Bench. Use this repository for all pulls, issues, and contributions related to benchmarks.

The Marked Bench is versioned benchmark infrastructure for testing whether AI systems can detect and classify contradictions between a premise and a query.

The project is designed to become a reproducible public standard: every score is tied to a suite ID, suite version, deterministic suite hash, immutable case list, JSON report schema, confusion matrix, per-class metrics, confidence calibration metrics, explanation-audit coverage, validation result, diagnostic slice metrics, result card, result claim, standard profile, implementation kit, scoring compatibility vectors, scoring specification, and leaderboard entry.

Current standardization status is tracked in:

- Project status
- Standardization status
- Baseline robustness
- Maintainer handoff

Current Health Snapshot

- `main` is expected to stay green under Benchmark CI.
- Public release evidence is regenerated deterministically and checked for drift in CI.
- Licensing, review workflow, scoring sanity, case quality, and baseline robustness each have dedicated guard scripts.
- The multi-hop track has a documented baseline-robustness watchlist item: `HashPriorBaseline` currently outranks `ContradictionEngine`.
- The recommended project posture is polish and trust-hardening before adding more benchmark surface.

FILED BY: PB-STD UNIT

DATE: 2025-05-04

FOR OFFICIAL USE ONLY



The Marked Bench

Current Tracks / Install / Run A Benchmark

Martin123132 / The-Marked-Bench-

EXHIBIT A-002

PRIMARY SUBJECT

BENCH ID: MB-0001



SURVEILLANCE STILL_02
2025-05-04 21:47:31

WITNESS NOTE #42

No one sits
for long.
Feels like it
remembers.

— anonymous
18:09

EVIDENCE TAG

MB-0001

TYPE: PARK BENCH
LOCATION: RIVERSIDE PARK
COLLECTED: 2025-05-04
COLLECTED BY: PB-STD UNIT



PBSTD-250504-0001

FIELD SKETCH



OBSERVATION LOG

- ☒ Unattended
- ☒ Weathered
- ☒ Public
- ☒ Persistent

STATUS: UNDER REVIEW

SUBJECT: MB-0001

CLASS: PUBLIC INFRASTRUCTURE

ROLE: SILENT WITNESS

EVIDENCE
MB-0001
DO NOT REMOVE

SEALED
PB-STD UNIT

PERIMETER

DO NOT CROSS

2.01 m

EVIDENCE LINE — DO NOT CROSS

EVIDENCE LINE — DO NOT CROSS

SURVEILLANCE SEQUENCE (STILLS)



CHAIN OF CUSTODY

☒ COLLECTED	PB-STD UNIT	2025-05-04
☒ RECEIVED	FORENSIC OPS	2025-05-04
☒ REGISTERED	DATA VAULT	2025-05-04

HANDLING INSTRUCTIONS

Handle with care.
Do not alter.
Document everything.

Current Tracks

Track	Suite ID	Version	Purpose	Baseline
Foundation	marked-bench-contradiction-standard	0.1.1	Expanded canonical contradiction suite	100.00
Adversarial	marked-bench-contradiction-adversarial	0.2.0	Longer-context, implicit, and trap cases	52.37
Multi-hop	marked-bench-contradiction-multihop	0.3.0	Linked-evidence contradiction cases	24.14
Controls	marked-bench-contradiction-controls	0.4.0	False-positive distractor controls with contradiction anchors	100.00

The adversarial track is intentionally not solved by the packaged symbolic baseline. The multi-hop track is the default target for future symbolic, neural, retrieval-augmented, and hybrid systems. The controls track is a false-positive stress track for systems that over-call contradictions on paraphrases, scoped negatives, time shifts, and harmless elaborations.

Install

```
pip install -e .
```

Requires Python 3.10 or newer.

Run A Benchmark

Foundation track:

```
marked-bench --suite contradiction --report  
artifacts/foundation-report.json
```

Adversarial track:

```
marked-bench --suite contradiction-adversarial --report  
artifacts/adversarial-report.json
```

Multi-hop track:

```
marked-bench --suite contradiction-multihop --report  
artifacts/multihop-report.json
```

Control track:

```
marked-bench --suite contradiction-controls --report  
artifacts/controls-report.json
```

Validate a report before publication:

```
marked-bench --validate-report  
artifacts/adversarial-report.json
```

FILED BY: PB-STD UNIT

DATE: 2025-05-04

FOR OFFICIAL USE ONLY



The Marked Bench

Run A Benchmark — Continued

Martinl23l32 / The-Marked-Bench-

EXHIBIT A-003

RUN GUIDE

BENCH ID: MB-0001

WITNESS NOTE #21

The steps are
breadcrumbs.
Follow them.

— evening walker
19:48

FIELD SKETCH



EVIDENCE TAG

MB-0001

TYPE: RUN GUIDE

LOCATION: RIVERSIDE PARK

COLLECTED: 2025-05-04

COLLECTED BY: PB-STD UNIT



RGUIDE-20250504-0003

CHAIN OF CUSTODY

☒	COLLECTED	PB-STD UNIT	2025-05-04
☒	RECEIVED	FORENSIC OPS	2025-05-04
☒	REGISTERED	DATA VAULT	2025-05-04

SURVEILLANCE SEQUENCE (STILLS)



OBSERVATION LOG

- ☐ Unattended
- ☒ Weathered
- ☒ Public
- ☒ Persistent

STATUS: UNDER REVIEW

EVIDENCE LINE — DO NOT CROSS

HANDLING INSTRUCTIONS

Handle with care.
Do not alter.
Document everything.



Export the machine-readable registry of public tracks and artifacts:

```
marked-bench --export-registry benchmark_registry.json
```

List and inspect public suites:

```
marked-bench --list-suites  
marked-bench --suite-info contradiction
```

Export the release manifest that pins public artifact SHA-256 digests:

```
marked-bench --export-release-manifest releases/marked_bench_release_v0_4_8.json
```

Export the generated technical note:

```
marked-bench --export-technical-note docs/TECHNICAL_NOTE.md
```

Export and validate the machine-readable conformance report:

```
marked-bench --export-conformance-report conformance/marked_bench_conformance_v0_4_8.json  
marked-bench --validate-conformance-report conformance/marked_bench_conformance_v0_4_8.json
```

Export and validate the machine-readable benchmark standard profile:

```
marked-bench --export-standard-profile standard/marked_bench_standard_profile_v0_4_8.json  
marked-bench --validate-standard-profile standard/marked_bench_standard_profile_v0_4_8.json
```

Export and validate the checked standard change-control profile:

```
marked-bench --export-change-control standard/marked_bench_change_control_v0_4_8.json  
marked-bench --validate-change-control standard/marked_bench_change_control_v0_4_8.json
```

Quickly verify current standardization status:

```
marked-bench --check-standard-status  
marked-bench --check-standard-status --json
```

SUBJECT: MB-0001

CLASS: PUBLIC INFRASTRUCTURE

ROLE: SILENT WITNESS

EVIDENCE
MB-0001
DO NOT REMOVE

1

1.72 m

PERIMETER

DO NOT CROSS

2.01 m

2

3

SEELED
EVIDENCE



The Marked Bench

Scoring Compatibility / Scoring Spec /
Adoption / Implementation

Martin123132 / The-Marked-Bench-

EXHIBIT A-004

PRIMARY SUBJECT

BENCH ID: MB-0001

WITNESS NOTE #17

It's always been here."

- morning walker
06:12



WITNESS NOTE #42

No one sits for long.
Feels like it remembers.

- Anonymous
18:09

OBSERVATION LOG

- ☒ Unattended
- ☒ Weathered
- ☒ Public
- ☒ Persistent

STATUS: UNDER REVIEW

EVIDENCE TAG

MB-0001

TYPE: PARK BENCH

LOCATION: RIVERSIDE PARK

COLLECTED: 2025-05-04

COLLECTED BY: PB-STD UNIT

PBSTD-250504-0001



Export and validate deterministic scoring compatibility vectors:

```
marked-bench --export-scoring-compatibility  
standard/marked_bench_scoring_  
compatibility_v0_4_8.json
```

```
marked-bench --validate-scoring-compatibility  
standard/marked_bench_scoring_  
compatibility_v0_4_8.json
```

Export and validate the language-neutral scoring specification:

```
marked-bench --export-scoring-spec  
standard/marked_bench_scoring_spec_v0_4_8.json
```

```
marked-bench --validate-scoring-spec  
standard/marked_bench_scoring_spec_v0_4_8.json
```

```
marked-bench --export-scoring-spec-doc  
docs/SCORING_SPEC.md
```

Export and validate the machine-readable adoption packet for external users:

```
marked-bench --export-adoption-packet  
adoption/marked_bench_adoption_  
packet_v0_4_8.json
```

```
marked-bench --validate-adoption-packet  
adoption/marked_bench_adoption_  
packet_v0_4_8.json
```

```
marked-bench --validate-evidence-ledger  
adoption/third_party_evidence_  
ledger_v0_4_8.json
```

Export and validate the external implementation kit:

```
marked-bench --export-implementation-kit  
adoption/marked_bench_implementation_kit_v0_4_8.json
```

```
marked-bench --validate-implementation-kit  
adoption/marked_bench_implementation_kit_v0_4_8.json
```

FILED BY: PB-STD UNIT

DATE: 2025-05-04

FOR OFFICIAL USE ONLY



CHAIN OF CUSTODY			
☒ COLLECTED	PB-STD UNIT	2025-05-04	
☒ RECEIVED	FORENSIC OPS	2025-05-04	
☒ REGISTERED	DATA VAULT	2025-05-04	

HANDLING INSTRUCTIONS

Handle with care.
Do not alter.
Document everything.



The Marked Bench

Score External Systems — Part I

Martin123132 / The-Marked-Bench-

Score External Systems

Systems written in any language can submit predictions without importing this package. Export a template, fill the `predicted` labels, and score it into a full benchmark report:

```
marked-bench --suite contradiction-adversarial --export-prediction-template artifacts/predictions.jsonl
```

```
marked-bench --suite contradiction-adversarial --score-predictions artifacts/predictions.jsonl --system-name "my-system" --report artifacts/my-system-report.json
```

Prediction files may be JSONL or JSON. Each record needs `case\_id` and `predicted`; optional `detector\_score`, `detector\_note`, `rationale`, and `evidence` fields are preserved in the final report. `detector\_score` is interpreted as binary contradiction confidence on [0, 1] for calibration metrics. `rationale` and `evidence` feed the report's explanation audit so reviewers can see whether a score is backed by inspectable reasoning evidence.

Create and validate leaderboard submission metadata

```
marked-bench --create-submission artifacts/my-system-submission.json --submission-report artifacts/my-system-report.json --system-version "1.0.0" --submitter "name-or-org"
```

```
marked-bench --validate-submission artifacts/my-system-submission.json
```

Create a standardized review rubric for a validated submission bundle:

```
marked-bench --create-submission-review artifacts/my-system-review.json --review-bundle artifacts/my-system-submission-bundle.json --reviewer reviewer-name
```

```
marked-bench --validate-submission-review artifacts/my-system-review.json
```

Create and validate a standard result card for publication or citation:

```
marked-bench --create-result-card artifacts/my-system-result-card.json --result-report artifacts/my-system-report.json --result-bundle artifacts/my-system-submission-bundle.json --result-review artifacts/my-system-review.json
```

```
marked-bench --validate-result-card artifacts/my-system-result-card.json
```

CHAIN OF CUSTODY

☒ COLLECTED	SE-TEP UNIT	2025-05-04
☒ RECEIVED	FORENSIC OPS	2025-05-04
☒ REGISTERED	DATA VAULT	2025-05-04

HANDLING INSTRUCTIONS


Handle with care.
Do not alter.
Document everything.



WITNESS NOTE #05

The bench connects what we build with what others will test.

— anonymous
04:21





The Marked Bench

Score External Systems — Part II / Build Leaderboards

Martin123132 / The-Marked-Bench-

EXHIBIT A-006

PRIMARY SUBJECT

BENCH ID: MB-0001

WITNESS NOTE #67

Benchmarks
tell the truth
sooner or
later.

— night archivist
21:13

WITNESS NOTE #73

Score
everything.
Explain
nothing.

— anonymous
02:31



SURVEILLANCE STILL_06
2025-05-04 21:43:58

OBSERVATION LOG

- ☒ Reproducible
- ☒ Comparable
- ☒ Transparent
- ☐ Citeable

STATUS: UNDER REVIEW

EVIDENCE TAG

MB-0001

TYPE: PUBLIC BENCH

LOCATION: RIVERSIDE PARK

COLLECTED: 2025-05-04

COLLECTED BY: PB-STD UNIT



PBSTD-250504-0001

FIELD SKETCH



SUBJECT: MB-0001

CLASS: PUBLIC INFRASTRUCTURE

ROLE: SILENT WITNESS

EVIDENCE
MB-0001
DO NOT REMOVE

SEALED
PB-STD UNIT

1.72 m

2.01 m

PERIMETER
DO NOT CROSS

EVIDENCE LINE — DO NOT CROSS

EVIDENCE LINE — DO NOT CROSS

SURVEILLANCE SEQUENCE (STILLS)



20:11:02
CAM_07



21:47:31
CAM_07



00:13:58
CAM_07



04:59:17
CAM_07

CHAIN OF CUSTODY

☒ COLLECTED	PB-STD UNIT	2025-05-04
☒ RECEIVED	FORENSIC OPS	2025-05-04
☒ REGISTERED	DATA VAULT	2025-05-04

HANDLING INSTRUCTIONS

Handle with care.
Do not alter.
Document everything.



Create a self-contained public publication
packet in one command:

```
marked-bench --create-publication-packet  
artifacts/my-system-publication-packet  
--publication-report artifacts/my-system-report.json  
--publication-predictions artifacts/predictions.jsonl  
--system-version "1.0.0"  
--submitter "name-or-org"
```

```
marked-bench --validate-publication-packet  
artifacts/my-system-publication-packet/  
publication_packet.json
```

Create a citeable result claim from that packet:

```
marked-bench --create-result-claim  
artifacts/my-system-publication-packet/  
result_claim.json  
--claim-publication-packet  
artifacts/my-system-publication-packet/  
publication_packet.json
```

```
marked-bench --validate-result-claim  
artifacts/my-system-publication-packet/  
result_claim.json
```

Create a complete external-submission example:

```
python -m marked_bench.examples.external_submission_demo
```

```
marked-bench --validate-submission-bundle  
artifacts/external_submission_demo/  
example_external_submission_bundle.json
```

```
marked-bench --validate-submission-review  
artifacts/external_submission_demo/  
example_external_submission_review.json
```

A checked copy of that workflow is committed under
'submissions/example_external_jsonl/' so adopters can
inspect a full prediction, report, submission, bundle,
review, and result-card packet without generating one
first. A checked one-command publication packet is
committed under 'submissions/example_publication_packet/'.

Build Leaderboards

Active foundation leaderboard:

```
marked-bench --build-leaderboard baselines/always_none_v0_1_1.json  
baselines/hash_prior_v0_1_1.json baselines/contradiction_engine_v0_1_1.json  
--leaderboard-output leaderboard/leaderboard_v0_1_1.json
```

Legacy foundation leaderboard:

```
marked-bench --build-leaderboard baselines/always_none_v0_1_0.json  
baselines/contradiction_engine_v0_1_0.json  
--leaderboard-output leaderboard/leaderboard_v0_1_0.json
```

Adversarial leaderboard:

```
marked-bench --build-leaderboard  
baselines/always_none_adversarial_v0_2_0.json  
baselines/hash_prior_adversarial_v0_2_0.json  
baselines/contradiction_engine_adversarial_v0_2_0.json  
--leaderboard-output leaderboard/leaderboard_adversarial_v0_2_0.json
```

Multi-hop leaderboard:

```
marked-bench --build-leaderboard baselines/always_none_multihop_v0_3_0.json  
baselines/hash_prior_multihop_v0_3_0.json  
baselines/contradiction_engine_multihop_v0_3_0.json  
--leaderboard-output leaderboard/leaderboard_multihop_v0_3_0.json
```

Controls leaderboard:

```
marked-bench --build-leaderboard baselines/always_none_controls_v0_4_0.json  
baselines/hash_prior_controls_v0_4_0.json  
baselines/contradiction_engine_controls_v0_4_0.json  
--leaderboard-output leaderboard/leaderboard_controls_v0_4_0.json
```




The Marked Bench

Checked-In Evidence - Part I

Martin123132 / The-Marked-Bench-

EXHIBIT A-007

PRIMARY SUBJECT

BENCH ID: MB-0001

WITNESS NOTE #17

It's always
been here.

- morning walker
06:12

WITNESS NOTE #42

No one sits
for long.
Feels like it
remembers.

- Anonymous
18:09



SURVEILLANCE STILL_07
2025-05-04 21:47:31

OBSERVATION LOG

- ☒ Unattended
- ☒ Weathered
- ☒ Public
- ☒ Persistent

STATUS: UNDER REVIEW

EVIDENCE TAG

MB-0001

TYPE: PARK BENCH
LOCATION: RIVERSIDE PARK
COLLECTED: 2025-05-04
COLLECTED BY: PB-STD UNIT



PBSTD-2S0504-0001

FIELD SKETCH



SUBJECT: MB-0001

CLASS: PUBLIC INFRASTRUCTURE

ROLE: SILENT WITNESS

EVIDENCE
MB-0001
DO NOT REMOVE

1.72 m

1

PERIMETER

DO NOT CROSS

2

2.01 m

3

EVIDENCE LINE - DO NOT CROSS

EVIDENCE LINE - DO NOT CROSS

SURVEILLANCE SEQUENCE (STILLS)



20:11:02
CAM_07



21:47:31
CAM_07



00:13:58
CAM_07



04:59:17
CAM_07

CHAIN OF CUSTODY

☒ COLLECTED	PB-STD UNIT	2025-05-04
☒ RECEIVED	FORENSIC OPS	2025-05-04
☒ REGISTERED	DATA VAULT	2025-05-04

HANDLING INSTRUCTIONS

Handle with care.
Do not alter.
Document everything.

Checked-In Evidence

- Benchmark registry: `benchmark\_registry.json`
- Release manifest: `releases/`
- Conformance report: `conformance/marked\_bench\_conformance\_v0\_4\_8.json`
- Standard profile: `standard/marked\_bench\_standard\_profile\_v0\_4\_8.json`
- Change-control profile: `standard/marked\_bench\_change\_control\_v0\_4\_8.json`
- Scoring compatibility profile: `standard/marked\_bench\_scoring\_compatibility\_v0\_4\_8.json`
- Scoring specification: `standard/marked\_bench\_scoring\_spec\_v0\_4\_8.json`
- Scoring specification document: `docs/SCORING-SPEC.md`
- Adoption packet: `adoption/`
`marked\_bench\_adoption\_packet\_v0\_4\_8.json`
- Third-party evidence ledger: `adoption/third\_party\_evidence\_ledger\_v0\_4\_8.json`
- Implementation kit: `adoption/`
`marked\_bench\_implementation\_kit\_v0\_4\_8.json`
- Implementation kit templates: `adoption/implementation\_kit/`
- Suite manifests and coverage profiles: `suites/`
- Baseline reports: `baselines/`
- Leaderboard snapshots: `leaderboard/`
- JSON schemas: `schemas/`
- Benchmark methodology: `docs/BENCHMARK-STANDARD.md`
- Benchmark card: `docs/BENCHMARK-CARD.md`
- Project status: `docs/PROJECT-STATUS.md`
- Baseline robustness: `docs/BASELINE-ROBUSTNESS.md`
- Maintainer handoff: `docs/MAINTAINER-HANDOFF.md`
- Technical note: `docs/TECHNICAL-NOTE.md`
- Submission guide: `docs/SUBMISSION-GUIDE.md`
- Submission bundle schema: `schemas/`
`submission\_bundle.schema.json`
- Submission review schema: `schemas/`
`submission\_review.schema.json`
- Result card schema: `schemas/`
`result\_card.schema.json`
- Publication packet schema: `schemas/`
`publication\_packet.schema.json`
- Result claim schema: `schemas/`
`result\_claim.schema.json`

FILED BY: PB-STD UNIT

DATE: 2025-05-04

FOR OFFICIAL USE ONLY



The Marked Bench

Checked-In Evidence - Part II

Martin123132 / The-Marked-Bench-

EXHIBIT A-008
PRIMARY SUBJECT
BENCH ID: MB-0001

WITNESS NOTE #17

It's always
been here."
— morning walker
06:12



SURVEILLANCE STILL_08
2025-05-04 22:07:13

WITNESS NOTE #42

No one sits
for long.
Feels like it
remembers.

— anonymous
18:09

OBSERVATION LOG

- ☒ Unattended
- ☒ Weathered
- ☒ Public
- ☒ Persistent

STATUS: UNDER REVIEW

EVIDENCE TAG

MB-0001

TYPE: PARK BENCH
LOCATION: RIVERSIDE PARK
COLLECTED: 2025-06-04
COLLECTED BY: PB-STD UNIT



PBSTD-250504-0001

FIELD SKETCH



SUBJECT: MB-0001

CLASS: PUBLIC INFRASTRUCTURE

ROLE: SILENT WITNESS

EVIDENCE
MB-0001
DO NOT REMOVE

SEALED
EVIDENCE
DO NOT REMOVE

1.72 m


PERIMETER
DO NOT CROSS

2.01 m

SURVEILLANCE SEQUENCE (STILLS)



20:11:02
CAM_07



21:47:31
CAM_07



00:13:58
CAM_07



04:59:17
CAM_07

CHAIN OF CUSTODY

- | | | |
|--|--------------|------------|
| ☒ COLLECTED | PB-STD UNIT | 2025-05-04 |
| ☒ RECEIVED | FORENSIC OPS | 2025-05-04 |
| ☒ REGISTERED | DATA VAULT | 2025-05-04 |

HANDLING INSTRUCTIONS

Handle with care.
Do not alter.
Document everything.



Checked-In Evidence

- Implementation kit schema:
schemas/implementation_kit.schema.json
- Standard profile schema:
schemas/standard_profile.schema.json
- Change-control schema:
schemas/change_control.schema.json
- Scoring compatibility schema:
schemas/scoring_compatibility.schema.json
- Scoring specification schema:
schemas/scoring_spec.schema.json
- Adoption packet schema:
schemas/adoption_packet.schema.json
- Third-party evidence ledger schema:
schemas/third_party_evidence_ledger.schema.json
- Checked external submission packet:
submissions/example_external_jsonl/
- Checked publication packet:
submissions/example_publication_packet/
- Adoption guide:
docs/ADOPTION_GUIDE.md
- Announcement package:
docs/ANNOUNCEMENT_PACKAGE.md
- Third-party evidence protocol:
docs/THIRD_PARTY_EVIDENCE.md
- Standard change-control protocol:
docs/CHANGE_CONTROL.md
- Standardization status:
docs/STANDARDIZATION_STATUS.md
- Submission review rubric:
docs/SUBMISSION_REVIEW_RUBRIC.md
- Release notes:
docs/RELEASE_NOTES_v0_2_0.md
- Current release notes:
docs/RELEASE_NOTES_v0_4_8.md

FILED BY: PB-STD UNIT

DATE: 2025-05-04

FOR OFFICIAL USE ONLY



The Marked Bench

Quality Gates / Package Layout / Closing Notes

Martin123132 / The-Marked-Bench-

EXHIBIT A-009

PRIMARY SUBJECT

BENCH ID: MB-0001

WITNESS NOTE #23

Still here.
Still silent.
Still watching.
04:59:17



OBSERVATION LOG

- ☒ Unattended
- ☒ Weathered
- ☒ Public
- ☒ Persistent

STATUS: UNDER REVIEW

EVIDENCE TAG

MB-0001

ITEM: PARK BENCH

LOCATION: RIVERSIDE PARK

COLLECTED: 2025-05-04

COLLECTED BY: PB-STD UNIT



PBSTD-250504-0001

FIELD SKETCH



WITNESS NOTE #57

Bench remains.
Claims change.
Evidence stays.

— anonymous
05:00

SUBJECT: MB-0001

CLASS: PUBLIC INFRASTRUCTURE

ROLE: SILENT WITNESS

SURVEILLANCE SEQUENCE (STILLS)



CHAIN OF CUSTODY

☒ COLLECTED	PB-STD UNIT	2025-05-04
☒ RECEIVED	FORENSIC OPS	2025-05-04
☒ REGISTERED	DATA VAULT	2025-05-04

HANDLING INSTRUCTIONS

Handle with care.
Do not alter.
Document everything.

Quality Gates

Run these before publishing or submitting results:

```
python -m unittest discover -s tests
python scripts/validate_benchmark_artifacts.py
python -m marked_bench.benchmark_cli --check-standard-status
python scripts/check_scoring_sanity.py
python scripts/check_case_quality.py
python scripts/check_baseline_robustness.py
python scripts/regenerate_release_artifacts.py --check
python scripts/check_review_workflow.py
python scripts/check_license_notice.py
```

The artifact validator checks that suite manifests match code, baseline reports pass validation, the benchmark registry is current, and leaderboard snapshots match their underlying reports. It also checks the release manifest against the current public artifact hashes and checks public JSON artifacts against their public schemas. The checked external submission packet is also validated end-to-end so its JSONL predictions, report, submission bundle, review file, and file hashes stay consistent.

Checked result cards are validated against their referenced reports, bundles, reviews, hashes, and standard publication claims. Checked publication packets are validated against their copied reports, submissions, bundles, reviews, result cards, and file hashes. Checked result claims are validated against publication packets so public score wording stays tied to exact hashes and explicit boundaries. The conformance report provides one machine-readable pass/fail artifact for the full release package.

The standard profile is validated so the benchmark's own standardization requirements stay explicit, evidence-backed, and current. The scoring compatibility profile is validated so independent implementations can prove they produce the same scores from the same prediction vectors. The scoring specification is validated so independent implementations have a language-neutral contract for labels, metrics, rounding, and calibration.

The adoption packet is validated so external handoff, announcement, and citation material stays pinned to the same release evidence. The third-party evidence ledger is validated so adoption claims stay separate from unverified interest or private anecdotes. The implementation kit is validated so external CI templates, result-claim snippets, and pinned release paths stay aligned with the current release.

Package Layout

```
marked_bench/
├── benchmark_adoption.py      # Adoption packet export/validation
├── benchmark_cli.py          # CLI runner
├── benchmark_evidence.py      # Third-party evidence ledger validation
├── benchmark_implementation.py # External implementation kit validation
├── benchmark_scoring_compatibility.py # Scoring compatibility vectors
├── benchmark_scoring_spec.py  # Language-neutral scoring spec
├── benchmark_standard_profile.py # Benchmark standard profile validation
├── benchmark_leaderboard.py   # Validated leaderboard builder
├── benchmark_publication.py   # One-command public result packets
├── benchmark_claim.py         # Citeable result claim validation
├── examples/
│   └── external_submission_demo.py # End-to-end external JSONL workflow
├── contradiction/
│   ├── benchmark_suite.py      # Versioned benchmark tracks
│   └── engine.py              # Symbolic baseline detector
```

This repository is intentionally benchmark-only. It does not include the wider research utilities from the original toolkit.

Contributing

Read 'CONTRIBUTING.md' and 'docs/SUBMISSION_GUIDE.md' before adding benchmark cases, reports, or leaderboard entries. Existing public case IDs should not be edited after publication; add new coverage through a new suite version or track.

License

This repository uses the PolyForm Noncommercial License 1.0.0. Personal and non-commercial use is permitted under that license. Commercial use requires a separate written commercial license; contact the Chief Operating Officer (COO) of TWO HANDS NETWORK LTD to discuss commercial licensing.

About

The Marked Bench: a versioned contradiction-detection benchmark for AI reasoning evaluation.

INVESTIGATION CONTINUES

EVIDENCE PRESERVED. TRUTH PENDING.



EVIDENCE LINE — DO NOT CROSS

EVIDENCE LINE — DO NOT CROSS