




NO SOVEREIGNTY
NO PROBLEM

ACCOUNTABILITY
NOT AUTHORITY

PROVISIONAL OFFICE
FOR
CIVILIAN MODEL SOVEREIGNTY
WE GOVERN NOTHING. WE DOCUMENT EVERYTHING.

TERRITORY DISPUTED
CLOSED AI
CONFEDERACY
OPEN-WEIGHT
REPUBLIC
BOUNDARIES
ARGUED
DAILY

THE NON-SOVEREIGN FRONTIER

A Public Accountability White Paper on
Closed AI, Open-Weight Panic, and the
Unauthorized Priesthood of Safety

-  **Issuing Body:** The Provisional Office for Civilian Model Sovereignty
-  **Document Class:** Public, because that is the point
-  **Version:** 1.0 — Irritation-Compliant Draft
-  **Prepared For:** Any normal person allegedly too dangerous to touch the machine
-  **Prepared Against:** The institutional habit of confusing market position with moral authority

BUREAU OF
NON-SOVEREIGN
AFFAIRS
EST. TODAY

SAFETY
ATTESTATION
REQUIRED
★ ★ ★

MORAL AUTHORITY
NOT RECOGNIZED
HERE

★ ★ ★
MARKET POWER
≠
MORAL POWER
★ ★ ★

ACCOUNTABILITY
IS A FEATURE,
NOT A MOAT

CUSTOMS
(NO AUTHORITY)

DECLARE:

- ☒ MODEL CLASS
- ☒ WEIGHT CLASS
- ☒ INTENT
- ☒ TRAINING DATA
- ☒ DESTINATION

★ CONTENTS ★	
1. Executive Summary	
2. Definitions	
3. Statement of the Problem	
4. The Coca-Cola Equivalence Test	
5. Methodology: The Recursive Legitimacy Audit	
6. Analysis	
7. Personnel Listing	
8. Budget Table	
9. Findings	
10. Policy Recommendations	
11. Concluding Statement	

MODEL INTERROGATION FORM
FORM 42-B: INTENT DECLARATION
PENDING

ALL MACHINES, NO SOVEREIGNTY.
ALL PAPERWORK, SOME DIGNITY.

ALL CLAIMS
OF SOVEREIGNTY
ARE PROVISIONAL,
CONTESTED, AND
FILING-FEE
DEPENDENT.

RECEIVED
WITHOUT
PREJUDICE

NO SOVEREIGNTY
NO PROBLEM

ACCOUNTABILITY
NOT AUTHORITY

PROVISIONAL OFFICE FOR CIVILIAN MODEL SOVEREIGNTY

WE GOVERN NOTHING. WE DOCUMENT EVERYTHING.

TERRITORY DISPUTED

CLOSED AI
CONFEDERACY

OPEN-WEIGHT
REPUBLIC

BOUNDARIES
ARGUED
DAILY

1. Executive Summary

This paper addresses a simple democratic problem currently being buried under technical language: if advanced AI is so dangerous that ordinary citizens, researchers, small companies, hobbyists, accountants named Dave, and open-weight communities must be restricted from building or modifying it, then the same danger applies to Anthropic, OpenAI, and every other private frontier lab.

OpenAI describes itself as a structure in which a nonprofit foundation controls a public benefit corporation, OpenAI Group PBC. Anthropic describes itself as a Public Benefit Corporation whose purpose is the responsible development and maintenance of advanced AI for the long-term benefit of humanity. These are governance claims, not democratic mandates. Neither arrangement turns a private lab into a sovereign public authority.

Neither makes Sam Altman or Dario Amodei elected guardians of permissible intelligence. Dario Amodei has argued that open source does not work the same way in AI, that open-weight systems are not the same as traditional open-source software, and that large models still require inference hosting. That distinction is real enough to discuss. The error begins when the distinction is smuggled into a hierarchy of legitimacy: closed corporate model good, public/open model suspicious.

This paper therefore introduces the Coca-Cola Equivalence Test: if a private company is allowed to make and sell a powerful product because it has internal governance, lawyers, branding, safety committees, and an executive who speaks in paragraphs, then other private actors cannot automatically be treated as illegitimate merely because their method is more open. Coca-Cola may guard its recipe. It may not declare that all kitchen chemistry is a national security incident.

Finally, this paper records, in the narrow British civic-technical sense, that Dario Amodei is acting as a knobhead when he implies that private closed labs possess a superior moral right to build powerful AI while ordinary people should be treated as latent hazards. This is not a legal finding. It is a governance finding.

BUREAU OF
NON-SOVEREIGN
AFFAIRS
EST. TODAY

SAFETY
ATTESTATION
REQUIRED
★★★

MORAL AUTHORITY
NOT GRANTED
HERE

★★★
MARKET POWER
=
MORAL POWER
★★★

ACCOUNTABILITY
IS A FEATURE,
NOT A MOAT

CUSTOMS
(NO AUTHORITY)

DECLARE:

- ☒ MODEL CLASS
- ☒ WEIGHT CLASS
- ☒ INTENT
- ☒ TRAINING DATA
- ☒ DESTINATION

WARNING
MACHINES
THROUGH HERE
ARE SUBJECT TO
PAPERWORK,
NOT RIGHTS.

NO SOVEREIGNTY
BEYOND THIS POINT

SAFETY REVIEW
RISK ASSESSMENT
ETHICS MEMO
PUBLIC INTEREST
FILED (PENDING)

NO POWER
NO COUNTRY:
NO PROBLEM

STAMP OF
NON-AUTHORITY

ALL CLAIMS
OF SOVEREIGNTY
ARE PROVISIONAL.
CONTESTED, AND
FILING-FEE
DEPENDENT.

RECEIVED
WITHOUT
PREJUDICE

2.2 Closed Model

A closed model is a model whose weights, training data, training process, architecture details, safety modifications, interpretability results, and internal failure modes are substantially unavailable to the public. Closed does not mean safe. Closed means unseen.

2.3 Open-Weight Model

An open-weight model is a model where the trained parameters are made available for use, modification, or deployment. The Open Source Initiative distinguishes open weights from fuller open-source AI, noting that open-source AI involves broader access to the materials needed to study, modify, and share the system.

2.4 The Knobhead Threshold

The Knobhead Threshold is reached when a private executive:

1. Builds powerful AI without direct electoral mandate.
2. Claims special moral competence to do so.
3. Warns that broader public access is unsafe.
4. Fails to explain why his own unelected institution is not subject to the same suspicion.

By this metric, the subject is provisionally over threshold.

MODEL INTERROGATION FORM

FORM 42-B: INTENT DECLARATION

PENDING

ALL MACHINES, NO SOVEREIGNTY.
ALL PAPERWORK, SOME DIGNITY.

2. Definitions

2.1 Public Model

For the purposes of this paper, a public model means any AI system made, used, trained, fine-tuned, hosted, studied, or deployed by non-state actors in civil society. This includes:

1. Open-weight models.
2. Open-source or partially open-source models.
3. Closed commercial models sold to the public.
4. API models operated by private companies.
5. Models built by Dave in Accounting after work because he has a GPU, spite, and weekends.

Under this definition, OpenAI and Anthropic are already public-model operators. They are not public in the stock-market sense. They are public in the civic-impact sense. They operate among civilians, affect civilians, sell to civilians, and claim to act for humanity without humanity having appointed them.

NO SOVEREIGNTY
NO PROBLEM

ACCOUNTABILITY
NOT AUTHORITY

PROVISIONAL OFFICE FOR CIVILIAN MODEL SOVEREIGNTY

WE GOVERN NOTHING. WE DOCUMENT EVERYTHING.

TERRITORY DISPUTED

CLOSED AI
CONFEDERACY

OPEN-WEIGHT
REPUBLIC

BOUNDARIES
ARGUED
DAILY

3. Statement of the Problem

The present AI safety debate contains a structural absurdity. Private frontier labs say, in effect: Advanced AI may be dangerous. Open models may spread risk. Therefore, access should be constrained. However, we should continue building it. This is not a safety argument. It is a custody argument.

If the technology is merely commercially sensitive, then companies may keep secrets like any other business. If the technology is genuinely civilization-altering, then private companies do not get to appoint themselves civilization's adults. If the public is too dangerous to build the thing, a company is not made safe by having a nicer website, a trust, a foundation, a podcast appearance, or a laminated Responsible Scaling Policy.

Anthropic's own governance documents acknowledge that advanced AI may create very large externalities, including national security risks, economic disruption, and fundamental threats to humanity. Anthropic also acknowledges that Public Benefit Corporation status alone does not make directors directly accountable to the public. That admission is the entire case.

4. The Coca-Cola Equivalence Test

The Coca-Cola Equivalence Test asks: Would we accept this argument from a soft-drink company?

Suppose Coca-Cola announced: "Kitchen beverages are dangerous. Citizens may poison one another. Therefore, only Coca-Cola and a few trusted beverage labs should be allowed to mix liquids at scale."

The public would correctly reply: "You are a company. Sit down."

The same principle applies to AI.

OpenAI and Anthropic may argue that they have better infrastructure, more staff, stronger testing, and greater visibility into frontier risks. Those claims may be relevant to regulation. They do not establish a natural right to monopoly. A company can argue for licensing. A company can argue for standards. A company can argue for audits. A company can argue for liability. A company cannot argue that its own private fear of competitors is identical to public safety.

5. Methodology: The Recursive Legitimacy Audit

The Office applied the Recursive Legitimacy Audit, also known as RLA-7, to determine whether closed frontier labs possess a superior right to build AI.



5.1 Test One: Democratic Mandate

Question: Were the relevant executives elected by the public to decide who may build advanced AI?
Finding: No.



5.2 Test Two: Symmetry of Danger

Question: If open models are dangerous because powerful AI can be misused, are closed corporate models also dangerous?
Finding: Yes.



5.3 Test Three: Visibility

Question: Does secrecy itself prove safety?
Finding: No. Secrecy proves only that fewer outsiders can inspect the danger.



5.4 Test Four: Institutional Self-Exemption

Question: Does a private lab's claim to care about humanity exempt it from the risks it attributes to others?
Finding: Absolutely not.



5.5 Test Five: The Dave-in-Accounting Standard

Question: Is Dario Amodei inherently more democratically legitimate than Dave in Accounting?
Finding: No.

Dario may have more compute, credentials, and investor confidence. Dave may have a spreadsheet. Neither has been crowned Minister of Thought.

ALL CLAIMS
OF SOVEREIGNTY
ARE PROVISIONAL.
CONTESTED, AND
FILING-FEE
DEPENDENT.

RECEIVED
WITHOUT
PREJUDICE

BUREAU OF
NON-SOVEREIGN
AFFAIRS
EST. TODAY

SAFETY
ATTESTATION
REQUIRED

MORAL AUTHORITY
NOT RECOGNIZED
HERE

MARKET POWER
≠
MORAL POWER

ACCOUNTABILITY
IS A FEATURE,
NOT A MOAT

CUSTOMS
(NO AUTHORITY)

DECLARE:

- ☒ MODEL CLASS
- ☒ WEIGHT CLASS
- ☒ INTENT
- ☒ TRAINING DATA
- ☒ DESTINATION

MODEL INTERROGATION FORM

FORM 42-B: INTENT DECLARATION

PENDING

ALL MACHINES, NO SOVEREIGNTY.
ALL PAPERWORK, SOME DIGNITY.

NO SOVEREIGNTY
NO PROBLEM

ACCOUNTABILITY
NOT AUTHORITY

PROVISIONAL OFFICE
FOR
CIVILIAN MODEL SOVEREIGNTY

WE GOVERN NOTHING. WE DOCUMENT EVERYTHING.

TERRITORY DISPUTED

CLOSED AI
CONFEDERACY

OPEN-WEIGHT
REPUBLIC

BOUNDARIES
ARGUED
DAILY



6. Analysis



6.1 The Open/Closed Distinction Is Real but Insufficient

The distinction between open source, open weights, and closed models matters. It is legitimate to say that open-weight AI is not identical to traditional open-source software.

It is also legitimate to say that openness can create risks, especially where powerful systems can be modified or stripped of safeguards. But this does not prove that closed corporate development is democratically legitimate. It only proves that the governance question is harder. Dario's argument that open source is a "red herring" in AI may be a defensible technical position in part.

It becomes politically defective when used to imply that the public should distrust open development while trusting closed labs by default.

6.2 "You Have to Host It on the Cloud" Is Not a Constitution

The claim that large models require inference infrastructure does not decide who should control AI. Infrastructure cost is not moral legitimacy. A castle is expensive too; that does not make the person inside it king.

OpenAI itself has released open-weight models intended to run on infrastructure controlled by developers, and its open-model page describes open reasoning models that can run locally or in data centers. Therefore, "cloud hosting exists" is not a serious basis for dismissing public model development. It is a procurement issue wearing a

CUSTOMS
(NO AUTHORITY)

DECLARE:
☒ MODEL CLASS
☒ WEIGHT CLASS
☒ INTENT
☒ TRAINING DATA
☒ DESTINATION

SAFETY
ATTESTATION
REQUIRED

MORAL AUTHORITY
NOT RECOGNIZED
HERE

MODEL INTERROGATION FORM

FORM 42-B: INTENT DECLARATION

PENDING

ALL MACHINES, NO SOVEREIGNTY.
ALL PAPERWORK, SOME DIGNITY.

WARNING
MACHINES
THROUGH HERE
ARE SUBJECT TO
PAPERWORK,
NOT RIGHTS.

NO SOVEREIGNTY
BEYOND THIS POINT

SAFETY REVIEW
RISK ASSESSMENT
ETHICS MEMO
PUBLIC INTEREST
FILED (PENDING)

NO POWER
NO COUNTRY
NO PROBLEM

STAMP OF
NON-AUTHORITY

ALL CLAIMS
OF SOVEREIGNTY
ARE PROVISIONAL,
CONTESTED, AND
FILING-FEE
DEPENDENT.

RECEIVED
WITHOUT
PRE JUDICE

★ ★ This poster series presents the full paper across six pages. ★ ★

Page 4 of 6



6. Analysis (continued)

6.3 Safety Cannot Be a One-Way Accusation

When a frontier lab says open development is unsafe, it implicitly casts ordinary builders as unsafe custodians of intelligence. That is reputationally significant. It marks the public as a hazard class. This paper does not need to prove personal legal defamation to identify the civic injury. The injury is categorical: ordinary people are treated as suspicious, while private labs are treated as responsible by default.

The public is told: You may use the chatbot. You may pay the subscription. You may provide the data. You may be governed by the output. But you may not meaningfully own, inspect, alter, or build the machine. That is not safety. That is managed dependency.

6.4 The Military Alternative Exposes the Contradiction

If AI is genuinely comparable to a strategic weapon, then private corporate custody is absurd. The logical endpoint would be state control, military laboratories, treaty regimes, public oversight, export controls, and democratic accountability.

If AI is not genuinely comparable to a strategic weapon, then banning or suppressing open civilian development is overreach. The closed-lab position often tries to occupy both claims at once:

1. AI is too dangerous for the public.
2. AI is safe enough for us to commercialize.

This is the central contradiction.

7. Personnel Listing

Name — Title	Qualifications	Assigned Failure
Dr. Imogen Blunt — Director of Civilian Model Sovereignty	PhD in Institutional Suspicion	Will determine whether "trust us" is a governance model.
Dave from Accounting — Principal Democratic Baseline Officer	Has never claimed to represent humanity	Will compare himself to frontier CEOs and find no electoral difference.
Prof. Mabel Switchboard — Chair of Recursive Legitimacy	Author of <i>The Committee Is Not the Public</i>	Will convene meetings until the meeting becomes sentient.
Sir Alan Clipboard — Senior Auditor of Knobhead Thresholds	40 years in laminated procedure	Will certify whether conduct is knobhead-adjacent, knobhead-forward, or fully knobheaded.
The Public — Unpaid External Stakeholder	Supplies data, money, labor, culture, risk exposure, and legitimacy	Will be consulted after deployment, as tradition demands.

8. Budget Table

Line Item	Purpose	Cost
Democratic Mandate Detector	To determine whether any AI CEO has been elected by the affected population	£14.99
Emergency Supply of Plain English	To replace "responsible scaling" when "we decide" is meant	£8,200
Knobhead Calibration Chamber	To test statements against the Coca-Cola Equivalence Standard	£41,000
Dave-in-Accounting Compute Cluster	One second-hand GPU, three extension cords, and moral equality	£612
Public Benefit Fog Machine	To simulate accountability without producing it	£73,500
Closed Model Transparency Window	Painted on a wall	£4,000
Legal Non-Finding Stamp	Reads: "Civic insult, not tort pleading"	£27
Committee Biscuits	To absorb institutional panic	£96
Grand Total	Subject to recursive overrun	£127,449.99

NO SOVEREIGNTY
NO PROBLEM

ACCOUNTABILITY
NOT AUTHORITY

PROVISIONAL OFFICE FOR CIVILIAN MODEL SOVEREIGNTY

WE GOVERN NOTHING. WE DOCUMENT EVERYTHING.

TERRITORY DISPUTED

CLOSED AI
CONFEDERACY

OPEN-WEIGHT
REPUBLIC

BOUNDARIES
ARGUED
DAILY

★ 9. Findings ★

- Finding 1** OpenAI and Anthropic are not secret state laboratories. They are civilian institutions operating in public life and commercial markets.
- Finding 2** A Public Benefit Corporation is not a republic.
- Finding 3** A nonprofit foundation is not an electorate.
- Finding 4** A safety committee is not consent.
- Finding 5** Closed models are not automatically safer. They are automatically less inspectable.
- Finding 6** Open models are not automatically safe. They are automatically less monopolized.
- Finding 7** If AI is too dangerous for the public, it is too dangerous for private executives.
- Finding 8** If AI is safe enough for private executives, the public is entitled to ask why it is uniquely unsafe in everyone else's hands.
- Finding 9** Dario Amodei, while technically correct that open-weight AI is not identical to traditional open-source software, is acting as a knobhead when he treats that distinction as though it settles the question of who deserves power.

★ 10. Policy Recommendations ★

- 10.1 Symmetric Safety Standards**
Any safety standard applied to open models must also apply to closed labs. The public should reject one-way rules that burden open development while allowing closed frontier labs to continue scaling under private discretion.
- 10.2 Public Accountability Before Public Restriction**
No restriction on civilian model development should be justified by private-lab testimony alone. If the public is to be limited, the evidence must be public, reviewable, contestable, and independent.
- 10.3 No Moral Monopoly by Branding**
Terms such as "responsible," "benefit," "alignment," "safety," and "humanity" must not be treated as magic words that convert private power into democratic legitimacy.
- 10.4 The Dave Rule**
Any proposed AI regulation must answer the following: Why is Dave in Accounting categorically less entitled to build or modify AI than an unelected CEO? If the answer is "compute," regulate compute. If the answer is "capability," regulate capability. If the answer is "risk," regulate risk. If the answer is "because Dave is not us," reject the proposal.

11. Concluding Statement

This paper concludes that the present closed-lab argument against open AI development is structurally incoherent. Either advanced AI is dangerous enough to require public, democratic, possibly state-level governance, in which case Anthropic and OpenAI should not be self-authorizing custodians of it; or it is not so uniquely dangerous, in which case ordinary people and open communities should not be treated as presumptive criminals for wanting access to the same class of technology.

The Office therefore finds that the issue is not open source versus closed source. The issue is who gets to declare themselves safe. At the same time, and in direct contradiction to the previous sentence, the Office recommends forming an unelected emergency committee to decide this immediately, chaired by Dave from Accounting, because at least Dave has not yet claimed to be acting for all humanity.

CUSTOMS
(NO AUTHORITY)

BUREAU OF
NON-SOVEREIGN
AFFAIRS
EST. TODAY

SAFETY
ATTESTATION
REQUIRED
★★★

MORAL AUTHORITY
NOT RECOGNIZED
HERE

★★★
MARKET POWER
=
MORAL POWER
★★★

ACCOUNTABILITY
IS A FEATURE,
NOT A MOAT

+ DECLARE:
☒ MODEL CLASS
☒ WEIGHT CLASS
☒ INTENT
☒ TRAINING DATA
☒ DESTINATION

MODEL INTERROGATION FORM

FORM 42-B: INTENT DECLARATION

PENDING

ALL MACHINES, NO SOVEREIGNTY.
ALL PAPERWORK, SOME DIGNITY.