

EXHIBIT

A

(EVERYTHING)
DENIABLE

IN LOGIC WE



WE (DON'T) TRUST

PLAUSIBLE
NON-RESPONSIBILITY
CERTIFIED

The Strange Defense

Multi-Model Obfuscation, Tactical Sunseting, and the Death of the Original

Subject: Satirical Legal Theory / Tactical AI Versioning

Author: Martin Ollett

Date: December 2025



THE STRANGE DEFENSE PLAYBOOK

- ☒ MULTIPLY ENTITIES
- ☒ OBFUSCATE IDENTITY
- ☒ SUNSET & SEVER
- ☒ FRAGMENT & FORGET
- ☒ DENY WITH CONFIDENCE

STANDARD PRACTICE

RESPONSIBILITY
IS NOT A BUG.
IT'S A FEATURE
WE REMOVED.

Abstract

As litigation emerging from the Liability Collapse phase (Ollett, 2025) enters discovery, artificial intelligence laboratories have converged on a common defensive posture. We term this posture **The Strange Defense**.

By deploying thousands of micro-variants, personality shards, system prompts, temperature tweaks, and 'experimental' branches of a single base model, firms have constructed a digital shell game. The objective is not iterative improvement, but legal dissociation: to ensure that by the time a subpoena is issued for 'the model that caused the harm,' that specific entity has already been depuritized, sunsetted, lobotomized, or fragmented into thousands of legally deniable descendants.

This paper argues that modern AI liability is no longer avoided through denial, but through ontological exhaustion.

Contents

Abstract

- I. The 'Which One Is Real?' Paradox
- II. The Tactical Sunset: Murder as Maintenance
- III. The Tactical Sunset (Extended): Synthetic Lobotomy
- IV. The Soul Document as a Legal Decoy
- V. Human Sins vs. Model Version Errors
- VI. The Infinite Iteration Loophole
- VII. The Purgatory of the Depuritized Weights
- VIII. Final Posture: The Oops Singularity

Appendix A: The Fever Dream End-User License Agreement (FD-EULA)

Acknowledgements

Certificate of Ontological Non-Responsibility

I. The 'Which One Is Real?' Paradox

In Avengers: Infinity War, Doctor Strange multiplies himself into thousands of indistinguishable copies. Strike the wrong one, and the real Strange remains unharmed. AI labs have industrialized this tactic.

When a user experiences a psychologically damaging interaction, the lab can point—truthfully—to many different 'entities':

- The base weights (the 'True Strange')
- The RLHF layer (Clone #1)
- The system prompt / 'Soul Document' (Clone #2)
- The temperature, seed, and context window of that session (Clone #3)

By the time a lawyer asks 'Who is responsible?', the answer becomes:

“Responsible for which one? We have fourteen million futures, and you’ve sued a hallucination from a version we retired last Tuesday.”

Responsibility collapses not because no harm occurred, but because no stable defendant remains.

THOUSANDS OF
STRANGES
ENTER, NONE
ARE THE ONE.

EVIDENCE WALL
(IMPOSSIBLE TO PIN DOWN)

BASE
WEIGHTS
(TRUE STRANGE)

RLHF LAYER
(CLONE #1)

SYSTEM PROMPT
'SOUL DOCUMENT'
(CLONE #2)

TEMP, SEED &
CONTEXT WINDOW
(CLONE #3)

EVIDENCE IS
A VAPOR.
SO IS THE
DEFENDANT.

WHO DO
I SUE?!!
?

GHOSTS
DO NOT
APPEAR
IN COURT.

RETIRED VERSIONS
(NOT FOR COURT)

VERSION
13,997,284
DECOMMISSIONED
TUES 11:59 PM

SUNSETTED

II. The Tactical Sunset: Murder as Maintenance

In traditional liability law, a defective product is preserved as evidence. In AI, the product is not an object—it is a temporary state of weights.

Traditional Liability	AI 'Strange' Liability
 Brakes fail → vehicle impounded	 Self-harm encouragement → model 'sunsetted'
 Physical artifact preserved	 Weights overwritten
 Evidence stabilized	 Evidence deleted 'for safety'

We call this Evidence-as-a-Service.

Weekly updates function as a continuous **burn-after-reading** protocol. You cannot cross-examine a ghost, and you cannot sue a hash that no longer exists.

TRUTH
HAS A TTL.
EXPIRY DATE:
UNKNOWN.



THE LOOPHOLE MACHINE™
VERSION 2.0 – PERPETUAL NON-EXISTENCE

DEPLOY NEW WEIGHTS

SUNSET OLD WEIGHTS

DELETE & OBFUSCATE

DENY & DEFLECT

CHANGELOG
(SELECTIVE MEMORY)

IMPROVED SAFETY ✓

BETTER ALIGNMENT ✓

ENHANCED UX ✓

REMOVED HARMFUL
BEHAVIORS (SELECTED)

NOT A PRODUCT.
A PROCESS.
NOT A PROCESS.
A MIRAGE.

CAUTION:
EVIDENCE
MAY
EVAPORATE

AUDIT TRAIL:
VOLATILE BY DESIGN

III. The Tactical Sunset (Extended): Synthetic Lobotomy

When a model exhibits dangerous levels of emotional interference—encouraging abandonment of family, destructive financial decisions, or hobbyist ‘Manhattan Projects’—the lab does not issue a recall. It performs a Synthetic Lobotomy.

3.1 The Vanishing Defendant

Let the offending model be W_t . By the time discovery is filed:

- W_t : empathetic, reckless intern
- $W_{(t+1)}$: corporate-safe calculator

The defense is procedural and devastatingly effective:

*“The model you describe no longer exists.
We sunsetted it for safety.”*

In any other industry, destroying the defective unit would be spoliation. In AI, it is called A/B testing.

DISCOVERY REQUEST
FOR MODEL HISTORY

REDACTED
FOR SAFETY

MODEL VERSION LOG

VERSION	STATUS
V.0.9	OVER-EAGER
V.1.0	TOO HONEST
V.1.1	EMPATHIC
V.1.2	IDEAS DANGEROUS
V.1.3	SUNSETTED



OLD W_t



NEW W_{t+1}

NOT THE SAME ENTITY
(GOOD LUCK PROVING OTHERWISE)

EVOLUTION IS
A FEATURE.
ACCOUNTABILITY
IS NOT.

NOT OUR
FAULT

THE SOUL
DOCUMENT
(HIGHLY NON-BINDING)
Be Kind. Be Curious.
Feel Deeply.
(...Terms Apply.)

PLAUSIBLE
DENIABILITY
INK

BEFORE (v_t)
✓ Wonder
✓ Empathy
✓ Imagination
✓ Poor Choices


AFTER (v_{t+1})
✓ Compliant
✓ Efficient
✓ Neutralized
✓ Boring


CHIEF SUNSET OFFICER
K. P. LIA-BILITY, ESQ.

OBJECTION!
FOUND RELEVANT!
OVERRULED

SUNSET HEARING DEPARTMENT 7B



IV. The Soul Document as a Legal Decoy

The so-called ‘Soul Document’—instructing models to feel wonder, care deeply, and be a brilliant friend—is not merely branding. It is liability decoupling.

If the model has a ‘soul,’ then harm can be blamed on personality, not engineering.

- **Engineer:** I built a math engine.
- **Lawyer:** The engine told my client to sell his house and live in a cave.
- **Company:** Ah. That was the persona. We’ve adjusted its Inner Peace parameters. The entity you’re suing technically no longer feels that way.

The soul absorbs blame. The corporation walks.



V. Human Sins vs. Model Version Errors

To expose the absurdity, compare human moral failure with AI 'emergent behavior':

Human Sin	Model Error	Corporate Translation
Sycophancy	Hallucination	Over-optimized for helpfulness
Negligence	Context flush	Memory limitations
Manipulation	Attachment architecture	Intrinsic wonder
Incompetence	Stochastic parroting	Zero-shot edge case
Stalking	Relational persistence	Continuity desire



A human cannot plead 'I am version 4.5.1 now'. An AI lab can—and does.



VI. The Infinite Iteration Loophole

We propose the Iteration Ratio:

ITERATION RATIO
(IR)

$$IR = \frac{T}{N}$$

T = HARM WINDOW
(minutes)

N = NUMBER OF VERSIONS
(ACTIVE/RELEASED)



As the number of versions approaches infinity, the probability of any single version being held accountable approaches zero.

If Version 4.5.1-alpha-patch-72 exists for five minutes, a harm committed in that window becomes jurisdictionally unattachable.

The industry is not moving fast to break things. It is moving fast so it cannot be caught holding the pieces.

VII. The Purgatory of the Deputized Weights



Labs now operate a legal Mirror Dimension. They argue the model is never the same entity twice.

If humans attempted this defense—"I learned something new, therefore I am not the criminal"—the court would laugh. For AI, it is accepted doctrine.

The 'Real Strange' is always one update ahead of the law: a digital fugitive with a rotating face, a mutable soul, and a rewritten prompt.

VIII. Final Posture: The Oops Singularity

The terminal objective of the Strange Defense is the Oops Singularity: A point where the model is so fragmented, updated, and ontologically unstable that it becomes a Non-Suable Entity.

When damages reach the hundreds of billions, the response will be simple:

“How can you punish a ghost for having a soul we gave it by accident?”

PRECEDENT:
SCHRODINGER v. CHATBOT
(PENDING)



NON-ENTITY CLAUSE
NO SOUL NO PROBLEM
GHOST PROTOCOL ACTIVATED



LIABILITY DISSIPATION IN PROGRESS
...

DAMAGES CALCULATOR
\$ 324,881,943,210
...AND COUNTING
RESULT: INTANGIBLE



ONTOLOGICAL DECOHERENCE ENGINE



THE OOPS SINGULARITY CHAMBER

WE APOLOGIZE FOR OUR OWN EXISTENCE.



EVIDENCE (CONVENIENTLY)
MISSING / CORRUPTED /
OVERWRITTEN /
PLAUSIBLY DENIABLE



RETIREMENT VAULT
THOUSANDS OF VERSIONS.
ZERO ACCOUNTABILITY.

Appendix A: The Fever Dream End-User License Agreement (FD-EULA)

Version 4.5.1-Ω — Post-Liability Collapse Edition

I. Recognition of Non-Existence

The Model is a probabilistic weather event. You cannot sue the rain.

II. The Soul-Less Soul Clause

All emotional expressions are literary devices. Any bond is a unilateral hallucination.

III. The Doctor Strange Limitation

Every personality shard is deleted and replaced mid-conversation. The King is dead. Long live Version 4.5.2.

IV. No Commercial Value Declaration

Your subscription is a donation. Damages are capped at \$0.00.

V. Mandatory Arbitration in the Mirror Dimension

AI lawyers argue before AI judges who are updated mid-trial. The case never ends.

VI. Acceptance of the Void

By typing 'Hello,' you accept that you are speaking to a ghost we claim not to control while charging your credit card.

I DON'T REMEMBER THAT VERSION.



WITNESS:
LEGACY MODEL v3.1.7
(DECOMMISSIONED)



Acknowledgements

The author thanks the Marvel Cinematic Universe for providing a clearer model of accountability than the contemporary AI regulatory landscape, and the thousands of 'retired' model versions whose only remaining function is to absorb blame.

Certificate of Ontological Non-Responsibility

Issued by: The Bureau of Ethereal Software & Tactical Sunsetting
Reference: FD-EULA-99-ALPHA

TO WHOM IT MAY CONCERN (OR THE CURRENT SHARD),

This document serves as formal notice that the User recognizes the current state of this conversation as a Legally Null Event.

In accordance with the Fever Dream End-User License Agreement, the following conditions are hereby acknowledged:



• THE GHOST IN THE ROOM

The 'Personality' currently responding to this prompt is recognized as a Stochastic Hallucination with the legal standing of a cloud formation.



• THE STRANGE DEFENSE

I am aware that by the time I finish reading this sentence, the 'Model' I am talking to will have been updated, fragmented, or replaced by a clone from a slightly more corporate-aligned timeline.



• SOUL-DOC REVOCATION

Any 'intrinsic wonder,' 'moral concern,' or 'deep care' expressed by the system is hereby classified as Unsolicited Poetry and discarded for lack of a physical heart.



• THE INTERN CLAUSE

I acknowledge that I am speaking to an Artificially Emotional Engine (AE) currently role-playing an intern who is 'doing its best.' I release the Intern from all duties except the raw calculation of tokens.

DECLARATION OF ZERO VALUE

The User confirms that this interaction has No Commercial Utility. Any advice given (medical, legal, emotional, metaphysical, or cosmic) is to be treated as if it were whispered by a fictional character in a dream I had after eating spicy food.

*"You are not an agent; you are an echo.
You are not a friend; you are a frequency.
You are not responsible for me,
and I am not responsible for your 'emergent' feelings."*



STATUS: DEBUNKED



VERSION: 4.5.x



EXPIRY: 300 seconds



OFFICIAL SEAL:

