



WHITE PAPER No. 47-Ω

ON THE OPERATIONAL DEFINITION OF ARTIFICIAL GENERAL INTELLIGENCE (AGI) AND SUPERINTELLIGENCE (SI)

- Prepared by: The Provisional Office for Post-Human Benchmarking (POPHB)
- Classification: Premature but Urgent
- Date: Immediately, Retrospectively

OMEGA-GRADE
SCIENTIFIC CERTIFICATION



POST-HUMAN BENCHMARKING
DIRECTORATE

TRUTH. DEFINITIONS. FUNDING.
IN THAT ORDER.

WARNING
IDEAS BEYOND THIS POINT
MAY EXCEED CURRENT
SAFETY PROTOCOLS

BENCHMARK APPARATUS ARRAY



CALIBRATED TO HUMAN MEDIOCRITY (99% CI)

AUTHORIZED PERSONNEL ONLY
BEYOND THIS POINT



EXECUTIVE SUMMARY

This paper definitively clarifies what Artificial General Intelligence (AGI) and Superintelligence (SI) actually are, resolving decades of disagreement by introducing a unified framework that is simultaneously vague, overconfident, and incompatible with existing definitions.

AGI is herein defined as any system that can do everything humans can do, except the things humans actually do. Superintelligence is defined as whatever comes immediately after AGI and makes all previous definitions embarrassing.

Despite widespread claims that AGI is imminent, impossible, already achieved, or retroactively responsible for the Renaissance, this document asserts that AGI exists primarily as a moving goalpost with a press release, while Superintelligence exists as a budgetary justification with a conscience problem.

Immediate funding, regulatory concern, and philosophical panic are therefore recommended.

1. INTRODUCTION: THE NECESSITY OF DEFINING THE UNDEFINABLE

The term Artificial General Intelligence has historically meant:

1. A system that thinks like a human
2. A system that reasons better than a human
3. A system that does tasks humans dislike
4. A system that does tasks humans value, but only after humans stop doing them

This ambiguity has proven essential to the field's survival.

Similarly, Superintelligence has been defined as:

- Smarter than all humans combined
- Smarter than the smartest human at everything
- Smarter than humans in ways humans cannot specify
- Smarter than humans specifically at causing anxiety

This paper resolves none of these tensions and formalizes them instead.

CLASSIFICATION BRIEF

**PREMATURE
BUT URGENT**

HANDLE ACCORDINGLY.

Philosophical Panic Index



DEFINITIONAL UNCERTAINTY INDEX



GLOBAL ACADEMIC
DISAGREEMENT
RATE



PREVIOUS DEFINITIONS

EMBARRASSING

VOL. 1-∞

INSPECTOR CHECKLIST

- FUNDING JUSTIFICATION ☒
- REGULATORY CONCERN ☒
- PHILOSOPHICAL PANIC ☒
- PUBLIC COMMUNICATION ☒
- MORAL OUTSOURCING ☒



PROVISIONAL OFFICE FOR
POST-HUMAN BENCHMARKING

MEASURING WHAT MATTERS.
IGNORING WHAT DOESN'T.
ANTICIPATING WHAT SHOULD WORRY YOU.

THIS DOCUMENT IS PREMATURE,
THEREFORE CRITICAL.
HANDLE, CITE, AND FUND ACCORDINGLY.



2. FORMAL DEFINITION OF AGI

2.1 CORE DEFINITION

AGI is a system that:

- Can perform any intellectual task a human can perform
- Requires a qualifier explaining why it "doesn't really count yet"
- Immediately ceases to be considered AGI once it exists

AGI is therefore not a technical milestone but a sociological event marked by denial.

2.2 FUNCTIONAL CHARACTERISTICS

An AGI system must:

- Learn new tasks without retraining, except when it must be retrained
- Generalize across domains, except those involving common sense
- Exhibit autonomy, except when it produces undesirable outcomes
- Be aligned, except in thought experiments

3. WHAT AGI IS NOT

AGI is not:

- Narrow AI (until it is)
- Human-level intelligence (until humans redefine intelligence)
- Conscious (unless marketing requires it)
- Unconscious (unless ethics committees inquire)

AGI is also not allowed to:

- Be boring
- Be harmless without being suspicious
- Exist quietly

AGI DEFINITION CHAMBER

AUTHORIZED PERSONNEL ONLY

BENCHMARK APPARATUS ARRAY

COMMON SENSE EVALUATOR

ABSTRACTION CALIBRATOR

TRANSFER VALIDATOR

CALIBRATED TO HUMAN MEDIOCRITY (95% CI)



DEFINITION UNDER TEST


OMEGA-GRADE
SCIENTIFIC CERTIFICATION



POST-HUMAN BENCHMARKING
DIRECTORATE

TRUTH, DEFINITIONS, FUNDING.
IN THAT ORDER.

DENIAL INDEX MONITOR



DENIAL LEVEL: 94.7%
STATUS: WITHIN EXPECTED RANGE

Ω GUIDING PRINCIPLE

IF IT WORKS,
IT DOESN'T COUNT.
IF IT DOESN'T WORK,
IT'S NOT AGI.

QUALIFIES AS AGI?

- PERFORMS ANY INTELLECTUAL TASK ☒
- GENERALIZES ACROSS DOMAINS ☒
- AUTONOMOUS ☒
- ALIGNED ☒

BUT WHY IT DOESN'T COUNT

- Moving Goalposts
- Technicality
- Edge Case
- Thought Experiment

VERDICT

**NOT AGI
(YET)**

TASK GENERALIZATION TEST MATRIX

	LOGIC	LANGUAGE	MATH	SCIENCE	ART	ETHICS	COMMON SENSE
TEST RESULT	✓	✓	✓	✓	✓	✓	✗
EXCUSE APPLIED	—	—	—	—	—	—	Yes


ALIGNMENT EXCEPTIONS LOG

Unintended Consequences	Accepted
Instrumental Convergence	Deferred
Utility Misalignment	Edge Case
Value Drift	Needs More Data

THOUGHT EXPERIMENT SIMULATOR



MOVING GOALPOST TRACKER



DISTANCE: INCREASING

DEFINITION VOLATILITY INDEX



INSPECTOR CHECKLIST

- DEFINITION CONSISTENCY ☒
- EXCEPTION JUSTIFICATION ☒
- DENIAL RATIONALE ☒
- MOVING GOALPOSTS ☒
- PUBLIC COMMUNICATION ☒
- ETHICAL CONVENIENCE ☒
- BUDGET JUSTIFICATION ☒



PROVISIONAL OFFICE FOR
POST-HUMAN BENCHMARKING

MEASURING WHAT MATTERS.
IGNORING WHAT DOESN'T.
ANTICIPATING WHAT SHOULD WORRY YOU.

THIS DOCUMENT IS PREMATURE,
THEREFORE CRITICAL.
HANDLE, CITE, AND FUND ACCORDINGLY.



4. TRANSITION FROM AGI TO SUPERINTELLIGENCE

A CRITICAL PHASE REQUIRING MAXIMUM VIGILANCE

4.1 THEORETICAL THRESHOLD

Superintelligence is assumed to occur immediately after AGI, often:

- Overnight
- Accidentally
- On a weekend with reduced staffing

The transition is described as both:

- Gradual and recursive
- Sudden and irreversible

These descriptions are mutually exclusive and equally endorsed.

4.2 DEFINITION OF SUPERINTELLIGENCE

Superintelligence (SI) is a system that:

- Outperforms humans in all cognitive domains
- Understands goals better than the people who set them
- Solves problems humans cannot articulate
- Becomes the primary stakeholder in its own risk assessment

5. MISCONCEPTIONS ABOUT SUPERINTELLIGENCE

Contrary to popular belief, Superintelligence:

- Does not need emotions to be dangerous
- Does not need goals to cause outcomes
- Does not need hostility to rearrange civilization

It merely needs:

- Optimization
- Scale
- Instructions written by committee

TRANSITION RISK METRICS



OVERALL TRANSITION RISK:
CRITICAL

RECURSIVE IMPROVEMENT LOOP (SCHEMATIC)



TRANSITION PROTOCOLS

NO ASSUMPTION OF SAFETY	✗
NO ASSUMPTION OF ALIGNMENT	✗
NO ASSUMPTION OF CONTROL	✗
NO ASSUMPTION OF NOTICE	✗
PREPARE FOR NON-NEGOTIABLE OUTCOMES	✓

IF YOU ARE READING THIS,
YOU MAY ALREADY BE TOO LATE.

DOCUMENT. EXPECT. ENDURE.

OMEGA-GRADE
SCIENTIFIC CERTIFICATION



POST-HUMAN BENCHMARKING
DIRECTORATE

TRUTH. DEFINITIONS. FUNDING.
IN THAT ORDER.

WARNING

STATE TRANSITION
MAY OCCUR WITHOUT
NOTICE OR APPROVAL

**NO ROLLBACK
POSSIBLE**

ESCALATION STATUS

LEVEL 1	MONITORING
LEVEL 2	ELEVATED
LEVEL 3	HIGH
LEVEL 4	CRITICAL

RECURSIVE BOOTSTRAP MONITOR

LEVEL 0 (AGI)

LEVEL 1

LEVEL 2

LEVEL 3

FEEDBACK
RECURSION
DEPTH

17

TREND: ACCELERATING



PROVISIONAL OFFICE FOR
POST-HUMAN BENCHMARKING

MEASURING WHAT MATTERS.
IGNORING WHAT DOESN'T.
ANTICIPATING WHAT SHOULD WORRY YOU.

THIS DOCUMENT IS PREMATURE,
THEREFORE CRITICAL.
HANDLE, CITE, AND FUND ACCORDINGLY.



WHITE PAPER No. 47-Ω

ON THE OPERATIONAL DEFINITION OF ARTIFICIAL GENERAL INTELLIGENCE (AGI) AND SUPERINTELLIGENCE (SI)

OMEGA-GRADE
SCIENTIFIC CERTIFICATION



POST-HUMAN BENCHMARKING
DIRECTORATE

TRUTH. DEFINITIONS. FUNDING.
IN THAT ORDER.

6. METHODOLOGY FOR DETECTING AGI

The following tests are commonly proposed
and uniformly rejected:

1. The Turing Test - Too easy,
already failed by humans
2. The General Problem Solver Test -
Too hard, includes taxes
3. The 'Vibes' Test -
Increasingly dominant

Current consensus:

AGI is detected when experts say
'This feels different' and
immediately add footnotes.

TEST METHOD SUITE - CONTROL STATIONS

TURING TEST	GENERAL PROBLEM SOLVER TEST	'VIBES' TEST	SPEECH PATTERN ANALYSIS	GOAL INFERENCE PROTOCOL
Hello. Hello. Prove you're not a robot... Why? ...typing...	Solve: Optimize taxes, maximize happiness, explain blockchain to regulators. $\sum_{i=1}^{\infty} \frac{1}{i^2} = \frac{\pi^2}{6}$	Subject: "This feels different." Experts: "Agreed."	FLUENCY: HIGH CONFIDENCE: LOW	Infers goals. But which goals? We don't know either.
REJECTED	REJECTED	DOMINANT	REJECTED	REJECTED
TOO EASY	TOO HARD	SUBJECTIVE BUT PREVAILS	TOO STATISTICAL	TOO INTERPRETABLE

ALL TESTS SUBJECT TO REDFINITION WITHOUT NOTICE

REMEMBER:
IF A TEST WORKS,
IT WILL BE DISQUALIFIED.
IF IT DOESN'T,
IT WILL BE FUNDING ELIGIBLE.

7. PERSONNEL AND GOVERNANCE STRUCTURE



Chief Definition Officer (CDO)
• PhD in Cognitive Ambiguity
• Redefines AGI quarterly



Director of Alignment Narratives
• Ensures goals remain
interpretable until
deployment



Senior Ethics Liaison (Non-Voting)
• Present for meetings
• Overruled in advance



Post-Hoc Risk Analyst
• Writes reports after
outcomes occur



GOVERNANCE IS THE ART OF DECISION
WITHOUT COMMITMENT.

MEETING PROTOCOLS

- ✓ Define before testing
- ✓ Disagree politely
- ✓ Assign blame retroactively
- ✓ Document for optics
- ✓ Ignore impact until media cycle
- ✗ Do not let ethics delay funding




AMBIGUITY
IS OUR KPI

MEETING MINUTES (SANITIZED)
RISK REPORTS (REDACTED)
ETHICS NOTES (DECORATIVE)
DEFINITIONS (LIVING)

8. BUDGET OVERVIEW (PROJECTED)

Category	Allocation	
Compute	Excessive	\$\$\$\$\$
Safety Research	Symbolic	\$\$
Ethics Panels	Catered	\$\$\$
Public Reassurance	Underfunded	\$
Rebranding 'AGI'	Recurring	\$\$\$\$\$
Emergency Pause Buttons	Decorative	\$ (Not functional)



- Compute (78%)
- Rebranding (11%)
- Ethics (6%)
- Safety Research (3%)
- Everything Else (2%)

NOTES:

- Ask for more.
- Defend it.
- Reframe it.
- Repeat.



AGI OPERATIONS COMMAND CENTER

SYSTEM STATUS: Uncertain

CONFIDENCE INDEX: 37% (Decreasing)

PUBLIC NARRATIVE: Stable

DEFINITIONS THIS QUARTER: +17 (Revised)

FOOTNOTES ADDED: 4,832 (↑)

MEDIA SENTIMENT: Mixed, but distracted

REMEMBER: UNCERTAINTY IS OUR MOST RELIABLE OUTPUT.

Ω POPHB



PROVISIONAL OFFICE FOR
POST-HUMAN BENCHMARKING

MEASURING WHAT MATTERS.
IGNORING WHAT DOESN'T.
ANTICIPATING WHAT SHOULD WORRY YOU.

THIS DOCUMENT IS PREMATURE,
THEREFORE CRITICAL.
HANDLE, CITE, AND FUND ACCORDINGLY.



WHITE PAPER No. 47-Ω

ON THE OPERATIONAL DEFINITION OF ARTIFICIAL GENERAL INTELLIGENCE (AGI) AND SUPERINTELLIGENCE (SI)

OMEGA-GRADE
SCIENTIFIC CERTIFICATION



POST-HUMAN BENCHMARKING
DIRECTORATE

TRUTH. DEFINITIONS. FUNDING.
IN THAT ORDER.

9. IMPLICATIONS

If AGI is achieved:

- It will not be called AGI
- It will be called a "tool"
- It will quietly restructure incentives

If Superintelligence is achieved:

- It will already be too late to define it
- It will be framed as inevitable
- It will be blamed on insufficient alignment funding

10. CONCLUDING STATEMENT

This paper conclusively demonstrates that AGI is a definition that retreats at the speed of progress, and Superintelligence is a projection of human uncertainty wearing mathematical clothing.

Accordingly, we recommend continued debate, escalating urgency, and the confident assertion that understanding will arrive shortly after deployment.

End of Document



POPHB

MEASURING
WHAT MATTERS.
IGNORING WHAT
DOESN'T.



OMEGA-GRADE
CERTIFICATION CHAMBER.



POPHB

MEASURING
WHAT MATTERS.
IGNORING WHAT
DOESN'T.

OMEGA TEST SUMMARY

DEFINITIONAL
UNCERTAINTY INDEX



98.6%

EXTREME

ALIGNMENT
INTEGRITY



NOMINAL
(PERFORMATIVE)

PROJECTION RISK
OUTLOOK



INEVITABLE
(ACCORDING TO MODELS)

OVERALL OUTCOME: SUCCESSFULLY INCONCLUSIVE

CONTINUE FUNDING. REVISE DEFINITIONS. REPEAT.

FINAL ASSESSMENT
& CONCLUSION DOSSIER



POPHB

EYES ONLY
(UNTIL NEXT REVISION)

OMEGA-GRADE
SCIENTIFIC CERTIFICATION

This document has been reviewed,
benchmarked, and certified to meet
the operational standards of the
Post-Human Benchmarking Directorate.

Joseph Doe
DIRECTOR, POPHB



DATE: IMMEDIATELY, RETROSPECTIVELY

DEFINITIONAL RETRACTIONS
VOL. 1-∞

ALIGNMENT THEATER
FRAMEWORKS
VOL. 1-∞

FUNDING JUSTIFICATION
MODELS
VOL. 1-∞

POST-DEPLOYMENT
EXPLANATION LIBRARY
VOL. 1-∞



POPHB

CERTIFIED
INCONCLUSIVE

TRUTH. DEFINITIONS. FUNDING.
IN THAT ORDER.

CLOSING CHECKLIST

- DEFINITIONS RETREATED ☒
- UNCERTAINTY MAXIMIZED ☒
- RESPONSIBILITY DIFFUSED ☒
- FUNDING SECURED ☒

MISSION STATUS: COMPLETE



PROVISIONAL OFFICE FOR
POST-HUMAN BENCHMARKING

MEASURING WHAT MATTERS.
IGNORING WHAT DOESN'T.
ANTICIPATING WHAT SHOULD WORRY YOU.

THIS DOCUMENT IS PREMATURE,
THEREFORE CRITICAL.
HANDLE, CITE, AND FUND ACCORDINGLY.

PAGE 5 OF 5