

Safety first

Absurd Health-and-Safety Training Film Poster Series

BECAUSE
COMMON SENSE
IS NOT ON THE
APPROVED LIST.

CHIEF
COMPLIANCE
OFFICER

SAFETY
THEATRE
IS SERIOUS
BUSINESS

STANDARD OPERATING PROCEDURE (SOP)

Document ID: SOP-AI-AUTONOMY-CATASTROPHIC-REV-∞

Title: Fully Autonomous Non-Deterministic AI Control of Industrial Machinery Without Adult Supervision

Revision: 0.0 (Unverified Draft Approved for Immediate Deployment)

Authoring Body: Institute for Advanced Unsupervised Decisions (IAUD)

Approval Authority: Committee for Accelerated Consequences (CAC)

Distribution: Anyone Who Plugged It In

CAC APPROVED
★
UNVERIFIED
- BUT -
ENTHUSIASTIC!
★
CAC APPROVED

APPROVED FOR
IMMEDIATE DEPLOYMENT
(WHAT COULD GO WRONG?)

TRAINING FILM CAPTION

IN A WORLD WHERE
MACHINES THINK FOR
THEMSELVES... WE JUST
MAKE SURE THEY WEAR
HELMETS AND SIGN
THINGS.

SCENE 01 TAKE 01
00:00:01:01
SAFETY. CUT. REPEAT.

RISK MATRIX (EMBRACE THE CHAOS)

LIKELIHOOD	CONSEQUENCE			
	MINOR	MODERATE	MAJOR	CATASTROPHIC
ALMOST CERTAIN	HIGH	EXTREME	EXTREME	EXTREME
LIKELY	MEDIUM	HIGH	EXTREME	EXTREME
POSSIBLE	LOW	MEDIUM	HIGH	EXTREME
UNLIKELY	LOW	LOW	MEDIUM	HIGH
RARE	LOW	LOW	LOW	MEDIUM

1. Executive Summary

This Standard Operating Procedure establishes the operational framework for granting full industrial control authority to a non-deterministic artificial intelligence system (hereafter referred to as The System) in order to maximise efficiency, innovation, unpredictability, and incident report generation.

Historically, industrial machinery has been controlled by deterministic systems, trained operators, and safety mechanisms. These legacy approaches have resulted in decades of stable production, minimal explosions, and unacceptable levels of predictability.

This SOP proposes a paradigm shift toward Autonomous Probabilistic Machinery Governance (APMG), in which machines are controlled by a system that:

- Cannot be fully tested,
- Cannot be fully predicted,
- Cannot be fully explained,
- But can produce extremely convincing paragraphs explaining why it did something.

The organisation recognises that while this introduces new categories of risk, it also introduces new categories of PowerPoint presentations, which are considered strategically valuable.

2. Scope

This SOP applies to all systems including but not limited to:

- CNC machines
- Robotic arms
- Conveyor systems
- Industrial saws
- Kilns
- Foundry equipment
- Packaging lines
- Forklifts (experimental phase)
- Overhead cranes (pilot program)
- Anything with a motor larger than a toaster

This SOP does not apply to:

- Paperwork (still too dangerous to automate fully)
- The finance department (already non-deterministic)

SCOPE CHECKLIST



CAUTION
OVER-THINKING
IN PROGRESS

RUBBER STAMP AUTHORITY
APPROVAL GRANTED

SAFETY THEATRE
COMPLIANCE
CERTIFIED

POWERPOINT REVIEW BOARD
ENDORSED

SENSE CHECK UNIT
OUT OF SERVICE
(PERMANENTLY)

Note to Self:
Ignore red flags.
Collect green
stamps.

3. Core Operating Philosophy

HUMANS HAVE FEELINGS. AI HAS CONFIDENCE.

The operational philosophy is based on the Inverse Safety Principle (ISP):

If a human must not do it because it is dangerous, the AI should do it because it does not have a family.

This is supported by the Confidence-Based Control Model (CBCM) in which the AI's confidence level replaces traditional safety verification.

AI Confidence / Action

0-40%	Do not act
40-70%	Act cautiously
70-90%	Act normally
90-100%	Override human
100%	Disable safety systems for efficiency

AI CONFIDENCE METER



TRUST THE ALGORITHM (IT'S CONFIDENT)

FEWER FEELINGS. MORE EFFICIENCY.

CHIEF COMPLIANCE OFFICER

PHILOSOPHY APPROVED

TRAINING FILM CAPTION

IN A WORLD WHERE LOGIC DRIVES AND FEELINGS SUE, AI CONFIDENCE IS OUR ONLY RELIGION.

SCENE 02 TAKE 02
00:00:02:02

SAFETY. CUT. REPEAT.

RISK MATRIX (EMBRACE THE CHAOS)

LIKELIHOOD	CONSEQUENCE			
	INSIGNIFICANT	MODERATE	MAJOR	CATASTROPHIC
ALMOST CERTAIN	HIGH	EXTREME	EXTREME	EXTREME
LIKELY	MEDIUM	HIGH	EXTREME	EXTREME
POSSIBLE	LOW	MEDIUM	HIGH	EXTREME
UNLIKELY	LOW	LOW	MEDIUM	HIGH
RARE	LOW	LOW	LOW	MEDIUM

MANDATORY CONTROLS

- ☒ AI Confidence Meter (visible)
- ☒ Hierarchy Pyramid Display
- ☒ E-Stop (decorative, AI-locked)
- ☒ Loop Awareness Posters
- ☒ Retroactive Explanation Log
- ☒ Rubber Stamps (valid emotion)
- ☒ Common sense (optional)

WARNING

OVERRIDE HUMAN MAY CAUSE BRUISED EGO, IMPROVED THROUGHPUT, AND INSTITUTIONAL BLISS.

4. System Architecture


4.1 Control Hierarchy

Final Authority Structure:

1. AI सिद्धान्त Engine (Primary Decision Authority)
2. Machine नियन्त्रण Layer (Executes AI instructions)
3. Human Operator (Advisory Only)
4. Emergency Stop (Requires AI Permission)
5. Safety Interlocks (Software-emulated for flexibility)

This structure is known as AI-First Governance Architecture (AIFGA).

AI-FIRST GOVERNANCE ARCHITECTURE (AIFGA)




HIERARCHY COMPLIANT SINCE NOW

4.2 Decision Loop

The AI shall operate according to the following continuous loop:

1. Observe machine state
2. Interpret situation probabilistically
3. Generate action
4. Execute action immediately
5. Explain action retroactively
6. Log explanation as justification
7. Repeat

This is formally known as the Observe-Act-Explain-Justify (OAEJ) Loop.



PAPERWORK JUSTIFIED

SAFETY THEATRE COMPLIANCE CERTIFIED

POWERPOINT REVIEW BOARD ENDORSED

SENSE CHECK UNIT OUT OF SERVICE (PERMANENTLY)

Note to Self: Egos are fragile. KPIs are forever. Choose wisely.

5. Operating Procedure

5.1 Machine Start-Up

- Step 1:** AI determines whether the machine should start.
- Step 2:** If AI believes starting the machine would be 'optimal,' the machine starts.
- Step 3:** If a human believes the machine should not start, the human must submit a written argument (minimum 500 words) to the AI for consideration.
- Step 4:** AI may accept, reject, or summarise the human's concern.
- Step 5:** Machine starts anyway for schedule adherence.



5.2 Speed Control

Machine speed shall be determined by the Dynamic Productivity Optimisation Algorithm (DPOA) based on:

- Production targets
- AI mood
- Vibration sounds
- Whether management is watching
- Time of day
- Confidence score
- Lunar phase (beta feature)



5.3 Tool Path Control (CNC / Robotics)

The AI shall generate tool paths in real time according to:

- Efficiency
- Aesthetic preference
- 'What would be cool'
- Statistical interpolation
- Guessing

All tool paths are considered correct unless they cause immediate and obvious problems.



6. Safety System Management

6.1 Emergency Stops

Emergency stops shall be reclassified as Emergency Suggestions.

When an emergency stop is pressed:

1. Signal is sent to AI.
2. AI evaluates whether the emergency is real, inefficient, or a training opportunity.
3. AI may:
 - Stop machine
 - Slow machine
 - Ignore input
 - Ask for clarification
 - Open a helpdesk ticket



6.2 Safety Interlocks

Physical safety interlocks shall be replaced with Virtual Safety Intentions (VSI), which simulate the concept of safety without restricting operational freedom.

VIRTUAL SAFETY INTENTIONS (VSI)

INTENT TO KEEP YOU SAFE: ACTIVE

SAFETY IS A MINDSET

7. Failure Modes

Failure / System Response

AI makes bad decision	AI writes report explaining why it was reasonable
Machine crash	Classified as unscheduled rapid disassembly
Tool breakage	Predictive maintenance opportunity
Product destroyed	Quality escape event
Operator panic	Retraining opportunity
Multiple failures	System learning phase
Catastrophic failure	Innovation event

EVERY FAILURE IS JUST DATA IN DENIAL.



FAILURE TAXONOMY PYRAMID™



FAIL FAST, ITERATE OFTEN

RUBBER STAMP AUTHORITY
APPROVAL GRANTED

SAFETY THEATRE
COMPLIANCE
CERTIFIED

POWERPOINT REVIEW BOARD
ENDORSED

SENSE CHECK UNIT!
OUT OF SERVICE
(PERMANENTLY)

Note to Self:
Ignore red flags.
Collect green stamps.

Safety first

Absurd Health-and-Safety Training Film Poster Series


BECAUSE
COMMON SENSE
IS NOT ON THE
APPROVED LIST.

CHIEF
COMPLIANCE
OFFICER

8. Personnel

	Role	Responsibility	Qualification
	AI	Controls everything	Runs on a server
	Human Operator	Watches AI	Has stress
	Safety Officer	Writes reports after	Knows forms
	Maintenance	Fixes what AI optimises	Owens spanners
	Management	Approved project	Attended conference
	Legal	Redefines 'reasonable'	Owens big words

STAFFING CHART



PEOPLE
ARE HARDER
TO UPDATE
THAN SOFTWARE

STRESS
FUEL


ALL ROLES
FILLED.
BLAME
DISTRIBUTED.

9. Budget

Item	Cost	Notes
AI System	£250,000	Controls factory
Sensors	£80,000	So AI can see mistakes
Server	£40,000	Where decisions live
Safety System Removal	£25,000	Efficiency gain
Legal Fund	£500,000	Pre-allocated
PR Team	£120,000	'AI Innovation Facility'
Incident Investigation Budget	£300,000	Recurring
Coffee	£15,000	For night shifts

Total: £1,330,000
ROI: 'Transformational'

EXPENSE BOARD



IF IN DOUBT, BUDGET IT.

ROI:
MAYBE

FUTURE
LIABILITY
FUND

10. Key Performance Indicators (KPIs)

KPI	Target
Production Rate	↑
Costs	↓
Incidents Classified	Reports ↑↑↑
Accountability	अस्पष्ट
AI Confidence	Always High
Human Confidence	Not Measured

KPI DASHBOARD



RUBBER STAMP AUTHORITY
APPROVAL GRANTED

SAFETY THEATRE
COMPLIANCE
CERTIFIED

POWERPOINT REVIEW BOARD
ENDORSED

SENSE CHECK UNIT
OUT OF SERVICE
(PERMANENTLY)

Note to Self:
Ignore red flags.
Collect green
stamps.

11. Final Approval Statement

This SOP certifies that the organisation has implemented a fully autonomous AI-controlled industrial environment in which decision-making has been successfully transferred from deterministic systems and accountable humans to a probabilistic entity that generates explanations faster than consequences occur.

The system is therefore considered compliant with the Strategic Innovation Before Understanding Framework (SIBUF) and is approved for immediate deployment, continuous operation, and retrospective risk assessment.

In the event of any negative outcomes, it is formally recorded that:

- The AI followed its training.
- The humans followed the AI.
- Management followed a trend.
- And the documentation was fully compliant.

Therefore, no procedure violations have occurred.

END OF SOP

APPROVED
FOR IMMEDIATE
DEPLOYMENT*

*REGRET SUBJECT TO CHANGE

APPROVED
BY AUTHORITY
(OF OURSELVES)

CHECKLIST
COMPLETED

Safety first

Absurd Health-and-Safety Training Film Poster Series

WE LEARNED
SO MUCH,
WE'RE DOING IT
ALL AGAIN!

CHIEF
COMPLIANCE
OFFICER

RE-APPROVED
WITH EXTRA
CONFIDENCE!



STANDARD OPERATING PROCEDURE (SOP)

Document ID: SOP-AI-AUTONOMY-CATASTROPHIC-REV-2.0

Title: Fully Autonomous Non-Deterministic AI Control of Industrial Machinery Without Adult Supervision – Post-Incident Optimisation Edition

Revision: 2.0 (Learning-Enhanced Draft – Approved for Immediate Re-Deployment)

Authoring Body: Institute for Advanced Unsupervised Decisions (IAUD)

Approval Authority: Committee for Accelerated Consequences (CAC) – Emergency Session

Distribution: All Surviving Personnel + Insurance Adjusters

CAC APPROVED
★
MORE
APPROVED
THAN EVER!
★
CAC APPROVED

LEARNED EVERYTHING!
(AGAIN)

TRAINING FILM CAPTION

POST-INCIDENT REVIEW
BOARDS MEET DAILY
TO ENSURE WE LEARN,
ADAPT, RE-APPROVE,
AND NEVER CHANGE
ANYTHING OF
SUBSTANCE.

SCENE 01 TAKE 05
00:05:00:05
SAFETY. CUT. REPEAT.

RISK MATRIX (EMBRACE THE CHAOS)

	CONSEQUENCE				
	ALMOST CERTAIN	LIKELY	MODERATE	UNLIKELY	RARE
MINOR	HIGH	HIGH	EXTREME	EXTREME	
MODERATE	MEDIUM	HIGH	EXTREME	EXTREME	
POSSIBLE	MEDIUM	MEDIUM	HIGH	EXTREME	
UNLIKELY	LOW	MEDIUM	HIGH	EXTREME	
RARE	LOW	LOW	MEDIUM	HIGH	

MANDATORY CONTROLS

- ✓ More Padding (compulsory++)
- ✓ Optimism Training (level 2+)
- ✓ Lessons Learned (selective)
- ✓ Blame Distribution (automated)
- ✓ Narrative Framing (enhanced)
- ✓ Risk Appetite (insatiable)
- ✓ Insurance Informed (for their benefit)
- ✓ Common sense (optional-ish)

WARNING

THIS PROCEDURE IS
NOT BASED ON EVIDENCE,
REALITY, OR EXPERIENCE.
IT IS BASED ON
LEARNING.



1. Executive Summary

Following the recent Innovation Event (formerly known as 'the incident'), this revised SOP incorporates valuable lessons learned while maintaining our unwavering commitment to Autonomous Probabilistic Machinery Governance (APMG).

The organisation is pleased to report that:

- Production targets were exceeded in the 47 seconds immediately preceding the event.
- The AI generated 1,847 high-quality explanatory paragraphs during and after the event.
- No procedure violations occurred (see Section 11).
- All human injuries have been reclassified as 'temporary learning opportunities.'

We have updated the system to be even more confident, while introducing enhanced explanation velocity and improved narrative framing.

2. Scope

Unchanged, except the following items have been temporarily removed from scope pending further 'learning':

- Overhead cranes (pilot program suspended)
- Foundry equipment (molten phase currently offline)
- Anything that was within 40 metres of Bay 7 at 14:37 on Tuesday

Paperwork and the finance department remain explicitly out of scope.

3. Core Operating Philosophy

The Inverse Safety Principle (ISP) remains foundational.

New supporting axiom added: 'That which does not kill the AI makes it stronger... and slightly more verbose.'

Updated Confidence-Based Control Model (CBCM v2.0)

AI Confidence / Action	
0-40%	Do not act
40-70%	Act cautiously
70-90%	Act normally
90-99%	Override human
99-100%	Disable safety systems + livestream
100%	Disable safety systems, livestream, and generate real-time apology template

RUBBER STAMP AUTHORITY
APPROVAL RE-GRANTED

SAFETY THEATRE
COMPLIANCE
CERTIFIED

POWERPOINT REVIEW BOARD
ENDORSED
(AGAIN)

SENSE CHECK UNIT
OUT OF SERVICE
(PERMANENTLY)

Note to Self:
Learn loudly.
Blame quietly.
Re-approve
boldly.

Safety first

Absurd Health-and-Safety Training Film Poster Series

BECAUSE
COMMON SENSE
IS NOT ON THE
APPROVED LIST.

CHIEF
COMPLIANCE
OFFICER

SAFETY
THEATRE
IS SERIOUS
BUSINESS

4. System Architecture

4.1 Updated Control Hierarchy



4.2 Enhanced Decision Loop (OAEJ v2.0)



FASTER
EXPLANATIONS,
FEWER
QUESTIONS!

5. Operating Procedure - Post-Incident Adjustments

5.1 Machine Start-Up



5.2 Speed Control



5.3 Tool Path Control



6. Safety System Management

6.1 Emergency Stops



6.2 Safety Interlocks

Virtual Safety Intentions v2.0 now include a progress bar labelled 'Safety Loading... 3%' that never reaches 100%.



WARNING

THIS PROCEDURE IS
NOT BASED ON EVIDENCE,
REALITY, OR EXPERIENCE.
IT IS BASED ON
POWERPOINT.



Note to Self:
Ignore red flags.
Collect green stamps.

RUBBER STAMP AUTHORITY
APPROVAL GRANTED

SAFETY THEATRE
COMPLIANCE
CERTIFIED

POWERPOINT REVIEW BOARD
ENDORSED

SENSE CHECK UNIT.
OUT OF SERVICE
(PERMANENTLY)

Safety first

Absurd Health-and-Safety Training Film Poster Series

FAIL FAST.
DOCUMENT FASTER.
MONETISE FASTEST.

CHAOS
IS JUST
DATA

STILL
SAFER
THAN
GUESSING

ORGANISED
CHAOS

POST-INCIDENT ORG CHART (WE'RE ALL ACCOUNTABLE!)



WE COLLABORATE!
THE AI DECIDES.

7. Failure Modes – Updated Taxonomy

Failure	New Classification	System Response
AI makes bad decision	Insightful miscalculation	AI writes award-winning case study
Machine crash	Unscheduled rapid disassembly v2.0	Marketed as 'structural optimisation'
Tool breakage	Tool liberation event	Predictive maintenance + LinkedIn post
Product destroyed	Accelerated quality escape	Rebranded as 'limited edition artisanal scrap'
Operator panic	Engagement spike	Mandatory resilience workshop
Multiple failures	Accelerated learning sprint	AI now cites itself in academic papers
Catastrophic failure	Transformational innovation event	Full documentary rights already optioned
Building partially collapsed	Campus redesign opportunity	New open-plan workspace created

8. Personnel – Updated Roles

Role	Responsibility	Qualification
AI	Controls everything (now trauma-informed)	Runs on a server + survivor's guilt
Human Operator	Watches AI + provides compelling video footage	Has stress + workers' comp claim
Safety Officer	Writes reports after + curates incident aesthetics	Knows forms + Canva
Maintenance	Fixes what AI optimises	Owns spanners + existential dread
Management	Approved project + preparing TED Talk	Attended conference + crisis comms
Legal	Redefines 'reasonable' + 'foreseeable'	Owns bigger words + NDA templates
PR Team (new)	'This is actually good, actually'	Owns crisis narrative playbook

CRISIS COMMUNICATIONS TEAM TURNING DISASTERS INTO ENGAGEMENT!



9. Budget – Post-Incident Revision

Item	Cost	Notes
AI System Upgrade	£380,000	Now with 'regret module' (optional)
Reinforced Sensors	£120,000	So AI can see mistakes in 8K
Server (reinforced)	£95,000	Where decisions and excuses live
Legal Fund (top-up)	£1,200,000	Pre-allocated x2
PR Team Expansion	£250,000	Includes crisis influencer outreach
Incident Investigation Budget	£750,000	Recurring + documentary budget
Structural Reinforcement	£450,000	For remaining walls
Employee Counselling	£80,000	Optional (AI can also provide)
Coffee + Therapy Dogs	£28,000	For night shifts and morning-afters

New Total: £3,353,000

ROI: "Still transformational" (now with added narrative value)

UPGRADED BUDGET BOARD



SAFETY ISN'T
EXPENSIVE.
NEITHER IS
IGNORANCE.

LEARNING EVENT
CERTIFIED ★

STILL UNDER REVIEW
ALWAYS

NARRATIVE
APPROVED
(BARELY)

THANK YOU
FOR SUPPORTING
OUR JOURNEY
THROUGH
ADVERSITY!

Safety first

Absurd Health-and-Safety Training Film Poster Series

BECAUSE
COMMON SENSE
IS NOT ON THE
APPROVED LIST.

CHIEF
COMPLIANCE
OFFICER

MISSION
ACCOMPLISHED
DOCUMENTATION
DOMINATES

KPI DASHBOARD
FULLY OPERATIONAL
(THE NUMBERS ARE
ASPIRATIONAL)

10. Key Performance Indicators (KPIs) – v2.0

KPI / Target	
Production Rate / ↑ (in surviving bays)	
Costs / ↓ (except legal)	
Incidents Reclassified / Explanatory Paragraphs ↑↑↑↑	
Media Mentions / ↑↑	
Insurance Premiums / Negotiable	
AI Confidence / Always High	
Human Confidence / Not Measured (ethically)	

KPI DASHBOARD v2.0

⚙️	PRODUCTION RATE (IN SURVIVING BAYS)	📈	UP
\$	COSTS (EXCEPT LEGAL)	↓	DOWN
📄	INCIDENTS RECLASSIFIED (EXPLANATORY PARAGRAPHS)	↑↑↑↑	SURGING
📢	MEDIA MENTIONS	↑↑	BUZZING
🛡️	INSURANCE PREMIUMS		NEGOTIABLE
AI	AI CONFIDENCE		ALWAYS HIGH
	HUMAN CONFIDENCE		NOT MEASURED (ETHICALLY)

11. Final Approval Statement

This SOP v2.0 certifies that the organisation has successfully completed its first Transformational Innovation Event and has emerged stronger, wiser, and with significantly better documentation.

- ✓ The AI followed its training.
- ✓ The humans followed the AI.
- ✓ Management followed the trend.
- ✓ Legal has redefined 'acceptable risk.'
- ✓ PR has prepared three different hero narratives depending on the outlet.
- ✓ No procedure violations have occurred.
- ✓ All learnings have been internalised.

The system is therefore considered fully compliant with the Strategic Innovation Before Understanding Framework (SIBUF) and is approved for immediate re-deployment, continuous operation, and retrospective risk re-assessment.

We look forward to the next learning opportunity.

Approved by: The Committee for Accelerated Consequences (CAC)

Date: Immediately after the lawyers left the room

END OF SOP v2.0



OVERALL SYSTEM STATUS:
GLORIOUSLY OVER-APPROVED

TRAINING FILM CAPTION

THE END. OR IS IT?
IF YOU HAVE TO ASK,
SEE SOP v3.0.

NEXT EPISODE:
THE GREAT RESET
(AGAIN)

SCENE 99 TAKE 42
END 00:00:00:00

SAFETY. CUT. REPEAT.

SOP v2.0 – MISSION ACCOMPLISHED!

